

Frozen methodology on a new backbone — NVDA VoiceChat-11B
same three reads, logistic C=1e-4; calib n=600 (vs MiniCPM's 2310)

