


(a) Accuracy on our eval splits

solid = calib OOF, faded = test, dashed = majority baseline



(b) Training vs eval loss

open = train, filled = calib OOF, square = test

