

Accuracy vs escalation rate — SD-QA — real human speech (VoiceBench, n=200)
gate + thresholds frozen from the 600-query calibration pool

