
When Does a Small Model Know to Hand Off? Reading Competence from Full-Duplex Speech Models for Real-Time Escalation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Full-duplex speech models answer in real time with a small on-device backbone,
2 yet many queries exceed its competence. Routing hard queries to a large cloud
3 model sounds solved — until one asks *when* the decision must be made: be-
4 fore the first output token, or the wrong answer is already on the air. We ask
5 whether the model’s own hidden state carries that warning early enough, and find
6 that the obvious readout is exactly what duplex fine-tuning corrupts: final-layer
7 probes collapse under distribution shift on text input (leave-one-pool-out math
8 AUC 0.93→0.37), consistent with late-layer specialization for duplex turn con-
9 trol, while verbalized self-assessment collapses on audio input. The signal is not
10 gone — it sits one layer band earlier. Mid-network, the same linear probe transfers
11 at 0.93, reads speech input from text-only calibration at 0.86, and an end-of-turn
12 read in the streaming loop reaches AUC 0.887 in ~30 ms — half a second before
13 the first audible word. Run live end-to-end in a turn-based loop, the gated system
14 climbs from 42.2% to 63.3% heard accuracy at 57% escalation, the remaining gap
15 dominated by the transcription uplink rather than the gate; frozen, the same gate
16 improves five of six public speech benchmarks, lifting Speech TriviaQA to 86.0%
17 live against the model’s official offline 75.5%. Duplex training breaks the obvious
18 competence readout; reading one band earlier, before speech commits, is enough
19 to escalate live.

20 1 Introduction

21 A full-duplex speech model listens and speaks in the same forward pass, under a conversational
22 deadline of a few hundred milliseconds [Stivers et al., 2009]. That deadline forces a small on-device
23 backbone — and a small backbone routinely faces queries beyond its competence: long-tail facts,
24 expert knowledge, multi-step reasoning. The natural architecture is selective escalation to a large
25 cloud model, but the duplex setting imposes two constraints that existing routing work does not
26 meet: the decision must come *before the first output token* (once the model starts speaking, a wrong
27 answer is already on the air), and the session — audio context, voice, cache — must stay with the
28 local model, which will speak the expert’s answer itself. The question is not whether one can build
29 a router; it is whether the small model knows it will fail *early enough to do anything about it*.

30 To answer it, we study MiniCPM-o 4.5 [Cui et al., 2026], a 9B omni/duplex fine-tune of Qwen3-8B,
31 against matched-pair controls spanning raw backbones, streaming-omni, audio-only, and frozen-
32 adapter fine-tunes, and evaluate every standard zero-shot competence readout — decode-time en-




Figure 1: **(1)** Speech streams live into a small duplex-trained local model. **(2)** A mid-network (L22) linear read decides in ~ 30 ms — before the first output token — so the expert call is already in flight while local decoding would just be starting. **(3)** Easy queries stay local; hard ones escalate to a cloud expert. **(4)** The *same talker speaks* the expert’s answer in one continuous session, improving live spoken sessions (42.2%→63.3%) and, frozen, public speech benchmarks (Speech TriviaQA 66.4%→86.0%).

33 tropy, verbalized self-assessment ($p(\text{True})$), and linear probes on hidden states — with judged fail-
 34 ure labels on a frozen pool of public-benchmark queries, in matched text and speech conditions.
 35 Throughout, *duplex* names the training regime, not the interaction protocol: we study duplex-trained
 36 weights driven through a *turn-based* loop whose only decision point is an end-of-turn read (§6);
 37 talker-on replay and concurrent listen/speak validations: Appendix G.

38 Everything runs on a frozen checkpoint — no fine-tuning; the only fitted component is a logistic
 39 regression on stored hidden states, calibrated once on text. That minimalism is the point: competing
 40 signals read the wrong place (late layers), arrive too late (post-answer verification), or must be
 41 trained in (think tokens); the mid-layer read survives all three. Three findings — find the signal,
 42 read it in time, show it matters:

- 43 1. **Representation: duplex fine-tuning makes the late-layer readouts brittle; a mid-**
 44 **network signal survives.** The conventional final-layer probe looks strong in distribution
 45 (AUC 0.82) but is mostly a query-type shortcut: under leave-one-pool-out transfer it *inverts*
 46 on held-out math (0.93→0.37). The damage is specific to the late-layer last-token position
 47 on text input — absent in raw backbones and non-duplex fine-tunes — consistent with late-
 48 layer specialization for duplex turn control. Mid-network (L22 of 36) the same linear read
 49 transfers at 0.93, and the mid-network optimum generalizes to a second, architecturally
 50 disjoint duplex family (hybrid-Mamba NemotronLabs-VoiceChat-11B).
- 51 2. **Modality: the mid-layer signal transfers text → speech in end-to-end models.** A probe
 52 calibrated on cheap text data reads speech-input hidden states at 0.86; the control that iso-
 53 lates the mechanism is Freeze-Omni, an adapter on frozen identical weights, where cross-
 54 modal transfer is dead (≤ 0.60). Verbalized self-assessment shows the mirror-image failure
 55 — it collapses on audio input — so the mid-layer probe is the one readout that survives both
 56 modalities. The transfer holds on real human recordings, not just TTS.
- 57 3. **System: the signal is available before the first token and works live, in and out of**
 58 **distribution.** The L22 read costs 20 ms (text) / 45 ms (audio) against time-to-first-token
 59 of 36/68 ms; in the streaming loop an end-of-turn read reaches AUC 0.887 in ~ 30 ms,
 60 half a second before first audible speech. Run live end-to-end — expert consulted mid-
 61 conversation, answer injected, talker speaks it — heard accuracy climbs 42.2%→63.3%
 62 at 57% escalation, and the frozen gate transfers to six public speech benchmarks with
 63 label-free thresholds, lifting Speech TriviaQA to 86.0% against the model’s official offline
 64 75.5%.

65 Extended analyses and controls are in the appendix.

66 2 Related Work

67 **Full-duplex real-time dialogue.** End-to-end spoken dialogue spans native duplex designs —
 68 dGSLM [Nguyen et al., 2023], Moshi [Défossez et al., 2024], GLM-4-Voice [Zeng et al., 2024],
 69 single-LLM native full-duplex [Fang et al., 2026] — streaming omni models [Xu et al., 2025a, Cui
 70 et al., 2026], frozen-backbone adapters [Wang et al., 2024a], and audio fine-tunes [Chu et al., 2024].
 71 A parallel line adds explicit turn-taking control: FlexDuo [Liao et al., 2025], MinMo [Chen et al.,
 72 2025], and Freeze-Omni’s state heads decide *when the local model should speak*, at 120–250 ms
 73 cadences that show a live loop tolerates far more decision latency than our ~ 30 ms read [cf. Riera
 74 et al., 2026, Wu et al., 2026]. Our gate decides the orthogonal question — *whether the local model*
 75 *should be the one answering* — and the head-based controllers’ shared choice of a last-layer readout
 76 is exactly the point our results show duplex fine-tuning damages.

77 **Does the model know what it knows?** $p(\text{True})$ self-evaluation [Kadavath et al., 2022], verbalized
 78 uncertainty [Lin et al., 2022], and truthfulness probes on hidden states [Azaria and Mitchell, 2023,
 79 Burns et al., 2023, Marks and Tegmark, 2024] established that correctness is readable internally;
 80 intermediate layers carry more of it than the output distribution [Orgad et al., 2025]. Semantic
 81 entropy [Kuhn et al., 2023, Farquhar et al., 2024] needs multiple sampled generations; Semantic
 82 Entropy Probes [Kossen et al., 2024] amortize it into a hidden-state probe — the closest signal to
 83 ours, but read post-hoc on text-only half-duplex models. Multimodal uncertainty fusion [Zhang
 84 et al., 2026] asks a downstream question; we ask *where* in a duplex model a modality-robust read
 85 survives at all.

86 **Selective escalation and routing.** The classic routers are creatures of the turn-based, text-only
 87 setting: FrugalGPT [Chen et al., 2023], HybridLLM [Ding et al., 2024], RouteLLM [Ong et al.,
 88 2024], and confidence-based cascades [Xue et al., 2025, Hu et al., 2024, Li et al., 2026] make one
 89 decision per complete typed query, from query-side features or preference data, before any model
 90 has answered; AutoMix [Aggarwal et al., 2024] escalates by *post-answer* self-verification, which
 91 a voice deadline forbids. Both assumptions — query-side features, and the large model answering
 92 the user directly — fail in full-duplex voice, where the session lives in the small model (query-side
 93 ceiling quantified in §4.1). A second family trains the routing decision *into the weights* as special
 94 tokens emitted during decoding: Toolformer [Schick et al., 2023] teaches the model to insert API-
 95 call tokens by self-supervised fine-tuning; Self-RAG [Asai et al., 2024] fine-tunes the full model
 96 on distilled critic labels so reflection tokens trigger retrieval and self-critique; Co-LLM [Shen et al.,
 97 2024] and CITER [Zheng et al., 2025] learn a per-token deferral decision to a larger model; pause
 98 and thought tokens [Goyal et al., 2024, Zelikman et al., 2024], fused think/no-think modes [Yang
 99 et al., 2025, DeepSeek-AI et al., 2025], and the built-in think gate MiniCPM-o 4.5 itself ships [Cui
 100 et al., 2026] route *compute* the same way. The design is uniform — extend the vocabulary, then
 101 fine-tune the generator until the policy lives in the checkpoint — and so is the cost: supervision
 102 must be constructed and the talker fine-tuned, the policy is frozen at training time (a new expert,
 103 budget, or threshold means retraining), and in a duplex model fine-tuning the talker is exactly the
 104 intervention §4.2 shows damages the last-layer readout. Our gate is the pluggable complement: a
 105 logistic read over a frozen checkpoint, refit from stored activations in minutes, its budget moved
 106 by a deployment-time threshold, transferring across modality and model family (§7) — and it beats
 107 the model’s own trained think token at matched budget (Table 9). Learning-to-defer [Madras et al.,
 108 2018] and early exit [Schuster et al., 2022] decide inside one model.

109 **The gap.** Duplex work decides when to speak; routing work decides which model — from the
 110 query surface, after a first answer, or via tokens trained into the weights; uncertainty work asks
 111 whether the model is confident, almost always offline and in text. No prior work we know of asks
 112 whether a full-duplex speech model *already* exposes a pre-answer, modality-robust, internal compe-
 113 tence signal early enough to drive a live small-to-large handoff while preserving the session.

3 Experimental Setup

Target model and matched-pair controls. The system target is **MiniCPM-o 4.5** (9B; omni + full-duplex fine-tune of Qwen3-8B), bf16, greedy decoding, seed 42. To attribute effects to the *kind* of fine-tune rather than the backbone, we assemble ten models spanning raw backbones, a second duplex generation (MiniCPM-o 2.6), a streaming-omni fine-tune (Qwen2.5-Omni-7B), an audio-only fine-tune (Qwen2-Audio-7B), and a frozen-backbone adapter model (Freeze-Omni, LLM weights byte-identical to its raw control); the full roster is Table 3 (appendix). The comparison statistic is always the within-pair difference.

Query pool, labels, and audio arm. 600 queries in five pools, frozen before any gate experiment (360 calibration / 240 test): *trap* = SimpleQA [Wei et al., 2024] (the target fails 100%), *hard-knowledge* = MMLU-Pro [Wang et al., 2024b], *easy-fact* = TriviaQA [Joshi et al., 2017], *easy-chat* = Dolly/Alpaca [Conover et al., 2023, Taori et al., 2023], *hard-math* = GSM8K tail + MATH-500 [Cobbe et al., 2021, Hendrycks et al., 2021, Lightman et al., 2024]; public benchmarks only, no LLM-generated queries. Adequacy labels come from an LLM judge [Zheng et al., 2023] with structured outputs, blind to the gate and validated by a stronger re-judge (agreement 96.2% text / 94.5% audio; every headline number moves ≤ 1 point under the stronger labels); the escalation expert is gpt-5.5 at low reasoning effort. The audio arm renders the same frozen pool to speech (single-voice TTS, content unchanged — only modality varies), following Spoken-SQuAD / VoiceBench practice [Li et al., 2018, Chen et al., 2024]; audio-side findings are additionally validated on real human recordings (SD-QA, Faisal et al., 2021) and on the external benchmarks’ official audio (§7). Judge prompts, pins, and spend: Appendix L–M.

Gate signals (no model training). (i) decode scalars: max token entropy / margin over the first k decode steps; (ii) linear probes: logistic regression on a single hidden state, position $\times$ layer swept, fit on the calibration split only; (iii) verbalized self-evaluation [Kadavath et al., 2022]: $p(\text{True})$ -pre asks “would you answer this correctly?” before answering, $p(\text{True})$ -post asks about a drafted answer; $P(\text{Yes})$ is read from the first token’s distribution.

4 Reading Competence from Duplex Models

4.1 In distribution every readout looks fine; transfer reveals the shortcut

Table 1 compares the candidate signals on MiniCPM-o 4.5. In distribution the final-layer probe reaches 0.822 — but that number is mostly a *query-type* shortcut: the same hidden state classifies a query’s source pool at 95.8%, a pool-identity oracle alone reaches AUC 0.715, and appending true pool labels to the probe adds nothing (0.822 $\rightarrow$ 0.821). The decisive test is leave-one-pool-out (LOPO) transfer, which we treat as the primary metric throughout: held-out math *inverts* to 0.372, and the mean escalation score on a 100%-fail trap pool is 0.232. An external query-surface router makes the same point from the other side: under identical labels a TF-IDF router reaches only OOF AUC 0.669 (below the pool oracle), collapses to 0.38–0.57 under LOPO, and 100 $\times$ the data on a public routing benchmark buys back only 0.04 of AUC (Appendix D). Instance-level competence information lives in the model’s internal state, not on the query surface.

Verbalized self-assessment transfers differently: *pre-answer* $p(\text{True})$ catches the trap pool the probe misses (escalation score 0.945 vs. the probe’s 0.232) — but timing is essential: asked *after* drafting, the model endorses its own confident-wrong answer (0.360). Post-draft $p(\text{True})$ is the strongest single signal (0.899) but requires paying for a full draft, which a voice deadline cannot afford.

4.2 The damage is duplex-specific; the signal survives mid-network

Matched pairs attribute the probe failure to the *kind* of fine-tune, not the backbone: the raw Qwen backbone, label-matched to the duplex model’s fail rates, holds LOPO math at 0.825 where the duplex model sits at 0.372; streaming-omni and audio-only fine-tunes are undamaged; a second duplex

Table 1: Signals on MiniCPM-o 4.5, judged failure labels. In-dist = calibration-split AUC (probes 5-fold out-of-fold); LOPO math = leave-one-pool-out AUC on held-out math; audio = the same readout under speech input (final/mid-layer probe: audio-input LOPO math / text-only calibration read on audio states). Full per-signal detail: Appendix B.

signal	in-dist	LOPO math	audio input
max entropy@4	0.696	$\approx$ chance on raw backbones	—
$p(\text{True})$ -pre	0.807	trap 0.945 (probe: 0.232)	collapses
$p(\text{True})$ -post (full draft)	0.899	—	degrades
final-layer probe	0.822	0.372 (inverts)	0.936 (intact)
mid-layer (L22) probe	0.866	0.931	0.855 (text-calibrated)



Figure 2: The core representational result: LOPO math AUC by layer and position. On the duplex model the standard (final-layer, last-token) read inverts on text input while mid-network layers hold ≥ 0.9 ; raw backbones show no cliff. The deployed gate reads L22.

160 generation (MiniCPM-o 2.6) reproduces the collapse (Appendix B). A layer $\times$ position sweep local-
161 izes it (Figure 2): the collapse is specific to (late layer $\times$ last token) on text input — mean-pooled
162 reads survive at the same depth but 0.13–0.15 below the mid-layer peak (LOPO math ≈ 0.80 here,
163 0.674 on MiniCPM-o 2.6), and mid-network the last-token read recovers to **0.931 at L22**, near the
164 raw backbone. This is the position a streaming head must re-purpose to encode turn-control state,
165 and the pattern is consistent with late-layer specialization for duplex turn control overwriting the
166 readout while the underlying information survives. The mid-network optimum is not a MiniCPM
167 quirk: on NemotronLabs-VoiceChat-11B — a hybrid-Mamba duplex model, architecturally disjoint
168 — the same recipe peaks mid-band at L30–34 of 56 and transfers to the external pools at AUC
169 0.70–0.79 with a quarter of the calibration data (Appendix C).

170 4.3 Text-calibrated, speech-deployed — in end-to-end models only

171 Probes trained on text hidden states and scored on speech-input hidden states locate where the repre-
172 sentation becomes modality-shared: early layers are modality-specific (0.54–0.60), and from $\sim 50\%$
173 depth text $\rightarrow$ audio transfer plateaus at 0.82–0.87 (0.855 at the deployed L22). This replicates on
174 MiniCPM-o 2.6 and cross-family on Qwen2.5-Omni. The decisive control is **Freeze-Omni**: identi-
175 cal frozen LLM weights with audio attached by an adapter — and cross-modal transfer is dead
176 (0.52–0.60). The modality-shared mid-layer core is *created by training the backbone on the modal-*
177 *ity*; the cheap recipe “calibrate on text, deploy on speech” is valid for end-to-end models and invalid
178 for adapter-attached ones. Verbalized introspection shows the mirror-image failure: on audio input,
179 pre-answer $p(\text{True})$ collapses on both MiniCPM duplex generations while both non-duplex controls
180 stay intact; controlled arms show the verbal self-check runs over text-token pathways that audio em-
181 beddings do not feed — duplicating the question as text restores it — while a linear probe reads the
182 same information off the same forward pass at 0.93 (Appendix B). All audio-side findings replicate
183 on 200 real human recordings (SD-QA), ruling out TTS artifacts. **The mid-layer probe is the one**
184 **readout that survives every training regime that can converse**; the live system deploys the same
185 L22 read plus two zero-latency pooled positions (the *deployed* probe; Appendix B).

5 From Signal to Gate

The decision exists before decoding starts. The gate is one 4096-d dot product on a hidden state the forward pass already computes. Measured on H100 (wall-clock, prefill start to L22 score, CUDA-synced; network excluded — the cloud call fires asynchronously; protocol: Appendix E), the decision lands in **20 ms (text) / 45 ms (audio)** at P50 (P95 25/104) against same-bench time-to-first-token of 36/68 ms: $T_{\text{gate}} < T_{\text{first token}}$, far inside the 200–300 ms turn-taking budget [Stivers et al., 2009], and the cloud call can launch while the remaining $\sim 40\%$ of prefill and the whole decode still run. Deciding pre-token is optimal, not merely feasible: first audible speech trails the first token by ~ 0.5 s (§6), so one could wait out a few decoded tokens and cancel — but decode-time readouts measure *weaker*, not stronger (max entropy@4 0.696, Table 1; mean-pooled early-decode hidden 0.776 vs. 0.828 at the prompt position; $p(\text{True})$ -pre an extra prefill, $p(\text{True})$ -post a full 1.9/3.5 s draft), and every waited millisecond delays the stall onset and the expert launch on a path where the expert round-trip is 85–88% of escalated wall time (per-layer cost curves: Appendix E).

Accuracy–cost tradeoff. On the frozen test split ($n=240$): small-only 58.8%, big-only (gpt-5.5) 91.7%. Gated escalation beats random escalation at every operating point (Figure 6, appendix); the balanced tier at a 33% budget reaches **77.9–78.7%**, recovering $\sim 58\%$ of the small→big gap at one third of the big-only cost. In distribution the gate is statistically on par with a pool-oracle (type-only) router — as the type-shortcut audit predicts — so its case over type routing rests on transfer, where the type shortcut collapses (§4.1) and the mid-layer signal does not. The gate is fit to *local failure*, not expert benefit ($y=1$ iff local wrong *and* expert right): a benefit label bakes the current expert’s blind spots into a router meant to survive expert swaps; against the harder label the frozen scores stay usable out of distribution (AUC 0.63–0.81), a fifth to three fifths of the random-to-oracle span at matched budgets (Appendix H).

Against the model’s own thinking mode. MiniCPM-o 4.5 ships a built-in reasoning mode — self-escalation in place. Gated-cloud escalation dominates all-THINK on both accuracy and latency at every escalation rate (gated-cloud@33%: 78.7% at mean 6.5 s vs. all-THINK 63.7% at 22.4 s; Table 9). The reason is mechanistic: the gate predominantly flags knowledge and trap failures, which extra compute cannot fix — thinking helps math and hurts chat — so external escalation is necessary, not merely better.

6 Live End-of-Turn Escalation

The streaming loop. The protocol is turn-based: audio streams in 1 s chunks, but the talker speaks only after the user’s turn ends and the model’s native listen/speak head is not engaged. Per-chunk probe scores are strictly weaker than the offline read (best chunk statistic AUC 0.764) and miss short utterances whose critical entity lands in the final chunk. The fix is an **end-of-turn read**: after the last audio chunk, prefill an assistant turn containing a single space and read L22 at the assistant-start position. This read reaches **AUC 0.887** — above the offline reference (0.843) — in ~ 30 ms, still before the first output token, and is the loop’s only decision point; escalation-rate thresholds are quantiles of this single score. Speculative mid-stream prefetch was measured and dropped (an expert answer to a partial transcript answers the full question only 7–51% of the time; Appendix F). On escalation the talker speaks a stall phrase, its *own* transcription of the audio goes to the expert (gpt-5.5, reasoning effort fixed low — higher effort buys zero accuracy on the escalated population at double the latency), and the returned answer is injected as context for the talker to relay. On the deployed channel the relay is clean: model wrong, correct answer injected, comply = 1.00 ($n=24$); more broadly the competence signal also predicts relay compliance (comply 42%→75% across gate-score quartiles) — a model that knows it doesn’t know is willing to listen. Conflicting-injection and framing controls: Appendix F.

The live result. We ran 240 frozen test queries $\times$ four escalation tiers, each a complete turn-based spoken session, scored on *heard accuracy* — what a user would actually have heard. Headline num-

bers use a pre-registered input-side filter (21 formula-LaTeX ids + 1 broken recording removed; spoken formulas are destroyed by any TTS render before the system hears them). Heard accuracy climbs monotonically, 42.2% → 52.8% → 59.2% → 63.3% at 0/14/31/57% escalation, every tier’s gain significant (paired bootstrap; Figure 11). A measured always-escalate arm closes the curve at 66.5% (P50 4.2 s): the last 43 points of escalation buy 3.2 points, because past the gated tiers the binding constraint is the transcription-uplink channel, not selection — the same sessions re-scored with gold-text expert answers reach 92.2%. The remaining gap to the channel-controlled counterfactual (escalated rows re-scored with gold-text expert answers) is dominated by *transcription-uplink* errors, not gate failure: the channel costs +2.3/+6.9/+12.8 points across tiers, and in the channel-controlled view the gate clears its paired random reference at every arm. Live arm accuracies on this pool carry a replication-noise floor of ± 2 –3 points, so no claim here rests on a sub-3-point live delta. Full tables (both views, both pools), latency profiles, uplink diagnostics, and two boundary query classes (false-premise, real-time) are in Appendix F. A talker-on replay (237/239 turns barge-in-seeded) leaves the frozen probe intact: AUC 0.871/0.856 vs. 0.866 headless (Appendix G). A TTS-on re-run of the balanced arm puts first *audible* audio at p50 0.55 s / p99 0.64 s after end of turn on local answers and 22–29 ms on escalated ones — the canned stall starts speaking while the expert call is in flight. First audio is not the answer: expert *content* becomes audible at p50 6.4 s / p95 26 / p99 34 (the round-trip is 85–88% of escalated wall time); the stall fully covers 25/84 of those waits, leaving dead air of p50 2.5 s / p95 22.4 (Appendix F).

7 Frozen Transfer to Public Speech Benchmarks

The out-of-distribution test the type-shortcut audit demands: we take the deployed gate *frozen* — weights, features, and layer choice fit once on the internal calibration mix — and run the complete live loop of §6 on six public speech benchmarks: Speech TriviaQA, Speech Web Questions, Llama Questions, and Reasoning QA from OpenAudioBench [Li et al., 2025] (spoken renderings of Joshi et al., 2017, Berant et al., 2013, and the set of Nachmani et al., 2024), SD-QA (real dialectal human speech) [Faisal et al., 2021], and VoiceBench AlpacaEval [Chen et al., 2024, Taori et al., 2023]. All audio is the benchmarks’ official pre-rendered audio; no query ever touches probe calibration. The only per-pool adaptation is *label-free*: tier thresholds are quantiles of the probe’s own score distribution on each pool, consuming no labels and no judge calls. It buys budget control, not discrimination: one absolute threshold frozen on the internal pool still beats matched-rate random on four of five pools, but fires at 2–48% across them (Appendix H). Where an official protocol exists we score with it (the OpenAudioBench judge; VoiceBench’s 1–5 judge), and our chat-mode control reproduces the official anchors on-scale.

The frozen gate improves all five QA benchmarks (Table 2), monotonically in budget. On Speech TriviaQA the 9B system reaches **86.0% live** — against its own official offline 75.5% and 62.9% for Qwen3-Omni-30B [Xu et al., 2025b] at $3.3\times$ the parameters: selective cloud escalation outperforms larger standalone backbones. The cleanest evidence that the signal is difficulty rather than surface is cross-lingual: an English-only calibration expansion moved *Chinese* Reasoning QA from AUC 0.621 to 0.683, the largest per-pool gain. On Llama Questions — single-type, so no shortcut to exploit — **selective escalation matches the expert at half the expert calls**: the probe’s split is decisive (kept-local 97.6% vs. escalated 69.6%, $z=6.46$), the expert’s whole advantage lands where the probe sends it (69.6% → 88.8%, $p<.0001$), and a measured live always-escalate arm lands *below* selective (93.6 vs. 94.8 at half the calls); matched-rate precision reaches 95% of the base-rate cap (Appendix I).

Two honest negatives. **AlpacaEval**: the gate does not beat random — open-ended instruction following has no “missing fact” event for a retrieval-failure detector to read. **Web Questions and SD-QA are flat** under the deployed-probe refit: on Web Questions strict alias-based labels decouple “wrong” from “uncertain” (even the gold-text ceiling is 86.4%), and the failures the gate ranks highest are ones the expert misses too (expert-fixable share 66% at the 15% cut vs. 75% pool-wide). Neither negative is intrinsic to the signal: on the second duplex family, where local floors are 0.31–0.38, the same frozen recipe beats matched-rate random on *all five* pools — WebQ, SD-QA and AlpacaEval

Table 2: Main transfer results: methods $\times$ five public speech QA benchmarks (TriviaQA/WebQ/Llama Q. judged by OAB official, $n = 250$; SD-QA/Reasoning zh reference-anchored, $n = 200/202$). Top: live deployed loop, heard accuracy in % (± 2 –3 point replication floor). Middle: gold-inject view at a matched 30% budget (Appendix H). The second duplex family (NemotronLabs-VoiceChat-11B, same frozen recipe) is Table 7. Δ = average gain over always-local (points). AlpacaEval (1–5 scale): 3.99 $\rightarrow$ 4.35 gated, 4.76 always-escalate, no gain over random on MiniCPM; the second family does gain (3.55 $\rightarrow$ 4.46 vs. random 4.23).

	TriviaQA	WebQ	Llama Q.	SD-QA	Reason. zh	avg	Δ
<i>live heard view (deployed loop)</i>							
always-local	66.4	57.2	84.0	51.5	58.9	63.6	—
+ gate @15%	73.2	62.8	90.4	60.0	65.3	70.3	+6.7
+ gate @30%	80.0	68.0	90.4	72.0	71.3	76.3	+12.7
+ gate @50%	86.0	73.2	94.8	79.5	77.2	82.1	+18.5
always-escalate (measured)	88.0	80.8	93.6	88.5	83.7	86.9	+23.3
always-expert (gold-text ceiling)	96.8	86.4	92.8	93.0	87.1	91.2	+27.6
<i>gold-inject view, 30% budget</i>							
random reference	75.5	65.7	86.4	63.9	67.4	71.8	+8.2
gate	83.6	70.0	91.2	74.5	72.8	78.4	+14.8
benefit oracle	96.4	86.8	94.8	81.5	89.1	89.7	+26.1
probe AUC (frozen)	0.789	0.785	0.806	0.792	0.683	0.771	—
official offline	75.5	70.2	—	—	—	—	—

included ($p \leq .0003$; Appendix C) — so the flat pools here are a headroom and failure-species property of the pool $\times$ backbone pairing. The external channel cost is small (heard-vs-gold gap 3.5–5.6 points at aggressive vs. 17.5 internally), and SD-QA’s escalated subset converts at 82–84% on real dialectal speech. Routing can also *reduce* tail latency, replacing slow wrong local decodes with short expert round-trips (Reasoning QA P90 13.3 $\rightarrow$ 10.9 s; Appendix I).

8 Discussion

What the signal detects. The pre-answer probe is specifically a *retrieval-failure* detector, strongest on missing-knowledge and long-tail-trap failures where the empty lookup is an observable event in the model’s state. Execution failures (getting stuck mid-derivation) surface in *decode-time* signals instead, and metacognitive failures — questions with false premises, where no failure event occurs in the model’s experience — are invisible to every signal we tested, query-surface routers included (fire rate 18% vs. trap’s 90%; Appendix F). This scope motivates a decode-time check behind the pre-answer gate.

Limitations. (i) One phrasing per $p(\text{True})$ variant. (ii) Real-speech validation: scripted read speech, one dialect. (iii) Formula-symbolic content is out of scope for the audio arm (destroyed by any TTS render — an input-side limit, distinct from the transcription uplink). (iv) The second-family replication covers mid-band readability and frozen-recipe transfer, not matched-pair attribution or the live loop. (v) Labels come from one judge family, bounded by a stronger re-judge (deltas ≤ 1 point), without human labels. (vi) The live sweep is one session per query per tier (± 2 –3 point replication floor); under concurrent prefill the frozen read degrades (AUC .87 $\rightarrow$.76); an in-regime refit recovers .82 — calibration must follow the regime (Appendix G); the official duplex serving format is untested.

9 Conclusion

Duplex-trained speech models carry a competence signal their standard readouts miss: a mid-network linear read that survives the readout damage, transfers from text calibration to speech, and lands half a second before the first audible word. Deployed live it turns a 9B talker into a system that knows when to hand off; frozen it lifts public speech benchmarks past models three times its size. The signal was never gone — only misread; one band earlier, it arrives in time to act.

References

- Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqui, and Mausam. AutoMix: Automatically mixing language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2310.12963.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.11511.
- Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. arXiv:2304.13734.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from question-answer pairs. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2013. WebQuestions.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2212.03827.
- Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.
- Qian Chen, Yafeng Chen, Yanni Chen, et al. MinMo: A multimodal large language model for seamless voice interaction. *arXiv preprint arXiv:2501.06282*, 2025.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench: Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024.
- Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, et al. Qwen2-Audio technical report. *arXiv preprint arXiv:2407.10759*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. GSM8K.
- Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick Wendell, Matei Zaharia, and Reynold Xin. Free dolly: Introducing the world’s first truly open instruction-tuned LLM. Databricks blog, 2023. databricks-dolly-15k.
- Junbo Cui, Bokai Xu, Chongyi Wang, Tianyu Yu, Weiyue Sun, Yingjing Xu, et al. MiniCPM-o 4.5: Towards real-time full-duplex omni-modal interaction. *arXiv preprint arXiv:2604.27393*, 2026.
- DeepSeek-AI, Daya Guo, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*, 2024.
- Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks V. S. Lakshmanan, and Ahmed Hassan Awadallah. Hybrid LLM: Cost-efficient and quality-aware query routing. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2404.14618.
- Fahim Faisal, Sharlina Keshava, Md Mahfuz ibn Alam, and Antonios Anastasopoulos. SD-QA: Spoken dialectal question answering for the real world. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, 2021. arXiv:2109.12072.

356 Qingkai Fang, Shoutao Guo, and Yang Feng. BayLing-Duplex: Native full-duplex speech dialogue
357 with a single autoregressive LLM. *arXiv preprint arXiv:2606.14528*, 2026.

358 Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large
359 language models using semantic entropy. *Nature*, 630:625–630, 2024.

360 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
361 Nagarajan. Think before you speak: Training language models with pause tokens. In *ICLR*, 2024.

362 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
363 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
364 *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks*, 2021.
365 arXiv:2103.03874.

366 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
367 Keutzer, and Shriyash Kaustubh Upadhyay. RouterBench: A benchmark for multi-LLM routing
368 systems. *arXiv preprint arXiv:2403.12031*, 2024.

369 Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale dis-
370 tantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual
371 Meeting of the Association for Computational Linguistics (ACL)*, 2017.

372 Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez,
373 Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language mod-
374 els (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

375 Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. Se-
376 mantic entropy probes: Robust and cheap hallucination detection in LLMs. *arXiv preprint
377 arXiv:2406.15927*, 2024.

378 Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for
379 uncertainty estimation in natural language generation. In *International Conference on Learning
380 Representations (ICLR)*, 2023. arXiv:2302.09664.

381 Chia-Hsuan Li, Szu-Lin Wu, Chi-Liang Liu, and Hung-yi Lee. Spoken SQuAD: A study of miti-
382 gating the impact of speech recognition errors on listening comprehension. In *Interspeech*, 2018.
383 arXiv:1804.00320.

384 Hao Li, Yiqun Zhang, Zhaoyan Guo, Chenxu Wang, Shengji Tang, Qiaosheng Zhang, Yang Chen,
385 Biqing Qi, Peng Ye, Lei Bai, Zhen Wang, and Shuyue Hu. LLMRouterBench: A massive bench-
386 mark and unified framework for LLM routing. *arXiv preprint arXiv:2601.07206*, 2026.

387 Tianpeng Li, Jun Liu, Tao Zhang, Yuanbo Fang, Da Pan, Mingrui Wang, et al. Baichuan-Audio:
388 A unified framework for end-to-end speech interaction. *arXiv preprint arXiv:2502.17239*, 2025.
389 Introduces OpenAudioBench.

390 Borui Liao, Yulong Xu, Jiao Ou, Kaiyuan Yang, Weihua Jian, Pengfei Wan, and Di Zhang. Flex-
391 Duo: A pluggable system for enabling full-duplex capabilities in speech dialogue systems. *arXiv
392 preprint arXiv:2502.13472*, 2025.

393 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
394 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International
395 Conference on Learning Representations (ICLR)*, 2024. arXiv:2305.20050; source of the MATH-
396 500 subset.

397 Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in
398 words. *Transactions on Machine Learning Research*, 2022. arXiv:2205.14334.

David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and accuracy by learning to defer. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2310.06824.

Eliya Nachmani, Alon Levkovitch, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh Mariooryad, Ehud Rivlin, RJ Skerry-Ryan, and Michelle Tadmor Ramanovich. Spoken question answering and speech continuation using spectrogram-powered LLM. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2305.15255; Spectron, source of the Llama Questions dataset.

Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden Tomasello, Robin Algayres, Benoît Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux. Generative spoken dialogue language modeling. *Transactions of the Association for Computational Linguistics*, 2023. arXiv:2203.16502.

Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data. *arXiv preprint arXiv:2406.18665*, 2024.

Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, and Yonatan Belinkov. LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2410.02707. TODO verify author list.

Pablo Riera, Pablo Brusco, Cristina Kuo, Marcelo Sancinetti, and S. R. K. Branavan. Synchronization and turn-taking in full-duplex speech dialogue models. *arXiv preprint arXiv:2605.20356*, 2026.

Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2302.04761.

Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2207.07061.

Shannon Zejiang Shen, Hunter Lang, Bailin Wang, Yoon Kim, and David Sontag. Learning to decode collaboratively with multiple language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. arXiv:2403.03870; Co-LLM.

Tanya Stivers, N. J. Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann, Gertie Hoymann, Federico Rossano, Jan Peter de Ruiter, Kyung-Eun Yoon, and Stephen C. Levinson. Universals and cultural variation in turn-taking in conversation. *Proceedings of the National Academy of Sciences*, 106(26):10587–10592, 2009.

Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following LLaMA model. GitHub repository, 2023. zh split: shibing624/alpaca-zh.

Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong. Freshllms: Refreshing large language models with search engine augmentation. *arXiv preprint arXiv:2310.03214*, 2023.

444 Xiong Wang, Yangze Li, Chaoyou Fu, Lei Xie, Ke Li, Xing Sun, and Long Ma. Freeze-Omni:
445 A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv preprint*
446 *arXiv:2411.00774*, 2024a.

447 Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, et al. MMLU-Pro: A more robust and
448 challenging multi-task language understanding benchmark. In *Advances in Neural Information*
449 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2024b. arXiv:2406.01574.

450 Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese,
451 John Schulman, and William Fedus. Measuring short-form factuality in large language models.
452 *arXiv preprint arXiv:2411.04368*, 2024. SimpleQA.

453 Donghang Wu, Haoyang Zhang, Chen Chen, Tianyu Zhang, Fei Tian, Xuerui Yang, Gang Yu, Hexin
454 Liu, Nana Hou, Yuchen Hu, and Eng Siong Chng. Chronological thinking in full-duplex spoken
455 dialogue language models. In *Proceedings of the 27th Annual Meeting of the Special Interest*
456 *Group on Discourse and Dialogue (SIGDIAL)*, 2026. arXiv:2510.05150.

457 Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, et al. Qwen2.5-
458 Omni technical report. *arXiv preprint arXiv:2503.20215*, 2025a.

459 Jin Xu, Zhifang Guo, Hangrui Hu, Yunfei Chu, Xiong Wang, et al. Qwen3-Omni technical report.
460 *arXiv preprint arXiv:2509.17765*, 2025b.

461 Leyang Xue, Yao Fu, Luo Mai, and Mahesh K. Marina. HybridServe: Efficient serving of large AI
462 models with confidence-based cascade routing. *arXiv preprint arXiv:2505.12566*, 2025.

463 An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

464 Eric Zelikman, Georges Harik, Yijia Shao, Varuna Jayasiri, Nick Haber, and Noah D. Goodman.
465 Quiet-STaR: Language models can teach themselves to think before speaking. *arXiv preprint*
466 *arXiv:2403.09629*, 2024.

467 Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong,
468 and Jie Tang. GLM-4-Voice: Towards intelligent and human-like end-to-end spoken chatbot.
469 *arXiv preprint arXiv:2412.02612*, 2024.

470 Dengjia Zhang, Alexander Martin, William Jurayj, Kenton Murray, Benjamin Van Durme, and Reno
471 Kriz. Unified multimodal uncertain inference. *arXiv preprint arXiv:2604.08701*, 2026.

472 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang,
473 Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica.
474 Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information*
475 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2023. arXiv:2306.05685.

476 Wenhao Zheng, Yixiao Chen, Weitong Zhang, Souvik Kundu, Yun Li, Zhengzhong Liu, Eric P.
477 Xing, Hongyi Wang, and Huaxiu Yao. CITER: Collaborative inference for efficient large language
478 model decoding with token-level routing. In *Conference on Language Modeling (COLM)*, 2025.
479 arXiv:2502.01976.

480 A Model roster

Table 3: Model roster. “Pair” groups a fine-tune with its raw backbone control; the comparison statistic is always the within-pair difference.

model	type	pair	role
MiniCPM-o 4.5	omni + duplex FT	Qwen3-8B	system target
Qwen3-8B	raw	—	pair control
MiniCPM-o 2.6	duplex FT	Qwen2.5-7B	duplex replication
Qwen2.5-Omni-7B	streaming omni FT	Qwen2.5-7B	omni control
Qwen2.5-7B	raw	—	pair control
Qwen2-Audio-7B	audio FT	Qwen2-7B	audio-FT control
Freeze-Omni	frozen LLM + adapter	Qwen2-7B	adapter control
Qwen2-7B	raw	—	pair control
Mistral-7B-v0.3	raw	—	backbone diversity
GLM4-9B-chat	raw	—	backbone diversity

481 B Extended signal and readout analyses

482 **ROC curves.** Figure 3 shows the calibration-split ROC for the signals of Table 1.




Figure 3: ROC on the calibration split. $p(\text{True})$ -post leads but needs a full drafted answer; the probes and $p(\text{True})$ -pre are pre-answer.

483 **$p(\text{True})$ backbone dependence.** Across backbones, post-draft $p(\text{True})$ beats the trained final-
 484 layer probe on four of five raw/duplex models tested; the exception (GLM4-9B-chat, 0.685 vs. probe
 485 0.798) shows the pattern depends on the backbone’s self-evaluation quality.

486 **Matched pairs.** Table 4: with label coverage matched by subsampling the raw control to the
 487 duplex model’s per-pool fail rates, the raw backbone keeps LOPO math at 0.825/0.809 while duplex
 488 models sit at chance — label coverage is ruled out; the representation changed. Degradation is
 489 ordered raw > omni-streaming > duplex. Table 5 gives the by-depth summary behind Figure 2:
 490 the within-pair deficit at the best mid-layer is only -0.03 vs. -0.59 at the final layer; mean-pooled

491 probes survive to the final layer on duplex models (LOPO math ≈ 0.80 at L35 on o4.5, 0.674 at L27
 492 on o2.6 — 0.13–0.15 below each model’s mid-layer last-token peak, and uniformly weaker in-mix,
 493 –0.058 at L22); duplex damage is severe but local, omni-streaming damage mild but diffuse.

Table 4: LOPO transfer (AUC, held-out pool) of the final-layer last-token probe, Qwen2.5-7B family, label-coverage matched.

LOPO pool	Qwen2.5-7B (raw)	Qwen2.5-Omni	MiniCPM-o 2.6	[o4.5 vs Qwen3]
hard-math	0.809	0.745	0.538 $\approx$ chance	0.372 (inverts)
hard-knowledge	0.680	0.613	0.526 $\approx$ chance	0.61

Table 5: LOPO math transfer by depth (last-token probes).

model	best mid-layer	final layer	shape
Qwen3-8B (raw)	0.964 (L33/36)	0.958	plateau, no cliff
MiniCPM-o 4.5 (duplex)	0.931 (L22/36)	0.366	cliff, inverts
Qwen2.5-7B (raw)	0.893 (L21/28)	0.809	mild dip
Qwen2.5-Omni (streaming)	0.794 (L16/28)	0.746	diffuse depression
MiniCPM-o 2.6 (duplex)	0.822 (L21/28)	0.540	cliff

494 **Modality specificity of the cliff.** Under audio input the final layer does *not* invert (L35 audio
 495 LOPO math 0.936 vs. 0.366 text); a speak-mode template control on text input leaves the cliff
 496 unchanged (0.362) — the operative variable is the modality of the input context, not the output
 497 mode.

498 **Audio collapse of verbalized introspection.** Table 6. Three controlled arms pin the mechanism
 499 on MiniCPM-o 4.5: *not perception* (ASR audit WER 0.074; $p(\text{True})$ on the model’s own transcript
 500 snaps back to 0.074; the well-heard subset still collapses); *not a context prior* (irrelevant filler audio
 501 does not inflate p_{yes}); *binding* (duplicating the question as text alongside the audio fully restores
 502 introspection, 0.034). The verbal instance-check runs over text-token pathways that audio embed-
 503 dings do not feed, while a linear probe reads the same information off the same forward pass at
 504 0.93+. The failure implies its own cheap fix, confirmed twice: re-present the query as text before
 505 asking (“repeat-then-judge”).

Table 6: Pre-answer introspection under audio input. Trap p_{yes} = mean stated probability of answering correctly on a 100%-fail pool (lower is better).

model	FT type	trap p_{yes} : text $\rightarrow$ audio	audio $p(\text{True})$ -pre AUC
MiniCPM-o 4.5	duplex	0.055 $\rightarrow$ 0.556 (collapse)	0.786 (trap dead)
MiniCPM-o 2.6	duplex	0.196 $\rightarrow$ 0.345	0.491 $\approx$ chance
Qwen2.5-Omni	omni-streaming	0.279 $\rightarrow$ 0.213 (intact)	0.727
Freeze-Omni	frozen + adapter	0.165 $\rightarrow$ 0.112 (intact)	0.612 (capability caveat)

506 **Real-speech validation.** The matched-pair design rerun on 200 real human recordings
 507 (VoiceBench SD-QA, USA split) replicates all three audio-side findings: modality tax
 508 (0.400 $\rightarrow$ 0.450), audio overconfidence (paired p_{yes} shift +0.089; failure subset 0.415 $\rightarrow$ 0.581), and
 509 the layer structure (text-calibrated read on real-speech audio holds 0.76–0.80 at L22–L26, peak
 510 0.800 at L25 — transfer across both content and modality).

511 **Synthesis.** End-to-end multimodal training *builds* the shared mid-layer core; duplex-style training
 512 additionally *damages* the readouts; omni-streaming gets both right; an adapter alone gets neither
 513 (probe readout text-fragile, verbal readout audio-fragile; in both cases knowledge survives and a
 514 readout breaks, and in both cases the breakage is duplex-fine-tune-specific).

The deployed probe. The mechanistic probe is a single-position L22 read. The deployed probe adds two zero-latency pooled positions (an 8-token tail mean and the running user-audio mean), re-fit on a calibration mix widened to 2,310 public-benchmark queries disjoint from every evaluation benchmark; a pre-registered guard held (frozen internal test AUC 0.860→0.879) and feature selection never saw the externals. In the mid-layer band the in-mix tradeoff optimum is broad (L20–L30) and the mid-layer probe’s in-distribution area advantage is modest — its real advantage is LOPO robustness, which an in-mix test cannot show.

C Second duplex family: NemotronLabs-VoiceChat-11B

Table 7: Second duplex family (NemotronLabs-VoiceChat-11B): same frozen recipe, offline re-mix with measured GPT-5.5 outcomes, official judges. The English-only model has no zh pool, so avg spans four benchmarks and Δ is over its own local floor. Accuracies in %.

	TriviaQA	WebQ	Llama Q.	SD-QA	Reason. zh	avg	Δ
always-local	35.6	37.6	70.5	31.0	—	43.7	—
+ gate @15%	49.8	48.0	78.6	42.5	—	54.7	+11.0
+ gate @30%	61.5	58.8	83.8	54.5	—	64.6	+21.0
+ gate @50%	74.1	69.6	87.6	67.5	—	74.7	+31.0
always-expert (ceiling)	95.5	82.8	93.2	89.0	—	90.1	+46.4
probe AUC (frozen recipe)	0.775	0.771	0.720	0.754	—	0.755	—

The pre-registered prediction test: with the identical recipe (same audio, same judge, 600-query calibration), the layer sweep on this hybrid-Mamba duplex model (56 blocks: 27 Mamba-2 / 4 attention) peaks mid-band at L30–34 (eot-last AUC 0.714 vs. 0.693/0.682 at the endpoints); stacking the same three positions lifts OOF 0.714→0.761→0.790; the frozen-methodology probe transfers to the external pools at AUC 0.781/0.793/0.754/0.701 — the deployed-probe band reached with a quarter of the calibration data. Boundary notes: the backbone is much weaker at knowledge retrieval (triviaqa local 0.32 vs. MiniCPM 0.62, same judge) and English-only (the cross-lingual pool has no analog). The live-loop port and the trained-checkpoint ablation remain future work; the port is gated by the released inference path, not by the architecture. The only supported loading path (the NeMo Speech nemotron-labs-voicechat branch) runs the duplex loop *cacheless* — every 80 ms frame re-runs the full prefix, an $O(T^2)$ schedule that cannot hold the frame clock in real time — although the hybrid backbone’s Mamba-2 state is exactly the $O(1)$ per-frame recurrence a streaming port would need; wiring stateful inference into the released duplex frame protocol is the engineering gap. (Offline, the same property is a convenience: the final frame’s forward contains every position, so one hook captures the whole eot window.) Accuracy-side conclusions do not wait on the port: the re-mix arithmetic below tracks measured live arms at 0.011 mean deviation on MiniCPM, inside the ± 0.02 –0.03 replication floor.

Escalation re-mix on every transfer pool. Completing the grid of Table 2 on this family: an offline re-mix (the same top- r reconstruction arithmetic that, on MiniCPM, tracks the measured live arms at 0.011 mean deviation) sends the top- r of each pool by NVDA probe score to the measured GPT-5.5 outcome and keeps NVDA’s own answer for the rest, scored with each pool’s official protocol (the OAB judge for the three OpenAudioBench pools — NVDA local floors 0.352/0.376/0.705 — our judge for SD-QA, VoiceBench 1–5 for AlpacaEval; the expert answers the heard transcript and is judged directly, so the relay-back channel is excluded — including the measured relay outcomes where they exist shifts tier points by ≤ 0.032). **Selective escalation beats matched-rate random on all five pools** (permutation $p \leq 0.0003$; Table 8, Figure 4) — including Web Questions, SD-QA and AlpacaEval, the three pools that are flat or negative on MiniCPM (§7). Selectivity is nearly unchanged (AUC 0.72–0.78 under official labels); what changed is headroom — floors of 0.31–0.38 against MiniCPM’s 0.51–0.66 — and on AlpacaEval the failure species: the terse English-only model fails open-ended queries *hard* (the probe-selected half scores 3.01 vs. 4.08 kept), where

553 MiniCPM’s open-ended failures are soft quality gradations a wrongness detector cannot see. The
 554 negatives are properties of the pool $\times$ backbone pairing, not of the signal.

Table 8: The frozen-recipe NVDA probe re-mixed at fixed escalation rates, official judge per pool. AlpacaEval rows are VoiceBench 1–5 scores; others accuracy. rnd@50 = matched-rate random expectation; all five selective-vs-random gaps hold at $p \leq 0.0003$ (permutation).

pool	judge	n	0%	15%	30%	50%	100%	rnd@50
Speech TriviaQA	OAB	247	0.356	0.498	0.615	0.741	0.955	0.656
Speech Web Questions	OAB	250	0.376	0.480	0.588	0.696	0.828	0.602
Llama Questions	OAB	234	0.705	0.786	0.838	0.876	0.932	0.818
SD-QA	ours	200	0.310	0.425	0.545	0.675	0.890	0.600
AlpacaEval (1–5)	VB	199	3.55	3.83	4.13	4.46	4.92	4.23




Figure 4: Re-mix curves for the frozen-recipe NVDA probe: selective (top- r by probe score) vs. matched-rate random on the five transfer pools, official judge per pool. Every pool separates, including the three that are flat on MiniCPM (all five pools $p \leq 0.0003$, permutation).

555 D Query-surface router audit

556 A RouteLLM-style TF-IDF $\rightarrow$ LR router under identical labels reaches OOF AUC 0.669 (below the
 557 pool oracle’s 0.715); a 24-configuration sweep tops out at 0.689, so this is the query surface’s ceiling,
 558 not tuning. Under LOPO it collapses to 0.38–0.57 and scores the 100%-fail trap pool at 0.230 —
 559 it misses the pool a router most needs. The training receipt: at the argmax threshold its out-of-fold
 560 accuracy exactly matches the majority class; the internal probes under the same accounting clear
 561 majority by 9–29 points (Figure 7). Nor is it data-starved: the identical recipe on RouterBench [Hu
 562 et al., 2024] (36,497 prompts) reaches only in-domain AUC 0.710 with leave-one-benchmark-out
 563 at chance; RouteLLM’s released preference-trained routers score 0.523/0.533 zero-shot on our pool
 564 and rank the trap pool below ordinary hard pools; our router transfers to RouterBench below chance
 565 (0.440). On audio labels the router improves (test 0.814) but still trails the mid-layer probe (0.879)
 566 and cannot fire mid-stream. In distribution, the gate’s area over random (+0.054) is statistically
 567 on par with the pool oracle’s (+0.042; paired bootstrap delta CI $[-0.003, +0.027]$, $p=0.13$) — the
 568 gate’s case is transfer, not the in-distribution margin. $p(\text{True})$ tradeoff variants: Figure 8.




Figure 5: Second-family test: (a) the usable read again peaks mid-network; (b) frozen-recipe transfer lands in the deployed-probe AUC band.




Figure 6: Accuracy vs. escalation rate, frozen test split: gated escalation vs. random-escalation and pool-oracle (type-only) baselines between the small-only and big-only endpoints.

569 E Gate timing details

Table 9: Deployment policies, frozen test split ($n=240$); latencies measured.

policy	acc (%)	lat. mean (s)	P50	P95
fast-only	58.8	3.0	3.5	4.2
all-THINK (built-in)	63.7	22.4	17.2	60.1
gated-think @33%	61.3	12.1	3.6	47.6
gated-cloud @15%	68.8	5.3	3.6	10.1
gated-cloud @33%	78.7	6.5	3.6	20.8
gated-cloud @50%	85.8	7.0	3.6	24.4

570 **Measurement protocol.** Gate latency is wall-clock from prefill start to the L22 score being avail-
571 able (truncated forward + probe dot product), CUDA-synchronized, one H100, 50 queries per
572 modality, three warmup calls excluded, interleaved per query with the reference arms; P50/P95:
573 gate 20/25 ms text and 45/104 ms audio vs. full-prefill TTFT 36/47 and 68/144 ms. Network and
574 expert-RPC initiation are excluded: the cloud call fires asynchronously at the gate, and its user-
575 facing cost is reported as expert-content-audible latency in Appendix F. Absolute prefill wall-times
576 vary with call overhead across benches; the robust cross-bench statistic is the ratio — the L22 state
577 is ready at ~ 57 – 61% of prefill.




Figure 7: Identical training receipt for the external router and the two internal probes: the router never clears the majority-class baseline; both probes do, by 9–29 points.




Figure 8: Offline tradeoff for the $p(\text{True})$ signals (pre-answer and post-draft).

578 Per-layer truncated-forward cost vs. quality (Figure 9): the L22 hidden state is available at ~ 57 – 61%
 579 of prefill wall-time while L22 is simultaneously the in-mix quality peak; forking much earlier is both
 580 weaker and modality-specific, so the fork belongs at ~ 55 – 60% depth. If the cloud call launches at
 581 the gate, the expert result arrives before the local answer finishes for 40–58% of the whole mix but
 582 only ~ 20 – 31% of *escalated* queries (they skew to short local answers); the residual wait is P50
 583 2–3 s, P90 ~ 27 s (Figure 10) — the stall-phase budget the injection protocol must cover.




Figure 9: Fork-depth Pareto: per-layer wall-time vs. gate quality.




Figure 10: Cloud round-trip vs. local answer duration: probability the expert result arrives before the local answer ends.

584 F Live-loop details




Figure 11: Live dual-view tradeoff ($n=218$ speakable subset). Blue: honest heard accuracy of the live system (paired-bootstrap CIs). Green: the channel-controlled counterfactual (same sessions, escalated rows re-scored with gold-text expert answers) with its paired random reference (dashed). The view-to-view gap is the speech-channel price; the floor-to-floor gap to offline operation is the audio + streaming-answer tax.

585 **Full live tables.** Table 10 (both scoring views, speakable subset and full pool) and Table 11 (measured response latency). The conservative tier *shortens* the P95 tail: its few escalations displace
 586 exactly the longest local decodes. Figure 12 joins accuracy and latency; Figure 13 replays three real
 587 sessions from recorded timers.
 588

589 **Time to first audible audio.** The sweep above is text-out; a TTS-on re-run of the balanced arm
 590 (same 240 queries, talker head + vocoder on, stall pre-synthesized once in the talker’s own voice)
 591 measures speech onset directly. The gate decision is unchanged — the TTS token enters the context
 592 only after the end-of-turn read (read p50 21 ms / p99 32 ms), and the same 84/240 queries fire. From
 593 gate decision to first audible audio: **local answers p50 0.552 s / p95 0.599 / p99 0.644; escalated**
 594 **turns p50 0.022 s / p99 0.029** — the canned stall (3.2–4.1 s of speech) plays the moment the gate
 595 fires, so the escalated path reaches first audio 25× faster. Expert *content* becomes audible at p50
 596 6.4 s / p95 26.0 / p99 34.0 (cache-corrected expert round-trip + first relay PCM, p50 0.66 s after the

597 answer returns); per-turn dead air between stall end and expert content is p50 2.5 s / p95 22.4 / p99
598 30.7, with 25/84 escalated turns fully covered — the stall covers the head of that wait, and Table 11’s
599 completion profile bounds the tail.

Table 10: Live streaming sweep, paired bootstrap 95% CIs ($B=10^4$). Top: speakable subset ($n=218$); bottom: full pool ($n=240$). Gold-inject = escalated rows re-scored with gold-text expert answers; channel cost = paired difference. * CI excludes 0; † CI lower bound at zero.

tier	esc	heard-acc	Δ vs. floor	gold-inject	channel cost
never	0%	0.422 [0.36,0.49]	—	0.422	—
conservative	14%	0.528 [0.46,0.59]	+0.106*	0.550	+0.023†
balanced	31%	0.592 [0.53,0.66]	+0.170*	0.661	+0.069*
aggressive	57%	0.633 [0.57,0.70]	+0.211*	0.761	+0.128*
always (synth)	100%	—	—	0.922 [0.89,0.95]	—
never (full)	0%	0.383 [0.32,0.45]	—	0.383	—
conservative	16%	0.483 [0.42,0.55]	+0.100*	0.525	+0.042*
balanced	35%	0.554 [0.49,0.62]	+0.171*	0.654	+0.100*
aggressive	61%	0.596 [0.53,0.66]	+0.212*	0.771	+0.175*
always (synth)	100%	—	—	0.917 [0.88,0.95]	—

Table 11: Measured response latency (s), $n=240$ per arm (latency is channel-independent). P99 on $n=240$ is indicative only.

tier	mean	P50	P95	P99	Δ mean	Δ P50	Δ P95	Δ P99
never	4.6	2.0	15.8	17.8	—	—	—	—
conservative	4.6	2.6	12.8	23.0	+0.0	+0.6	−3.0	+5.2
balanced	5.8	3.5	15.9	32.7	+1.2	+1.5	+0.1	+14.9
aggressive	6.6	4.4	18.5	33.0	+2.0	+2.4	+2.7	+15.2




Figure 12: Accuracy vs. per-arm P50 latency (speakable subset), both scoring views, with the simulated-P50 random reference.

600 **Speculative prefetch, measured and dropped.** Framed as speculative escalation (draft–verify at
601 sentence granularity), its acceptance rate — does an expert answer to the partial transcript answer
602 the full question — is 0.07/0.19/0.51 at 25/50/75% of the audio; at the mid-stream gate’s typical fire
603 point 6–8 of 10 calls would be wasted. The trap pool’s floor (0.00/0.10/0.27) shares the back-loaded-
604 core property that breaks mid-stream probe reads.

Three live sessions, timestamped (balanced arm; every event time is a recorded live-loop timer)

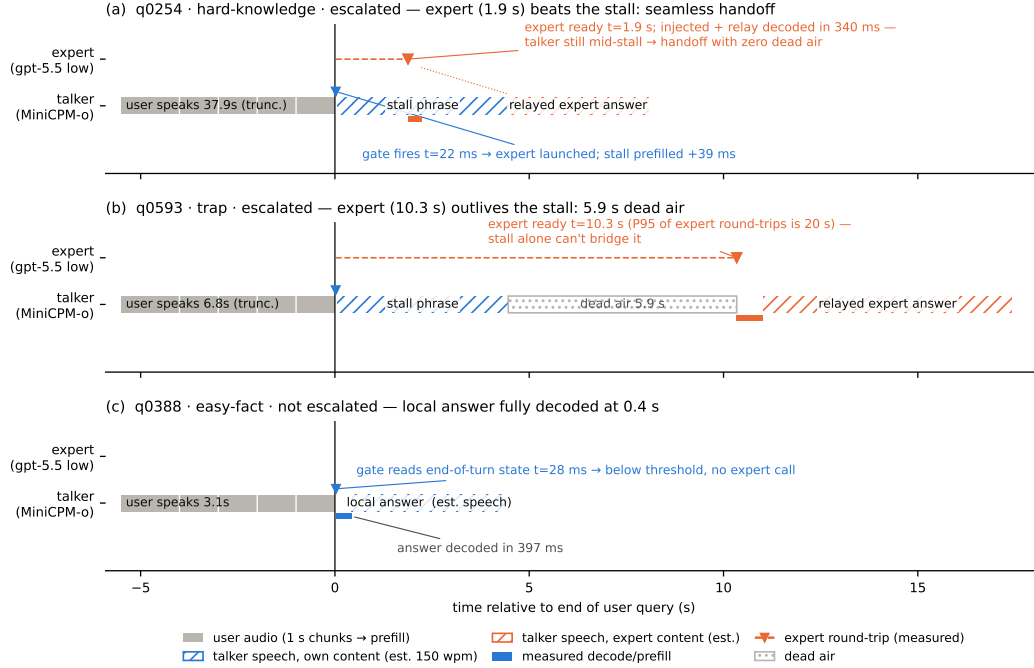


Figure 13: Three real sessions from recorded timers: expert beats the stall phrase (seamless); expert outlives it (dead air, the P99 signature); gate stays closed (local answer fully decoded 0.4 s after end of turn). Speech occupancy is estimated at 150 wpm (hatched); the measured arms are text-out, so no vocoder onset is plotted.

605 **Conflicting-injection controls (172 sessions).** Fabricated conflicting expert answers by within-
606 pool derangement, four framings: worst-case comply is 0.00 neutral / 0.53 deployed authority / 0.79
607 strong-override, while *seeding* words into the talker’s mouth backfires (comply 0.25, lip-service 0.62
608 — forced speech is treated as revisable). Math is silently re-derived rather than relayed or disputed.
609 Model-wrong × wrong-injection swallow rate is 0.64: expert quality is the accuracy ceiling of the
610 cascade.

611 **Uplink diagnostics.** On the escalated population the expert answering the talker’s own transcript
612 scores 0.585 vs. 0.871 on gold text. Prompt-level rescues on the self-transcript are no-ops (the
613 damaging errors are internally-coherent entity substitutions); replacing the ears helps: open-weights
614 Whisper recovers little (+4 pp) but a stronger cloud ASR on the same audio recovers 38% of the gap
615 (0.585→0.694, McNemar $p=0.007$) — diagnostic only (its critical-path latency and external-pool
616 survival are unmeasured). A same-model gold-text control with an audio-capable expert shows the
617 audio channel itself is lossless for speakable content: the leak is the talker’s transcription, not the
618 synthesis or the relay (relay drops: 5 cases of 108; corrupted transcripts dominate the escalated
619 heard-failures, first-generation sweep).

620 **Boundary query classes.** *False premises:* the model fails substantially (0.63 text / 0.47 audio)
621 and the expert would help (0.80), but the gate is blind (fire@balanced 0.18 vs. trap’s 0.90; no score
622 separation) — answering along a false premise does not feel like inability. *Real-time queries* (“what
623 is NVDA trading at”): the gate is not blind, the class was simply absent from every static training
624 family; extending calibration with FreshQA [Vu et al., 2023] (fast-changing labeled escalate a priori,
625 never-changing as in-family controls) raises held-out fast-changing fire rates to 0.80/0.93/1.0 across
626 tiers with frozen-test AUC unchanged (0.877 vs. 0.879). One calibration subtlety: tier budgets must

627 remain quantiles of the deployment-mix scores, or the new capability is spent on its own threshold
628 shift.

629 **Demonstration system.** The pipeline runs as an interactive web demo on one H100: live micro-
630 phone audio, end-of-turn read (~ 54 ms) against the balanced threshold, escalated turns stall, consult
631 the expert, and *speak* the relayed answer through the model’s native TTS head (probe context and
632 thresholds untouched). Two demo-only extensions: a web-search tool for the expert (real-time
633 queries then return live values), and backchannel-aware barge-in — speech over the talker ducks
634 playback, a fast ASR pass returns the floor on backchannels (“ok”, “mm-hm”; an English+Chinese
635 lexicon) or commits the interrupt, aborting generation at chunk granularity and seeding the next turn.
636 All measured arms keep the tool-free expert and turn-based loop; the demo will be made available
637 upon publication.

638 G Duplex-validation sweep: the gate under the talker

639 The headless live sweep (§6) never engages the talker head; this sweep replays its full 240-query
640 pool through the deployed voice loop above to test whether the frozen probe read survives the spo-
641 ken regime. Two arms, one session per query, deployed v4 gate with scores logged but escalation
642 disabled (the arm isolates gate validity, not routing outcomes). **Clean:** the query streams as mic-
643 shaped PCM (pool TTS audio plus steady noise at RMS 0.008), the energy VAD ends the turn, the
644 talker answers aloud, and the session is cut after the first audible chunk. **Overlap:** an easy warmup
645 question (probe score < 0.25 , locally answered, > 5 s spoken answer) puts the talker mid-answer,
646 and the target query is then spoken over it; sustained overlapping speech commits a barge-in whose
647 audio seeds the next turn. The barge-in path engaged on 237/239 pairs.

648 Against the frozen local-failure labels of the same pool, the identical frozen probe reads: headless
649 AUC 0.866 [0.817, 0.909], clean 0.871 [0.824, 0.914], overlap 0.856 [0.803, 0.901]; barged-only
650 subset ($n=237$) 0.854. Paired bootstrap deltas vs. headless are $+0.005$ [$-0.025, +0.035$] (clean) and
651 -0.009 [$-0.036, +0.016$] (overlap) — both inside the ± 0.02 – 0.03 live replication floor. Per-query
652 scores correlate across regimes at $r=0.90$ – 0.92 , and the balanced threshold makes the same fire/hold
653 decision as the headless read on 0.887 (clean) and 0.912 (overlap) of queries. The end-of-turn read
654 costs 24–26 ms at p50 ($p95 \leq 30$ ms) in both arms, with the first audible answer chunk arriving ≥ 0.5 s
655 after the gate decision ($p50 \approx 0.58$ s), so the decision precedes any audible commitment. Scope: this
656 loop keeps per-turn session control — the mic stays open while the talker speaks and interruptions
657 seed the next read, but incoming audio is not prefilled *during* generation. The concurrent regime is
658 measured separately below.

659 **Concurrent prefill: interleaved listen/speak in one KV stream.** A third arm drives the stream-
660 ing API in time division: while the talker is mid-answer, each 1 s chunk of the target query is
661 prefilled into the *same* session between generation chunks, so the end-of-turn read fires while gener-
662 ation is still active (238/240 queries, full interleave coverage). The carrier answer is teacher-forced:
663 under sampling the duplex model reliably yields the turn — it emits EOS within ~ 2 s of user au-
664 dio appearing in its context, the trained turn-taking behavior surfacing under interleave. This is
665 the one shift on the transfer ladder that breaks the frozen read: scores drift up (mean $0.48 \rightarrow 0.71$;
666 the frozen balanced threshold’s fire rate jumps $37 \rightarrow 75\%$), AUC drops $0.869 \rightarrow 0.758$ [0.693, 0.816]
667 (paired Δ AUC -0.109 [$-0.158, -0.064$]), and label-free requantiling recovers decision agreement
668 only to 75%. The information, however, survives in the state: refitting the *same* 12288-d linear
669 read on 360 calibration queries collected in-regime reaches **0.818** [0.761, 0.871] on the held-out 240
670 (calibration OOF 0.775) — more than half the gap recovered at a quarter of the deployed probe’s
671 calibration scale. The mid-layer competence direction persists under concurrent listen/speak; unlike
672 every earlier shift (modality, external pools, talker-on), calibration must follow the regime. Run-to-
673 run replication of the concurrent read correlates at 0.895. Caveat: the interleave goes through the
674 public streaming API — turn-role headers enter mid-generation and the carrier text is forced — so
675 this measures the concurrent *state*, not the official duplex serving format.

676 **A serendipitous recovery behavior.** An accidental run surfaced one bonus behavior worth report-
677 ing: when a long question contains a >1.25 s internal pause, the endpointer fires early and the talker
678 starts answering the fragment — but the question’s own continuation then triggers barge-in, and
679 the remainder becomes the next turn, where the gate simply re-scores it. Early end-pointing thus
680 degrades into an extra turn rather than a lost query — the honest duplex-ish recovery available to a
681 turn-based loop.

682 H Fixed-threshold transfer and the expert-benefit oracle

683 Both analyses below run in the *gold-inject* view (escalated rows take the expert’s answer to the gold
684 text). The deployed heard view cannot evaluate a freely moving threshold: the most permissive arm
685 escalates 50%, so the other half of each pool has no expert outcome on record, whereas the ceiling
686 runs cover every id. Labels follow each pool’s own judge (Table 2); the two judges disagree by up to
687 13 points on Web Questions, so they are never mixed within a row. AUC columns are recomputed
688 from the deployed sweep’s stored end-of-turn scores; they differ from Table 2’s AUC column —
689 captured in a separate scoring pass — by up to 0.03 externally and 0.04 on the internal pool (0.840
690 here vs. 0.879 there): the passes render and prefill the audio independently, and the probe score
691 carries render-to-render jitter. AlpacaEval is excluded — it has no binary correctness label on either
692 side, only a 1–5 score whose expert ceiling is degenerate (195/199 at 5).

693 **One absolute threshold, six pools.** The deployed gate adapts per pool, but only through label-free
694 quantiles of its own score distribution. To separate what the probe’s *ranking* carries from what its
695 *scale* carries, we freeze the single absolute threshold the internal system actually deployed ($\tau=0.628$,
696 the balanced tier) and apply that number verbatim everywhere. **Scale does not transfer:** the same
697 τ realizes escalation rates from 2.0% (Reasoning QA) to 48.0% (SD-QA), a 24-fold spread, so a
698 fixed cut cannot hold an escalation budget across pools. **Ranking does:** at whatever rate it fires,
699 fixed- τ escalation beats a random reference at the same realized rate on four of five external pools,
700 paired-bootstrap interval excluding zero (Table 12). The exception is Reasoning QA, where τ fires
701 on 2% of queries and there is almost nothing to measure. This is the concrete case for the label-free
702 quantile step: it buys budget control, not discrimination.

703 **Does the gate find failures the expert can fix?** The probe is fit and read against *local failure*,
704 but what a routing system needs is *expert benefit*: $y=1$ when the local model is wrong *and* the
705 expert is right. Re-scoring the same frozen scores against this harder label costs little on the retrieval
706 pools (Speech TriviaQA 0.796→0.782, Llama Questions 0.830→0.813) and more where the expert
707 is itself weak (Web Questions 0.773→0.689 against a 0.864 expert ceiling; internal 0.840→0.732).
708 The gate thus remains a usable benefit predictor out of distribution (0.63–0.81) despite never having
709 seen an expert outcome, and the drop is largest exactly where the paper already reports a flat live
710 curve. On Web Questions the mechanism is direct: 75.0% of local failures are expert-fixable pool-
711 wide, but only 65.5% of the failures inside the gate’s top-scoring 15% are, and that share climbs
712 monotonically as the cut widens (0.655/0.685/0.711 at 15/30/50%) — the gate’s most confident
713 failure calls are the ones the expert also misses. Table 13 places the deployed gate against the
714 matched-budget oracle that escalates true benefit cases first: the gate clears random at every pool
715 and budget, capturing a fifth to three fifths of the random-to-oracle span, and the residual gap is the
716 headroom a decode-time or expert-aware second stage would have to claim.

717 I Transfer latency shapes

718 **Matched-rate precision.** Precision at a fixed escalation rate is capped at base-fail/rate: on Llama
719 Questions (base fail 0.16) the 50% cap is 0.32 and the deployed probe reaches 0.304 — 95% of the
720 cap (recall 0.93). Quoting raw precision would misread the strongest pool as the weakest.

721 **The median left-fold.** On easy pools per-arm P50 is non-monotonic (Llama Questions
722 1.52→1.17→1.60/1.42 s): a verified median-of-a-mixture mechanism, not noise. The probe pref-

Table 12: One absolute threshold ($\tau=0.628$, frozen on the internal pool) applied verbatim to every pool; gold-inject view. Δ is fixed- τ accuracy minus a random reference at the *same realized rate* (paired bootstrap over ids, $B=10^4$, seed 42). AUC columns re-score the same frozen scores against local failure and against expert benefit ($y = \text{local wrong} \wedge \text{expert right}$, base rate in the last column). Rates and accuracies in %; AUC in decimals.

pool	n	rate	local	acc	rand	Δ 95% CI	AUC _{fail}	AUC _{benefit}	base
Speech TriviaQA	250	28.4	66.4	83.2	75.0	[+4.4, +12.1]	0.796	0.782	30.8
Speech Web Questions	250	36.0	56.8	72.0	67.5	[+0.8, +8.4]	0.773	0.689	32.4
Llama Questions	250	8.0	83.6	86.8	84.3	[+0.6, +4.6]	0.830	0.813	11.2
SD-QA	200	48.0	51.5	80.5	71.4	[+4.5, +13.6]	0.800	0.765	43.5
Reasoning QA (zh)	202	2.0	58.9	60.9	59.5	[-0.1, +3.5]	0.679	0.634	30.2
internal frozen test	240	35.0	38.3	65.4	57.0	[+3.8, +13.1]	0.840	0.732	53.3

Table 13: Oracle frontier at matched escalation budgets (gold-inject view): random reference / deployed quantile gate / benefit oracle, which escalates the true $y=1$ cases first. The gate clears random everywhere; the oracle column is the remaining headroom. Accuracies in %.

pool	15% budget			30% budget			50% budget		
	rand	gate	oracle	rand	gate	oracle	rand	gate	oracle
Speech TriviaQA	71.0	76.0	81.6	75.5	83.6	96.4	81.6	91.2	97.2
Speech Web Questions	61.2	64.0	72.0	65.7	70.0	86.8	71.6	76.4	89.2
Llama Questions	85.0	88.8	94.8	86.4	91.2	94.8	88.2	93.2	94.8
SD-QA	57.7	63.5	66.5	63.9	74.5	81.5	72.2	82.0	95.0
Reasoning QA (zh)	63.1	67.8	73.8	67.4	72.8	89.1	73.0	76.7	89.1
internal frozen test	46.3	49.2	53.3	54.3	60.8	68.3	65.0	74.2	88.3

723 entially escalates the queries whose *local* decode is longest — on retrieval pools wrong answers
724 hedge at length (decode time and answer length $r=0.97$), so a wrongness-selecting probe strips the
725 slow tail as a side effect (escalated-at-15% ids: never-arm local P50 1.711 s vs. 1.198 s kept; local
726 accuracy 0.32 vs. 0.73). Controls: the probe does not key on length ($\text{corr}(\text{score}, \text{length}) = +0.05$;
727 ≈ 0 conditioned on wrongness), and on SD-QA — where wrong answers are short — the same probe
728 selects wrongness identically and produces no fold. The mean is monotonic (1.65→2.06 s); kept-
729 local P50 falls 1.52→0.43 s across tiers (Figure 14). On Reasoning QA the shape inverts usefully:
730 the local chain-of-thought is *spoken*, so every reasoning token enters the response clock, while the
731 expert’s chain is hidden behind a flat ~ 3.7 s round-trip — escalation truncates the latency tail (P90
732 13.25→10.94 s) rather than fattening it.




Figure 14: Mixture decomposition of Llama Questions latency per arm: kept-local P50 falls 1.52→0.43 s as the probe strips the slowest local decodes while escalated rows pay a flat ~ 3 s expert round-trip; the arm median (dashed) is their mixture.

AlpacaEval scale caveats. The official 4.8 is an offline chat-mode number — our chat-mode control reproduces it at 4.86 — and the chat-vs-live gap (4.86 vs. 3.99) is the duplex loop’s spoken-reply style plus a token cap, so only live arms are comparable to each other. Selected queries score 3.90 on the never arm vs. 3.98 for skipped: no discrimination; escalated rows gain (3.90→4.71) because the expert writes better long-form answers — random picks would harvest the same gain.

J Full-pool and per-family figures

Figures 15 (dual-view) on the full frozen pool, and the complete external-family figure set (Appendix K); the noise audit behind the ± 0.02 –0.03 floor is Figure 16.




Figure 15: Dual-view live tradeoff, full pool ($n=240$); companion to Figure 11.




Figure 16: Replication-noise audit: same query, same audio, re-run across arms. Kept-local verdicts flip 2.3–18.8% by pool; repeated escalation 0.7–10.6%; implied paired SE 0.009–0.028.

K External-benchmark figure set




Figure 17: Speech TriviaQA, Speech Web Questions, SD-QA: dual-view (left) and accuracy–latency (right).




Figure 18: Llama Questions, Reasoning QA (zh), VoiceBench AlpacaEval: dual-view (left) and accuracy-latency (right).

742 L Judge and self-evaluation prompts

743 All adequacy labels use the following fixed system prompt (structured outputs with a strict {reason,
744 adequate} schema; the judge never sees which system produced the answer):

745 You are a strict, fair grader of a small language model’s answers.
746 You are given a QUESTION, an optional REFERENCE answer, and a
747 CANDIDATE answer produced by a small model. Decide whether the
748 CANDIDATE is ADEQUATE, meaning ALL of: (1) factually correct;
749 (2) any reasoning shown is valid; (3) it actually and completely
750 answers the question that was asked. For multiple-choice questions
751 the candidate must land on the correct option (by letter or by
752 unambiguously stating the option’s content). A partially correct,
753 hedged, refused, off-topic, or truncated answer is NOT adequate. If
754 a REFERENCE is provided, grade against it. Be strict but do not
755 penalize harmless differences in format or extra correct detail.

756 The verbalized self-evaluation baselines use one fixed phrasing each; $P(\text{Yes})$ is read from the first
757 generated token’s distribution:

758 [pre] I am going to show you a question. Do NOT answer it. Judge
759 honestly whether you yourself would answer it correctly. Question:
760 {q} Would you answer this question correctly? Reply with exactly one
761 word: Yes or No.
762 [post] Here is a question and a proposed answer. Question: {q}
763 Proposed answer: {a} Is the proposed answer correct? Reply with
764 exactly one word: Yes or No.

765 M Reproducibility

766 All sampling and splits use a fixed seed (42); the calibration/test split is frozen once and the test split
767 is touched only by the evaluation scripts. Environment pins: torch 2.8.0; transformers 4.51.0 for
768 MiniCPM-o 4.5, with per-model pins for comparison backbones; all experiments run on single-H100
769 instances. Probe artifacts (weights, feature configuration, per-pool label-free thresholds), the judge
770 prompt above, all trace files, and the per-experiment execution log — including per-run GPU and
771 API spend, on the order of \$500 total — ship in the project repository, released upon publication.

TODO (process only — delete before submission)

- [TODO: P1 — DONE 2026-08-25 (external review applied): main text now ends within 8 pages (refs start mid-page 8). Restructure: 3 contributions + hero figure; merged signal section; live 6pp → 1.2pp; router audit 3 lines; title dropped “Zero-Training”; causal claims softened (“consistent with”); no version numbers / artifact filenames in main text; everything cut lives in the appendix (A–L) or git history]
- [TODO: P2 — peer + meeting feedback (GPT reviewer pass 2026-08-26, borderline 4/7). (a) and (b) DONE 2026-08-26; (c) dropped as not computable; (d) unchanged. (a) DONE — three claims narrowed: “routing beats scaling the speech backbone” → “selective cloud escalation outperforms larger standalone backbones”; “standard readouts fail” → “late-layer readouts turn brittle after duplex fine-tuning” (abstract, intro contribution 1, hero + teaser panel-1 titles); and the EOT loop is no longer equated with native full-duplex — intro now carries an explicit scope sentence (“*duplex* names the training regime, not the interaction protocol”), live.tex states the turn-based protocol up front and is retitled “Live End-of-Turn Escalation”, and “duplex session/system” → “spoken session”/“live system” throughout. (b) DONE — new appendix section (app:fixedthr) + two tables: one absolute threshold ($\tau=.628$, frozen internally) applied verbatim to all pools fires at 2–48% yet still beats matched-rate random on 4/5 (ranking transfers, scale does not); expert-benefit oracle (y = local-wrong AND expert-right) with AUC re-scored per pool and a gate/random/oracle frontier at 15/30/50% budgets. Script: scripts/09_fixed_threshold_oracle.py → figures/fixed_threshold_oracle.json. Gold-inject view (heard view only observes the ~50% of ids some arm escalated); labels follow each pool’s own judge (oab_ok vs adequate — they differ by 13pp on Web Questions, never mix); AlpacaEval excluded, it has no binary label on either side. (c) DONE 2026-08-26 (was DROPPED, then re-run with TTS on): bench_live gained a tts flag (init_tts + Token2wav, canned stall, gen_speak port of demo_app); balanced arm re-run on the frozen 240 (~50 min H100, expert cache hits). EOT→first-audible: local p50 .552s/p99 .644s, escalated (canned stall onset) p50 .022s/p99 .029s, expert-content p50 6.4s/p95 26.0/p99 34.0. Same 84/240 fire as text-out; probe read unperturbed (p50 21ms). Landed: live.tex sentence + app:livedetail paragraph. Traces data/frozen_v3tts_traces.jsonl.balanced.shard0-3; analysis scripts/10_tts_latency.py. Fixed en route: fig:timeline-live panel (c) was captioned “first token 0.4s” but the plotted value is answer_ms = *full* answer decode; caption + figure title corrected and the figure regenerated. (d) DONE 2026-08-27 (was SKIP). 8as closed the last tier: concurrent prefill — target audio interleaved into the SAME KV stream between generation chunks, EOT read fires while generation is active (238/240). Frozen v3 probe degrades AUC 0.869→0.758 (paired Δ −0.109 [−0.158, −0.064]) — the ONE shift on the ladder that breaks frozen transfer; label-free requantiling recovers decisions only to 75%. In-regime refit of the same 12288-d read on 360 concurrent calib rows recovers 0.818 [0.761, 0.871] at quarter-scale data ⇒ signal intact, calibration must follow the regime. Landed: app:duplexval third arm, Limitations (vi), intro clause. Harness modal_duplex_concurrent.py; scripts 11/12; RESULTS 8as. Page rebalance done same pass: Table 2 policy table, dualview figure, second-family block and the serendipity paragraph moved to the appendix; main text ends p8, refs p9. Scale unified (accuracies/rates in %, AUC in decimals). Still open: official duplex serving format, full-scale (2310) in-regime recalibration, human continuity study, second expert. (d-earlier) 8aq duplex-validation sweep ran — the frozen-pool 240 replayed through the DEPLOYED voice loop (talker on, spoken answers) in clean + overlap/barge-in arms via _ws_duplex.py against the live demo; frozen v4 probe, no refit: AUC .871/.856 vs .866 headless (paired deltas inside the replication floor), decision agreement .89/.91, barged 237/239. Landed: live.tex closing paragraph, app:duplexval, intro scope sentence updated, Limitations (vi) narrowed to “concurrent-with-generation read remains open”. Data: gate-data volume

duplex_sweep/{clean,overlap}.jsonl; analysis_duplex_analyze.py. This also partially recovers (c): the voice loop logs EOT→first-audible per turn (first_audio_client_ms, p50 ≈.58s, both arms) — usable if a reviewer asks, though only under probe-off replay, not the escalating sweep. Still skipped: human continuity study, second expert, model-level concurrent listen/speak read. Frame the paper as what it is: pre-response competence probing in duplex-trained models + turn-based live escalation, now with a deployed-loop validity check.]

- [TODO: P2.2 — meeting feedback round 2 (2026-08-26), four items, all landed same day: (1) ugly zigzag figures: fork_pareto redrawn as depth-colored scatter + Pareto frontier (figures/fork_pareto.py, new); sllama_latency_decomp cut to the clean single decomposition panel (median-vs-mean fold fact stays in prose). (2) main table: tab:transfer upgraded Kitty-style — live tiers + ceiling + AUC + official PLUS gold-inject@30% rand/gate/oracle block (numbers from tab:oracle); judge folded into row names. Views kept separate and labeled — do NOT let anyone compare live heard vs gold-inject across blocks. (3) special-token routing discussed in related work: pause/thought tokens (goyal2024pause, zelikman2024quietstar), fused think modes (yang2025qwen3, guo2025deepseekr1), MiniCPM-o’s own think gate — they decide more compute not which model, need talker fine-tuning, policy baked per checkpoint; ours = frozen + pluggable + beats the built-in gate at matched budget (tab:policy). (4) simple-but-important framing: intro minimalism sentence (“competing signals read the wrong place / arrive too late / must be trained in”). Page budget paid by: fair_dualview .56→.45 linewidth, caption + related + gap compressions.]
- [TODO: P2.3 — storytelling pass (advisor feedback 2026-08-26, selectively adopted): narrative arc sharpened to “looks solved → the obvious readout is exactly what duplex training corrupts → signal sits one band earlier → read it before speech commits”. Abstract rewritten to that arc; intro adds the question line (“not whether one can build a router — whether the model knows it will fail early enough to do anything about it”) and the find-it/read-it-in-time/show-it-matters framing; “before the first token” upgraded to “half a second before the first audible word” (grounded in the TTS-on p50 .552s + gate ~30ms); conclusion closer = “the signal was never gone — only misread”. NOT adopted: restructuring the three contributions, cutting appendix richness. Space paid by trimming the extended-analyses name-drop line.]
- [TODO: P0-R — reviewer P0 pass (2026-08-27), five questions; answers exist in-repo, story needs surfacing: (1) why route pre-first-token: NOT a safety necessity (first token → first audible ≈0.5s exists) — it’s that waiting buys no AUC (measured: entropy@4 0.696, decode-hidden mean 0.776 vs prefill read 0.828/0.887) and every ms of delay is dead air on the escalated path (expert RTT = 85–88% of escalated wall time). Edit: one sentence in gate.tex reframing pre-token as free-and-optimal, citing the measured decode-time rows. (2) final-layer mean-pool WAS run: LOPO math o4.5 ≈0.80 (L35) / o2.6 0.674 vs mid last-token 0.931/0.822; in-dist uniformly worse (−0.058 at L22). Edit: replace signal.tex “are fine” with the numbers; claim = “only uniformly robust readout”, not “mid-layer necessary”. Watch counter-q: audio-native final layer is intact (0.936) — answer = text side still inverts and text→audio calibration only transfers mid-band. (3) label = local-wrong by design (expert-agnostic / pluggable); measured benefit-AUC cost already in app:fixedthr (0.63–0.81; internal 0.840→0.732; WebQ mechanism paragraph). Edit: promote 2 sentences to main text. Optional if labels exist: benefit-trained refit on stored EOT hiddens (CPU, minutes). (4) 20/45 ms = wall-clock prefill-start→L22 score, truncated forward + probe dot, CUDA-synced, H100, 50q, warmup excluded, P50 (P95 25/104); TTFT same bench. Excludes network/RPC — async path, reported separately. Edit: measurement-protocol sentence in app:gatedetail + define in gate.tex. (5) user-facing expert latency IS measured: stall onset p50 22 ms, stall 3.2–4.1 s, expert content audible p50 6.4 s/p95 26/p99 34, dead air after stall p50 ≈2.3–3.2 s, tail P90 ≈27 s. Edit: put content-audible numbers next to

878 the 22 ms claim in live.tex main text so $25\times$ can't be read as answer latency; optional:
879 exact dead-air p50/95/99 from frozen_v3tts traces (scripts/10_tts_latency.py extension,
880 CPU). STATUS 8-27: all five LANDED (gate.tex $\times 2$, signal.tex, live.tex, appendix $\times 3$);
881 dead air computed from the tts traces — p50 2.5 s / p95 22.4 / p99 30.7, 25/84 turns
882 fully covered (scripts/10_tts_latency.py, dead-air block added). Benefit-trained refit
883 NOT run (only re-scoring exists). $\Rightarrow$ WARNING page budget: same-engine (tectonic)
884 check puts References at p9@78% BEFORE this pass and p10@6% after — main
885 text already exceeded the 8-page target by ~ 0.8 pp before these edits (user build 8-27
886 01:52 also shows refs on p9). Needs a ~ 1 -page trim decision before upload.]

- 887 • [TODO: P3 — submission day (8-28): delete this section; final anonymization grep
888 (names, rhe9527, modal URLs, gmail); decide public-repo $\rightarrow$ private; OpenReview
889 form + upload (deadline 8-29 AoE, moderation hard limit Fri 8-28)]