
When Does a Small Model Know to Hand Off? Reading Competence from Full-Duplex Speech Models for Real-Time Escalation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Full-duplex speech models answer in real time with a small on-device backbone,
2 yet many queries exceed its competence. Routing to a large cloud model sounds
3 solved — until one asks *when* the decision must be made: before the first output
4 token, or the wrong answer is already on the air. We ask whether the model’s
5 own hidden state carries that warning early enough, and find that the obvious
6 readout is exactly what duplex fine-tuning corrupts: final-layer probes collapse
7 under distribution shift on text input (leave-one-pool-out math AUC 0.93→0.37),
8 consistent with late-layer specialization for duplex turn control, while verbalized
9 self-assessment collapses on audio input. The signal is not gone — it sits one layer
10 band earlier. Mid-network, the same linear probe transfers at 0.93, reads speech
11 input from text-only calibration at 0.86, and an end-of-turn read in the streaming
12 loop reaches AUC 0.887 in ~30 ms — half a second before the first audible word.
13 Run live end-to-end in a turn-based loop, the gated system climbs from 42.2%
14 to 63.3% heard accuracy at 57% escalation, the remaining gap dominated by the
15 transcription uplink rather than the gate; frozen, the same gate improves five of
16 six public speech benchmarks, lifting Speech TriviaQA to 86.0% live against the
17 model’s official offline 75.5%. Reading one band earlier, before speech commits,
18 is enough to escalate live.

19 1 Introduction

20 A full-duplex speech model listens and speaks in the same forward pass, under a conversational
21 deadline of a few hundred milliseconds [Stivers et al., 2009]. That deadline forces a small on-device
22 backbone — and a small backbone routinely faces queries beyond its competence: long-tail facts,
23 expert knowledge, multi-step reasoning. The natural architecture is selective escalation to a large
24 cloud model, but the duplex setting imposes two constraints that existing routing work does not
25 meet: the decision must come *before the first output token* (once the model starts speaking, a wrong
26 answer is already on the air), and the session — audio context, voice, cache — must stay with the
27 local model, which will speak the expert’s answer itself. The question is not whether one can build
28 a router; it is whether the small model knows it will fail *early enough to do anything about it*.

29 To answer it, we study MiniCPM-o 4.5 [Cui et al., 2026], a 9B omni/duplex fine-tune of Qwen3-8B,
30 against matched-pair controls spanning raw backbones, streaming-omni, audio-only, and frozen-
31 adapter fine-tunes, and evaluate every standard zero-shot competence readout — decode-time en-
32 trophy, verbalized self-assessment ($p(\text{True})$), and linear probes on hidden states — with judged fail-




Figure 1: **(1)** Speech streams live into a small duplex-trained local model. **(2)** A mid-network (L22) linear read decides in ~ 30 ms — before the first output token — so the expert call is already in flight while local decoding would just be starting. **(3)** Easy queries stay local; hard ones escalate to a cloud expert. **(4)** The *same talker speaks* the expert’s answer in one continuous session, improving live spoken sessions (42.2%→63.3%) and, frozen, public speech benchmarks (Speech TriviaQA 66.4%→86.0%).

ure labels on a frozen pool of public-benchmark queries, in matched text and speech conditions. Throughout, *duplex* names the training regime, not the interaction protocol: we study duplex-trained weights driven through a *turn-based* loop whose only decision point is an end-of-turn read (§6); talker-on replay and concurrent listen/speak validations: Appendix H.

Everything runs on a frozen checkpoint — no fine-tuning; the only fitted component is a logistic regression on stored hidden states, calibrated once on text. That minimalism is the point: competing signals read the wrong place (late layers), arrive too late (post-answer verification), or must be trained in (think tokens); the mid-layer read survives all three. Three findings — find the signal, read it in time, show it matters:

1. **Representation: duplex fine-tuning makes the late-layer readouts brittle; a mid-network signal survives.** The conventional final-layer probe looks strong in distribution (AUC 0.82) but is mostly a query-type shortcut: under leave-one-pool-out transfer it *inverts* on held-out math (0.93→0.37). The damage is specific to the late-layer last-token position on text input — absent in raw backbones and non-duplex fine-tunes — consistent with late-layer specialization for duplex turn control. Mid-network (L22 of 36) the same linear read transfers at 0.93, and the mid-network optimum generalizes to a second, architecturally disjoint duplex family (hybrid-Mamba NemotronLabs-VoiceChat-11B).
2. **Modality: the mid-layer signal transfers text → speech in end-to-end models.** A probe calibrated on cheap text data reads speech-input hidden states at 0.86; the control that isolates the mechanism is Freeze-Omni, an adapter on frozen identical weights, where cross-modal transfer is dead (≤ 0.60). Verbalized self-assessment shows the mirror-image failure — it collapses on audio input — so the mid-layer probe is the one readout that survives both modalities. The transfer holds on real human recordings, not just TTS.
3. **System: the signal is available before the first token and works live, in and out of distribution.** The L22 read costs 20 ms (text) / 45 ms (audio) against time-to-first-token of 36/68 ms; in the streaming loop an end-of-turn read reaches AUC 0.887 in ~ 30 ms, half a second before first audible speech. Run live end-to-end — expert consulted mid-conversation, answer injected, talker speaks it — heard accuracy climbs 42.2%→63.3% at 57% escalation, and the frozen gate transfers to six public speech benchmarks with label-free thresholds, lifting Speech TriviaQA to 86.0% against the model’s official offline 75.5%.

Extended analyses and controls are in the appendix.

65 2 Related Work

66 **Full-duplex real-time dialogue.** End-to-end spoken dialogue spans native duplex designs —
67 dGSLM [Nguyen et al., 2023], Moshi [Défossez et al., 2024], GLM-4-Voice [Zeng et al., 2024],
68 single-LLM native full-duplex [Fang et al., 2026] — streaming omni models [Xu et al., 2025a, Cui
69 et al., 2026], frozen-backbone adapters [Wang et al., 2024b], and audio fine-tunes [Chu et al., 2024].
70 Inside these models the standard control interface for turn-taking is a policy *trained into the gener-*
71 *ator*: a neural-FSM scheme has the LLM emit explicit listen/speak transition tokens [Wang et al.,
72 2024a]; duplex time-slicing learns an idle token that keeps the model silent [Zhang et al., 2024]; Syn-
73 cLLM aligns the two speech streams with synchronization tokens [Veluri et al., 2024]; LSLM halts
74 generation on a detected-interruption token [Ma et al., 2024]; VITA and SALMONN-omni emit
75 state tokens around the speak/listen decision [Fu et al., 2024, Yu et al., 2024]. A neighboring line
76 moves the same decision into small classifier heads — FlexDuo [Liao et al., 2025], MinMo [Chen
77 et al., 2025], Freeze-Omni’s state heads — at 120–250 ms cadences that show a live loop tolerates
78 far more decision latency than our ~ 30 ms read [cf. Riera et al., 2026, Wu et al., 2026]. All of these
79 decide *when the local model should speak*; our gate decides the orthogonal question — *whether the*
80 *local model should be the one answering* — and their shared choice of a trained last-layer readout
81 is exactly the intervention our results show duplex fine-tuning damages.

82 **Does the model know what it knows?** $p(\text{True})$ self-evaluation [Kadavath et al., 2022], verbalized
83 uncertainty [Lin et al., 2022], and truthfulness probes on hidden states [Azaria and Mitchell, 2023,
84 Burns et al., 2023, Marks and Tegmark, 2024] established that correctness is readable internally;
85 intermediate layers carry more of it than the output distribution [Orgad et al., 2025]. Semantic
86 entropy [Kuhn et al., 2023, Farquhar et al., 2024] needs multiple sampled generations; Semantic
87 Entropy Probes [Kossen et al., 2024] amortize it into a hidden-state probe — the closest signal to
88 ours, but read post-hoc on text-only half-duplex models. Multimodal uncertainty fusion [Zhang
89 et al., 2026] asks a downstream question; we ask *where* in a duplex model a modality-robust read
90 survives at all.

91 **Selective escalation and routing.** The classic routers are creatures of the turn-based, text-only
92 setting: FrugalGPT [Chen et al., 2023], HybridLLM [Ding et al., 2024], RouteLLM [Ong et al.,
93 2024], and confidence-based cascades [Xue et al., 2025, Hu et al., 2024, Li et al., 2026] make one
94 decision per complete typed query, from query-side features or preference data, before any model
95 has answered; AutoMix [Aggarwal et al., 2024] escalates by *post-answer* self-verification, which
96 a voice deadline forbids. Both assumptions — query-side features, and the large model answering
97 the user directly — fail in full-duplex voice, where the session lives in the small model (query-side
98 ceiling quantified in §4.1). A second family trains the routing decision *into the weights* as special
99 tokens emitted during decoding: Toolformer [Schick et al., 2023] teaches the model to insert API-
100 call tokens by self-supervised fine-tuning; Self-RAG [Asai et al., 2024] fine-tunes the full model
101 on distilled critic labels so reflection tokens trigger retrieval and self-critique; Co-LLM [Shen et al.,
102 2024] and CITER [Zheng et al., 2025] learn a per-token deferral decision to a larger model; pause
103 and thought tokens [Goyal et al., 2024, Zelikman et al., 2024], fused think/no-think modes [Yang
104 et al., 2025, DeepSeek-AI et al., 2025], and the built-in think gate MiniCPM-o 4.5 itself ships [Cui
105 et al., 2026] route *compute* the same way. The design is uniform — extend the vocabulary, then
106 fine-tune the generator until the policy lives in the checkpoint — and so is the cost: supervision
107 must be constructed and the talker fine-tuned, the policy is frozen at training time (a new expert,
108 budget, or threshold means retraining), and in a duplex model fine-tuning the talker is exactly the
109 intervention §4.2 shows damages the last-layer readout. Our gate is the pluggable complement: a
110 logistic read over a frozen checkpoint, refit from stored activations in minutes, its budget moved
111 by a deployment-time threshold, transferring across modality and model family (§7) — and it beats
112 the model’s own trained think token at matched budget (Table 9). Learning-to-defer [Madras et al.,
113 2018] and early exit [Schuster et al., 2022] decide inside one model.

The gap. The prior work factors cleanly along two axes — {text-based, full-duplex} $\times$ {policy trained in as special tokens, probe read out of frozen internals}. Three cells are dense: text models both train routing tokens and read probes; duplex models train their control policies in as tokens (or heads). The fourth cell is empty: no prior work we know of reads a *frozen* full-duplex model’s internal state for a pre-answer, modality-robust competence signal — and fills it at the moment our results say token-training cannot be trusted there, because fine-tuning the talker is what damages the readout the trained policies sit on. That empty cell is this paper: a probe early enough to drive a live small-to-large handoff while preserving the session.

3 Experimental Setup

Target model and matched-pair controls. The system target is **MiniCPM-o 4.5** (9B; omni + full-duplex fine-tune of Qwen3-8B), bf16, greedy decoding, seed 42. To attribute effects to the *kind* of fine-tune rather than the backbone, we assemble ten models spanning raw backbones, a second duplex generation (MiniCPM-o 2.6), a streaming-omni fine-tune (Qwen2.5-Omni-7B), an audio-only fine-tune (Qwen2-Audio-7B), and a frozen-backbone adapter model (Freeze-Omni, LLM weights byte-identical to its raw control); the full roster is Table 2 (appendix). The comparison statistic is always the within-pair difference.

Query pool, labels, and audio arm. 600 queries in five pools, frozen before any gate experiment (360 calibration / 240 test): *trap* = SimpleQA [Wei et al., 2024] (the target fails 100%), *hard-knowledge* = MMLU-Pro [Wang et al., 2024c], *easy-fact* = TriviaQA [Joshi et al., 2017], *easy-chat* = Dolly/Alpaca [Conover et al., 2023, Taori et al., 2023], *hard-math* = GSM8K tail + MATH-500 [Cobbe et al., 2021, Hendrycks et al., 2021, Lightman et al., 2024]; public benchmarks only, no LLM-generated queries. Adequacy labels come from an LLM judge [Zheng et al., 2023] with structured outputs, blind to the gate and validated by a stronger re-judge (agreement 96.2% text / 94.5% audio; every headline number moves ≤ 1 point under the stronger labels); the escalation expert is gpt-5.5 at low reasoning effort. The audio arm renders the same frozen pool to speech (single-voice TTS, content unchanged — only modality varies), following Spoken-SQuAD / VoiceBench practice [Li et al., 2018, Chen et al., 2024]; audio-side findings are additionally validated on real human recordings (SD-QA, Faisal et al., 2021) and on the external benchmarks’ official audio (§7). Judge prompts, pins, and spend: Appendix M–N.

Gate signals (no model training). (i) decode scalars: max token entropy / margin over the first k decode steps; (ii) linear probes: logistic regression on a single hidden state, position $\times$ layer swept, fit on the calibration split only; (iii) verbalized self-evaluation [Kadavath et al., 2022]: $p(\text{True})$ -pre asks “would you answer this correctly?” before answering, $p(\text{True})$ -post asks about a drafted answer; $P(\text{Yes})$ is read from the first token’s distribution.

4 Reading Competence from Duplex Models

4.1 In distribution every readout looks fine; transfer reveals the shortcut

Table 3 compares the candidate signals on MiniCPM-o 4.5. In distribution the final-layer probe looks healthy at AUC 0.822 — but that number is mostly a *query-type* shortcut. The same hidden state all but names which pool a query came from, a pool-identity oracle alone recovers most of the probe’s in-distribution margin, and appending true pool labels to the probe’s features changes nothing, because the label is already in the representation. The decisive test is therefore leave-one-pool-out (LOPO) transfer, which we treat as the primary metric throughout. Under LOPO the readout does not merely weaken: on held-out math it *inverts*, to AUC 0.372, and it scores a trap pool the model fails 100% of the time at 0.232 — ranking the queries it is certain to get wrong as the safest to keep local. An external query-surface router fails the same way from the other side — never clearing the pool oracle in distribution, collapsing under LOPO, barely moving with two orders of



Figure 2: The core representational result: LOPO math AUC by layer and position. On the duplex model the standard (final-layer, last-token) read inverts on text input while mid-network layers hold ≥ 0.9 ; raw backbones show no cliff. The deployed gate reads L22.

160 magnitude more data: instance-level competence lives in the model’s state, not on the query surface
 161 (Appendices B, E).

162 Verbalized self-assessment transfers differently. Asked *before* answering, $p(\text{True})$ catches exactly
 163 the trap pool the probe misses, scoring it 0.945. But timing is everything: asked *after* drafting, the
 164 model endorses its own confident-wrong answer and the signal inverts. Post-draft $p(\text{True})$ is the
 165 strongest single readout in distribution, at 0.899, and also the one a voice deadline cannot afford —
 166 it bills a full draft before the gate can fire.

167 4.2 The damage is duplex-specific; the signal survives mid-network

168 Matched pairs attribute the probe failure to the *kind* of fine-tune, not the backbone: the raw Qwen
 169 backbone, label-matched to the duplex model’s fail rates, holds LOPO math at 0.825 where the du-
 170 plex model sits at 0.372; streaming-omni and audio-only fine-tunes are undamaged; a second duplex
 171 generation (MiniCPM-o 2.6) reproduces the collapse (Appendix B). A layer $\times$ position sweep local-
 172 izes it (Figure 2): the collapse is specific to (late layer $\times$ last token) on text input — mean-pooled
 173 reads survive at the same depth but 0.13–0.15 below the mid-layer peak (LOPO math ≈ 0.80 here,
 174 0.674 on MiniCPM-o 2.6), and mid-network the last-token read recovers to **0.931 at L22**, near the
 175 raw backbone.

176 *What breaks at the output.* The natural reading — that duplex training re-purposes the last position
 177 for turn-control state — is half right at best, and we tested it rather than asserting it (Appendix D).
 178 Four candidate mechanisms fail: the cliff does not move when the model is prefilled in speak mode,
 179 the inverting direction is orthogonal to the model’s own audio-vs-text axis, source pools stay equally
 180 decodable at every depth in both the duplex model and its backbone, and once features are stan-
 181 dardized the duplex model’s representational drift from that backbone does not localize late at all.
 182 What does separate them is how the probe’s read is *carried* (Figure 6): through mid-network the
 183 transferable signal is distributed across directions, and over the final layers of both duplex models
 184 — but neither raw backbone — that distributed component decays and inverts, on MiniCPM-o 4.5
 185 while the read concentrates into a handful of dominant directions. Duplex tuning does not delete
 186 the competence signal; it stops delivering it to the output position the standard readout uses. Direct
 187 tests of the turn-control reading sharpen this without overturning it: the duplex output layer alone
 188 stays biased toward emitting control tokens, yet the inverting direction aligns with neither the con-
 189 trol vocabulary nor an explicit listen/speak axis, so which duplex objective causes the loss remains
 190 open. The mid-network optimum is not a MiniCPM quirk: on NemotronLabs-VoiceChat-11B — a
 191 hybrid-Mamba duplex model, architecturally disjoint — the same recipe peaks mid-band at L30–34
 192 of 56 and transfers to the external pools at AUC 0.70–0.79 with a quarter of the calibration data
 193 (Appendix C).

194 4.3 Text-calibrated, speech-deployed — in end-to-end models only

195 Probes trained on text hidden states and scored on speech-input hidden states locate where the repre-
196 sentation becomes modality-shared: early layers are modality-specific (0.54–0.60), and from ~50%
197 depth text→audio transfer plateaus at 0.82–0.87 (0.855 at the deployed L22). This replicates on
198 MiniCPM-o 2.6 and cross-family on Qwen2.5-Omni. The decisive control is **Freeze-Omni**: identi-
199 cal frozen LLM weights with audio attached by an adapter — and cross-modal transfer is dead
200 (0.52–0.60). The modality-shared mid-layer core is *created by training the backbone on the modal-*
201 *ity*; the cheap recipe “calibrate on text, deploy on speech” is valid for end-to-end models and invalid
202 for adapter-attached ones. Verbalized introspection shows the mirror-image failure: on audio input,
203 pre-answer $p(\text{True})$ collapses on both MiniCPM duplex generations while both non-duplex controls
204 stay intact; controlled arms show the verbal self-check runs over text-token pathways that audio em-
205 beddings do not feed — duplicating the question as text restores it — while a linear probe reads the
206 same information off the same forward pass at 0.93 (Appendix B). All audio-side findings replicate
207 on 200 real human recordings (SD-QA), ruling out TTS artifacts. **The mid-layer probe is the one**
208 **readout that survives every training regime that can converse**; the live system deploys the same
209 L22 read plus two zero-latency pooled positions (the *deployed* probe; Appendix B).

210 5 From Signal to Gate

211 **The decision exists before decoding starts.** The gate is one 4096-d dot product on a hidden
212 state the forward pass already computes. Measured on H100 (wall-clock, prefill start to L22 score,
213 CUDA-synced; network excluded — the cloud call fires asynchronously; protocol: Appendix F), the
214 decision lands in **20 ms (text) / 45 ms (audio)** at P50 (P95 25/104) against same-bench time-to-first-
215 token of 36/68 ms: $T_{\text{gate}} < T_{\text{first token}}$, far inside the 200–300 ms turn-taking budget [Stivers et al.,
216 2009], and the cloud call can launch while the remaining ~40% of prefill and the whole decode still
217 run. Deciding pre-token is optimal, not merely feasible: one could wait out a few decoded tokens
218 and cancel, but decode-time readouts measure *weaker* (max entropy@4 0.696; early-decode hiddens
219 0.776 vs. 0.828 at the prompt), and every waited millisecond delays the stall onset and the expert
220 launch (Appendix F).

221 **Accuracy–cost tradeoff.** On the frozen test split ($n=240$): small-only 58.8%, big-only (gpt-5.5)
222 91.7%. Gated escalation beats random escalation at every operating point (Figure 7, appendix); the
223 balanced tier at a 33% budget reaches **77.9–78.7%**, recovering ~58% of the small→big gap at one
224 third of the big-only cost. In distribution the gate is statistically on par with a pool-oracle (type-only)
225 router — as the type-shortcut audit predicts — so its case over type routing rests on transfer, where
226 the type shortcut collapses (§4.1) and the mid-layer signal does not. The gate is fit to *local failure*,
227 not expert benefit: a benefit label would bake the current expert’s blind spots into a router meant to
228 survive expert swaps, and against the harder label the frozen scores stay usable out of distribution
229 (Appendix I). Against the model’s own built-in reasoning mode, gated-cloud escalation dominates
230 on both accuracy and latency at every rate (Appendix F).

231 6 Live End-of-Turn Escalation

232 **The streaming loop.** The protocol is turn-based: audio streams in 1 s chunks, but the talker speaks
233 only after the user’s turn ends and the model’s native listen/speak head is not engaged. Per-chunk
234 probe scores are strictly weaker than the offline read (best chunk statistic AUC 0.764) and miss short
235 utterances whose critical entity lands in the final chunk. The fix is an **end-of-turn read**: after the
236 last audio chunk, prefill an assistant turn containing a single space and read L22 at the assistant-
237 start position. This read reaches **AUC 0.887** — above the offline reference (0.843) — in ~30 ms,
238 still before the first output token, and is the loop’s only decision point; escalation-rate thresholds
239 are quantiles of this single score. Speculative mid-stream prefetch was measured and dropped (an
240 expert answer to a partial transcript answers the full question only 7–51% of the time; Appendix G).
241 On escalation the talker speaks a stall phrase, its *own* transcription of the audio goes to the expert

(gpt-5.5, reasoning effort fixed low — higher effort buys zero accuracy on the escalated population at double the latency), and the returned answer is injected as context for the talker to relay. On the deployed channel the relay is clean: model wrong, correct answer injected, comply = 1.00 ($n=24$); more broadly the competence signal also predicts relay compliance (comply 42%→75% across gate-score quartiles) — a model that knows it doesn’t know is willing to listen. Conflicting-injection and framing controls: Appendix G.

The live result. We ran 240 frozen test queries $\times$ four escalation tiers, each a complete turn-based spoken session, scored on *heard accuracy* — what a user would actually have heard. Headline numbers use a pre-registered input-side filter (21 formula-LaTeX ids + 1 broken recording removed; spoken formulas are destroyed by any TTS render before the system hears them). Heard accuracy climbs monotonically, 42.2% → 52.8% → 59.2% → 63.3% at 0/14/31/57% escalation, every tier’s gain significant (paired bootstrap; Figure 12). A measured always-escalate arm closes the curve at 66.5% (P50 4.2 s): the last 43 points of escalation buy 3.2 points, because past the gated tiers the binding constraint is the transcription-uplink channel, not selection — the same sessions re-scored with gold-text expert answers reach 92.2%. The remaining gap to the channel-controlled counterfactual (escalated rows re-scored with gold-text expert answers) is dominated by *transcription-uplink* errors, not gate failure: the channel costs +2.3/+6.9/+12.8 points across tiers, and against the matched-rate random reference *internal to each view* the gate wins at every arm — in the deployed heard view by 7.2–9.5 points, in the channel-controlled view by 5.4–8.4 (Figure 12). Live arm accuracies on this pool carry a replication-noise floor of ± 2 –3 points, so no claim here rests on a sub-3-point live delta. Full tables (both views, both pools), latency profiles, uplink diagnostics, and two boundary query classes (false-premise, real-time) are in Appendix G. A talker-on replay leaves the frozen probe intact (AUC 0.871/0.856 vs. 0.866 headless); TTS-on, first audible audio lands p50 0.55 s on local turns and 22–29 ms on escalated ones (the canned stall), expert content at p50 6.4 s with the stall covering the head of the wait (Appendix H, G). Re-run end to end in the *concurrent* regime (audio interleaved while the talker speaks, in-regime gate) the loop reproduces: monotone gains on all seven pools, speakable 44.5→71.1%, selective still beating always-escalate (Table 12).

7 Frozen Transfer to Public Speech Benchmarks

The out-of-distribution test the type-shortcut audit demands: we take the deployed gate *frozen* — weights, features, and layer choice fit once on the internal calibration mix — and run the complete live loop of §6 on six public speech benchmarks: Speech TriviaQA, Speech Web Questions, Llama Questions, and Reasoning QA from OpenAudioBench [Li et al., 2025] (spoken renderings of Joshi et al., 2017, Berant et al., 2013, and the set of Nachmani et al., 2024), SD-QA (real dialectal human speech) [Faisal et al., 2021], and VoiceBench AlpacaEval [Chen et al., 2024, Taori et al., 2023]. All audio is the benchmarks’ official pre-rendered audio; no query ever touches probe calibration. The only per-pool adaptation is *label-free*: tier thresholds are quantiles of the probe’s own score distribution on each pool, consuming no labels and no judge calls. It buys budget control, not discrimination: one absolute threshold frozen on the internal pool still beats matched-rate random on four of five pools, but fires at 2–48% across them (Appendix I). Where an official protocol exists we score with it (the OpenAudioBench judge; VoiceBench’s 1–5 judge), and our chat-mode control reproduces the official anchors on-scale.

The frozen gate improves all five QA benchmarks (Table 1), monotonically in budget. On Speech TriviaQA the 9B system reaches **86.0% live** — against its own official offline 75.5% and 62.9% for Qwen3-Omni-30B [Xu et al., 2025b] at $3.3\times$ the parameters: selective cloud escalation outperforms larger standalone backbones. The cleanest evidence that the signal is difficulty rather than surface is cross-lingual: an English-only calibration expansion moved *Chinese* Reasoning QA from AUC 0.621 to 0.683, the largest per-pool gain. On Llama Questions — single-type, so no shortcut to exploit — **selective escalation matches the expert at half the expert calls**: the probe’s split is decisive (kept-local 97.6% vs. escalated 69.6%, $z=6.46$), the expert’s whole advantage lands where the probe sends it (69.6%→88.8%, $p<.0001$), and a measured live always-escalate arm lands *below*

Table 1: Main transfer results: methods $\times$ five public speech QA benchmarks (TriviaQA/WebQ/Llama Q. OAB official judge, $n = 250$; SD-QA/Reasoning zh reference-anchored, $n = 200/202$); heard accuracy in %, $\pm 2\text{--}3$ point replication floor. Bottom: second duplex family (NemotronLabs-VoiceChat-11B, same frozen recipe, offline re-mix, official judges, Appendix C); English-only, so its avg spans four benchmarks over its own floor, and its gate clears matched-rate random on every pool at every budget (per-budget random references and permutation tests: Table 7). Gold-inject view at matched budgets: Table 14. Δ = avg gain over always-local (points). AlpacaEval (1–5): 3.99 $\rightarrow$ 4.35 gated, 4.76 always-escalate, no gain over random on MiniCPM; the second family does gain (3.55 $\rightarrow$ 4.46 vs. random 4.23).

	TriviaQA	WebQ	Llama Q.	SD-QA	Reason. zh	avg	Δ
<i>live heard view (deployed loop)</i>							
always-local	66.4	57.2	84.0	51.5	58.9	63.6	—
+ gate @15%	73.2	62.8	90.4	60.0	65.3	70.3	+6.7
+ gate @30%	80.0	68.0	90.4	72.0	71.3	76.3	+12.7
+ gate @50%	86.0	73.2	94.8	79.5	77.2	82.1	+18.5
always-escalate (measured)	88.0	80.8	93.6	88.5	83.7	86.9	+23.3
always-expert (gold-text ceiling)	96.8	86.4	92.8	93.0	87.1	91.2	+27.6
<i>second family: NemotronLabs-VoiceChat-11B (offline re-mix, official judges)</i>							
always-local	35.6	37.6	70.5	31.0	—	43.7	—
+ gate @15%	49.8	48.0	78.6	42.5	—	54.7	+11.0
+ gate @30%	61.5	58.8	83.8	54.5	—	64.6	+21.0
+ gate @50%	74.1	69.6	87.6	67.5	—	74.7	+31.0
matched random @50%	65.6	60.2	81.8	60.0	—	66.9	+23.2
always-expert (ceiling)	95.5	82.8	93.2	89.0	—	90.1	+46.4
probe AUC (frozen)	0.789	0.785	0.806	0.792	0.683	0.771	—
probe AUC (2nd family)	0.775	0.771	0.720	0.754	—	0.755	—
official offline	75.5	70.2	—	—	—	—	—

selective (93.6 vs. 94.8 at half the calls); matched-rate precision reaches 95% of the base-rate cap (Appendix J).

Two honest negatives: **AlpacaEval** has no “missing fact” event to read, and **Web Questions and SD-QA are flat** (strict alias labels; the failures the gate ranks highest are ones the expert misses too). Neither is intrinsic — on the second duplex family the same recipe beats matched-rate random on *all five* pools, AlpacaEval included ($p \leq .0004$), so the flat pools are a headroom property of the pool $\times$ backbone pairing (Appendices C, J).

8 Discussion

What the signal detects. The pre-answer probe is a *retrieval-failure* detector: strongest where an empty lookup is observable, blind to false premises, and complementary to decode-time signals, which motivates a decode-time check behind it (Appendix G).

Limitations. (i) One phrasing per $p(\text{True})$ variant; real-speech validation is scripted read speech in one dialect; formula-symbolic content is out of scope for the audio arm (any TTS render destroys it). (ii) The second-family replication covers mid-band readability and frozen-recipe transfer only — not matched-pair attribution, the live loop, or the concurrent regime. (iii) Labels come from one judge family (a stronger re-judge bounds deltas at ≤ 1 point), no human labels. (iv) One session per query per tier ($\pm 2\text{--}3$ point floor); under concurrent prefill the frozen read degrades (AUC .87 $\rightarrow$.76) and an in-regime refit recovers .82 — calibration must follow the regime (Appendix H).

Conclusion. Duplex-trained speech models carry a competence signal their standard readouts miss — a mid-network read, available half a second before the first audible word. Never gone; only misread.

References

Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay,

316 Manaal Faruqui, and Mausam. AutoMix: Automatically mixing language models. In *Advances*
317 *in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2310.12963.

318 Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to
319 retrieve, generate, and critique through self-reflection. In *International Conference on Learning*
320 *Representations (ICLR)*, 2024. arXiv:2310.11511.

321 Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings*
322 *of the Association for Computational Linguistics: EMNLP 2023*, 2023. arXiv:2304.13734.

323 Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from
324 question-answer pairs. In *Proceedings of the 2013 Conference on Empirical Methods in Natural*
325 *Language Processing (EMNLP)*, 2013. WebQuestions.

326 Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in lan-
327 guage models without supervision. In *International Conference on Learning Representations*
328 *(ICLR)*, 2023. arXiv:2212.03827.

329 Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while
330 reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.

331 Qian Chen, Yafeng Chen, Yanni Chen, et al. MinMo: A multimodal large language model for
332 seamless voice interaction. *arXiv preprint arXiv:2501.06282*, 2025.

333 Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench:
334 Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024.

335 Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, et al. Qwen2-Audio technical
336 report. *arXiv preprint arXiv:2407.10759*, 2024.

337 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al.
338 Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. GSM8K.

339 Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick
340 Wendell, Matei Zaharia, and Reynold Xin. Free dolly: Introducing the world’s first truly open
341 instruction-tuned LLM. Databricks blog, 2023. databricks-dolly-15k.

342 Junbo Cui, Bokai Xu, Chongyi Wang, Tianyu Yu, Weiyue Sun, Yingjing Xu, et al. MiniCPM-o 4.5:
343 Towards real-time full-duplex omni-modal interaction. *arXiv preprint arXiv:2604.27393*, 2026.

344 DeepSeek-AI, Daya Guo, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via
345 reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

346 Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou,
347 Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dia-
348 logue. *arXiv preprint arXiv:2410.00037*, 2024.

349 Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks
350 V. S. Lakshmanan, and Ahmed Hassan Awadallah. Hybrid LLM: Cost-efficient and quality-
351 aware query routing. In *International Conference on Learning Representations (ICLR)*, 2024.
352 arXiv:2404.14618.

353 Fahim Faisal, Sharlina Keshava, Md Mahfuz ibn Alam, and Antonios Anastasopoulos. SD-QA:
354 Spoken dialectal question answering for the real world. In *Findings of the Association for Com-*
355 *putational Linguistics: EMNLP 2021*, 2021. arXiv:2109.12072.

356 Qingkai Fang, Shoutao Guo, and Yang Feng. BayLing-Duplex: Native full-duplex speech dialogue
357 with a single autoregressive LLM. *arXiv preprint arXiv:2606.14528*, 2026.

358 Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large
359 language models using semantic entropy. *Nature*, 630:625–630, 2024.

360 Chaoyou Fu, Haojia Lin, Zuwei Long, Yunhang Shen, Yuhang Dai, Meng Zhao, et al. VITA: To-
361 wards open-source interactive omni multimodal LLM. *arXiv preprint arXiv:2408.05211*, 2024.

362 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
363 Nagarajan. Think before you speak: Training language models with pause tokens. In *ICLR*, 2024.

364 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
365 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
366 *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks*, 2021.
367 arXiv:2103.03874.

368 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
369 Keutzer, and Shriyash Kaustubh Upadhyay. RouterBench: A benchmark for multi-LLM routing
370 systems. *arXiv preprint arXiv:2403.12031*, 2024.

371 Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale dis-
372 tantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual*
373 *Meeting of the Association for Computational Linguistics (ACL)*, 2017.

374 Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez,
375 Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language mod-
376 els (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

377 Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. Se-
378 mantic entropy probes: Robust and cheap hallucination detection in LLMs. *arXiv preprint*
379 *arXiv:2406.15927*, 2024.

380 Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for
381 uncertainty estimation in natural language generation. In *International Conference on Learning*
382 *Representations (ICLR)*, 2023. arXiv:2302.09664.

383 Chia-Hsuan Li, Szu-Lin Wu, Chi-Liang Liu, and Hung-yi Lee. Spoken SQuAD: A study of miti-
384 gating the impact of speech recognition errors on listening comprehension. In *Interspeech*, 2018.
385 arXiv:1804.00320.

386 Hao Li, Yiqun Zhang, Zhaoyan Guo, Chenxu Wang, Shengji Tang, Qiaosheng Zhang, Yang Chen,
387 Biqing Qi, Peng Ye, Lei Bai, Zhen Wang, and Shuyue Hu. LLMRouterBench: A massive bench-
388 mark and unified framework for LLM routing. *arXiv preprint arXiv:2601.07206*, 2026.

389 Tianpeng Li, Jun Liu, Tao Zhang, Yuanbo Fang, Da Pan, Mingrui Wang, et al. Baichuan-Audio:
390 A unified framework for end-to-end speech interaction. *arXiv preprint arXiv:2502.17239*, 2025.
391 Introduces OpenAudioBench.

392 Borui Liao, Yulong Xu, Jiao Ou, Kaiyuan Yang, Weihua Jian, Pengfei Wan, and Di Zhang. Flex-
393 Duo: A pluggable system for enabling full-duplex capabilities in speech dialogue systems. *arXiv*
394 *preprint arXiv:2502.13472*, 2025.

395 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
396 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International*
397 *Conference on Learning Representations (ICLR)*, 2024. arXiv:2305.20050; source of the MATH-
398 500 subset.

399 Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in
400 words. *Transactions on Machine Learning Research*, 2022. arXiv:2205.14334.

401 Ziyang Ma, Yakun Song, Chenpeng Du, Jian Cong, Zhuo Chen, Yuping Wang, Yuxuan Wang, and
402 Xie Chen. Language model can listen while speaking. *arXiv preprint arXiv:2408.02622*, 2024.

David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and accuracy by learning to defer. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2310.06824.

Eliya Nachmani, Alon Levkovitch, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh Mariooryad, Ehud Rivlin, RJ Skerry-Ryan, and Michelle Tadmor Ramanovich. Spoken question answering and speech continuation using spectrogram-powered LLM. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2305.15255; Spectron, source of the Llama Questions dataset.

Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden Tomasello, Robin Algayres, Benoît Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux. Generative spoken dialogue language modeling. *Transactions of the Association for Computational Linguistics*, 2023. arXiv:2203.16502.

Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data. *arXiv preprint arXiv:2406.18665*, 2024.

Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, and Yonatan Belinkov. LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2410.02707. TODO verify author list.

Pablo Riera, Pablo Brusco, Cristina Kuo, Marcelo Sancinetti, and S. R. K. Branavan. Synchronization and turn-taking in full-duplex speech dialogue models. *arXiv preprint arXiv:2605.20356*, 2026.

Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2302.04761.

Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2207.07061.

Shannon Zejiang Shen, Hunter Lang, Bailin Wang, Yoon Kim, and David Sontag. Learning to decode collaboratively with multiple language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. arXiv:2403.03870; Co-LLM.

Tanya Stivers, N. J. Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann, Gertie Hoymann, Federico Rossano, Jan Peter de Ruiter, Kyung-Eun Yoon, and Stephen C. Levinson. Universals and cultural variation in turn-taking in conversation. *Proceedings of the National Academy of Sciences*, 106(26):10587–10592, 2009.

Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following LLaMA model. GitHub repository, 2023. zh split: shibing624/alpaca-zh.

Bandhav Veluri, Benjamin N. Peloquin, Bokai Yu, Hongyu Gong, and Shyamnath Gollakota. Beyond turn-based interfaces: Synchronous LLMs as full-duplex dialogue agents. In *Proceedings of EMNLP (Main)*, 2024. arXiv:2409.15594.

448 Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan
449 Sung, Denny Zhou, Quoc Le, and Thang Luong. Freshllms: Refreshing large language models
450 with search engine augmentation. *arXiv preprint arXiv:2310.03214*, 2023.

451 Peng Wang, Songshuo Lu, Yaohua Tang, Sijie Yan, Wei Xia, and Yuanjun Xiong. A full-duplex
452 speech dialogue scheme based on large language models. In *Advances in Neural Information*
453 *Processing Systems (NeurIPS)*, 2024a. arXiv:2405.19487.

454 Xiong Wang, Yangze Li, Chaoyou Fu, Lei Xie, Ke Li, Xing Sun, and Long Ma. Freeze-Omni:
455 A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv preprint*
456 *arXiv:2411.00774*, 2024b.

457 Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, et al. MMLU-Pro: A more robust and
458 challenging multi-task language understanding benchmark. In *Advances in Neural Information*
459 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2024c. arXiv:2406.01574.

460 Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese,
461 John Schulman, and William Fedus. Measuring short-form factuality in large language models.
462 *arXiv preprint arXiv:2411.04368*, 2024. SimpleQA.

463 Donghang Wu, Haoyang Zhang, Chen Chen, Tianyu Zhang, Fei Tian, Xuerui Yang, Gang Yu, Hexin
464 Liu, Nana Hou, Yuchen Hu, and Eng Siong Chng. Chronological thinking in full-duplex spoken
465 dialogue language models. In *Proceedings of the 27th Annual Meeting of the Special Interest*
466 *Group on Discourse and Dialogue (SIGDIAL)*, 2026. arXiv:2510.05150.

467 Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, et al. Qwen2.5-
468 Omni technical report. *arXiv preprint arXiv:2503.20215*, 2025a.

469 Jin Xu, Zhifang Guo, Hangrui Hu, Yunfei Chu, Xiong Wang, et al. Qwen3-Omni technical report.
470 *arXiv preprint arXiv:2509.17765*, 2025b.

471 Leyang Xue, Yao Fu, Luo Mai, and Mahesh K. Marina. HybridServe: Efficient serving of large AI
472 models with confidence-based cascade routing. *arXiv preprint arXiv:2505.12566*, 2025.

473 An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

474 Wenyi Yu, Siyin Wang, Xiaoyu Yang, Xianzhao Chen, Xiaohai Tian, Jun Zhang, Guangzhi Sun,
475 Lu Lu, Yuxuan Wang, and Chao Zhang. SALMONN-omni: A codec-free LLM for full-duplex
476 speech understanding and generation. *arXiv preprint arXiv:2411.18138*, 2024.

477 Eric Zelikman, Georges Harik, Yijia Shao, Varuna Jayasiri, Nick Haber, and Noah D. Goodman.
478 Quiet-STaR: Language models can teach themselves to think before speaking. *arXiv preprint*
479 *arXiv:2403.09629*, 2024.

480 Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong,
481 and Jie Tang. GLM-4-Voice: Towards intelligent and human-like end-to-end spoken chatbot.
482 *arXiv preprint arXiv:2412.02612*, 2024.

483 Dengjia Zhang, Alexander Martin, William Jurayj, Kenton Murray, Benjamin Van Durme, and Reno
484 Kriz. Unified multimodal uncertain inference. *arXiv preprint arXiv:2604.08701*, 2026.

485 Xinrong Zhang, Yingfa Chen, Shengding Hu, Xu Han, Zihang Xu, Yuanwei Xu, Weilin Zhao,
486 Maosong Sun, and Zhiyuan Liu. Beyond the turn-based game: Enabling real-time conversations
487 with duplex models. In *Proceedings of EMNLP*, 2024. arXiv:2406.15718.

488 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang,
489 Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica.
490 Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information*
491 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2023. arXiv:2306.05685.

492 Wenhao Zheng, Yixiao Chen, Weitong Zhang, Souvik Kundu, Yun Li, Zhengzhong Liu, Eric P.
493 Xing, Hongyi Wang, and Huaxiu Yao. CITER: Collaborative inference for efficient large language
494 model decoding with token-level routing. In *Conference on Language Modeling (COLM)*, 2025.
495 arXiv:2502.01976.

496 A Model roster

Table 2: Model roster. “Pair” groups a fine-tune with its raw backbone control; the comparison statistic is always the within-pair difference.

model	type	pair	role
MiniCPM-o 4.5	omni + duplex FT	Qwen3-8B	system target
Qwen3-8B	raw	—	pair control
MiniCPM-o 2.6	duplex FT	Qwen2.5-7B	duplex replication
Qwen2.5-Omni-7B	streaming omni FT	Qwen2.5-7B	omni control
Qwen2.5-7B	raw	—	pair control
Qwen2-Audio-7B	audio FT	Qwen2-7B	audio-FT control
Freeze-Omni	frozen LLM + adapter	Qwen2-7B	adapter control
Qwen2-7B	raw	—	pair control
Mistral-7B-v0.3	raw	—	backbone diversity
GLM4-9B-chat	raw	—	backbone diversity

497 B Extended signal and readout analyses

Table 3: Signals on MiniCPM-o 4.5, judged failure labels. In-dist = calibration-split AUC (probes 5-fold out-of-fold); LOPO math = leave-one-pool-out AUC on held-out math; audio = the same readout under speech input (final/mid-layer probe: audio-input LOPO math / text-only calibration read on audio states). Full per-signal detail: Appendix B.

signal	in-dist	LOPO math	audio input
max entropy@4	0.696	≈chance on raw backbones	—
$p(\text{True})$ -pre	0.807	trap 0.945 (probe: 0.232)	collapses
$p(\text{True})$ -post (full draft)	0.899	—	degrades
final-layer probe	0.822	0.372 (inverts)	0.936 (intact)
mid-layer (L22) probe	0.866	0.931	0.855 (text-calibrated)

498 **The in-distribution shortcut, quantified.** The numbers behind §4.1’s claim that the final-layer
 499 probe’s in-distribution AUC of 0.822 is mostly pool identity: a linear read of the same hidden state
 500 classifies a query’s source pool at 95.8% accuracy; a pool-identity oracle alone reaches AUC 0.715;
 501 and appending true pool labels to the probe’s features moves it 0.822→0.821. In-distribution area
 502 over random is statistically on par between the gate and the pool oracle (Appendix E); the two
 503 separate only under LOPO, where the probe inverts to 0.372 and the oracle is undefined.

504 **ROC curves.** Figure 3 shows the calibration-split ROC for the signals of Table 3.

505 $p(\text{True})$ **backbone dependence.** Across backbones, post-draft $p(\text{True})$ beats the trained final-
 506 layer probe on four of five raw/duplex models tested; the exception (GLM4-9B-chat, 0.685 vs. probe
 507 0.798) shows the pattern depends on the backbone’s self-evaluation quality.

508 **Matched pairs.** Table 4: with label coverage matched by subsampling the raw control to the
 509 duplex model’s per-pool fail rates, the raw backbone keeps LOPO math at 0.825/0.809 while duplex
 510 models sit at chance — label coverage is ruled out; the representation changed. Degradation is
 511 ordered raw > omni-streaming > duplex. Table 5 gives the by-depth summary behind Figure 2:
 512 the within-pair deficit at the best mid-layer is only −0.03 vs. −0.59 at the final layer; mean-pooled
 513 probes survive to the final layer on duplex models (LOPO math ≈0.80 at L35 on o4.5, 0.674 at L27
 514 on o2.6 — 0.13–0.15 below each model’s mid-layer last-token peak, and uniformly weaker in-mix,
 515 −0.058 at L22); duplex damage is severe but local, omni-streaming damage mild but diffuse.




Figure 3: ROC on the calibration split. $p(\text{True})$ -post leads but needs a full drafted answer; the probes and $p(\text{True})$ -pre are pre-answer.

Table 4: LOPO transfer (AUC, held-out pool) of the final-layer last-token probe, Qwen2.5-7B family, label-coverage matched.

LOPO pool	Qwen2.5-7B (raw)	Qwen2.5-Omni	MiniCPM-o 2.6	[o4.5 vs Qwen3]
hard-math	0.809	0.745	0.538 $\approx$ chance	0.372 (inverts)
hard-knowledge	0.680	0.613	0.526 $\approx$ chance	0.61

516 **Modality specificity of the cliff.** Under audio input the final layer does *not* invert (L35 audio
517 LOPO math 0.936 vs. 0.366 text); a speak-mode template control on text input leaves the cliff
518 unchanged (0.362) — the operative variable is the modality of the input context, not the output
519 mode.

520 **Audio collapse of verbalized introspection.** Table 6. Three controlled arms pin the mechanism
521 on MiniCPM-o 4.5: *not perception* (ASR audit WER 0.074; $p(\text{True})$ on the model’s own transcript
522 snaps back to 0.074; the well-heard subset still collapses); *not a context prior* (irrelevant filler audio
523 does not inflate p_{yes}); *binding* (duplicating the question as text alongside the audio fully restores
524 introspection, 0.034). The verbal instance-check runs over text-token pathways that audio embed-
525 dings do not feed, while a linear probe reads the same information off the same forward pass at
526 0.93+. The failure implies its own cheap fix, confirmed twice: re-present the query as text before
527 asking (“repeat-then-judge”).

528 **Real-speech validation.** The matched-pair design rerun on 200 real human recordings
529 (VoiceBench SD-QA, USA split) replicates all three audio-side findings: modality tax
530 (0.400 $\rightarrow$ 0.450), audio overconfidence (paired p_{yes} shift +0.089; failure subset 0.415 $\rightarrow$ 0.581), and
531 the layer structure (text-calibrated read on real-speech audio holds 0.76–0.80 at L22–L26, peak
532 0.800 at L25 — transfer across both content and modality).

533 **Synthesis.** End-to-end multimodal training *builds* the shared mid-layer core; duplex-style training
534 additionally *damages* the readouts; omni-streaming gets both right; an adapter alone gets neither

Table 5: LOPO math transfer by depth (last-token probes).

model	best mid-layer	final layer	shape
Qwen3-8B (raw)	0.964 (L33/36)	0.958	plateau, no cliff
MiniCPM-o 4.5 (duplex)	0.931 (L22/36)	0.366	cliff, inverts
Qwen2.5-7B (raw)	0.893 (L21/28)	0.809	mild dip
Qwen2.5-Omni (streaming)	0.794 (L16/28)	0.746	diffuse depression
MiniCPM-o 2.6 (duplex)	0.822 (L21/28)	0.540	cliff

Table 6: Pre-answer introspection under audio input. Trap p_{yes} = mean stated probability of answering correctly on a 100%-fail pool (lower is better).

model	FT type	trap p_{yes} : text $\rightarrow$ audio	audio $p(\text{True})$ -pre AUC
MiniCPM-o 4.5	duplex	0.055 $\rightarrow$ 0.556 (collapse)	0.786 (trap dead)
MiniCPM-o 2.6	duplex	0.196 $\rightarrow$ 0.345	0.491 $\approx$ chance
Qwen2.5-Omni	omni-streaming	0.279 $\rightarrow$ 0.213 (intact)	0.727
Freeze-Omni	frozen + adapter	0.165 $\rightarrow$ 0.112 (intact)	0.612 (capability caveat)

535 (probe readout text-fragile, verbal readout audio-fragile; in both cases knowledge survives and a
 536 readout breaks, and in both cases the breakage is duplex-fine-tune-specific).

537 **The deployed probe.** The mechanistic probe is a single-position L22 read. The deployed probe
 538 adds two zero-latency pooled positions (an 8-token tail mean and the running user-audio mean), re-
 539 fit on a calibration mix widened to 2,310 public-benchmark queries disjoint from every evaluation
 540 benchmark; a pre-registered guard held (frozen internal test AUC 0.860 $\rightarrow$ 0.879) and feature selec-
 541 tion never saw the externals. In the mid-layer band the in-mix tradeoff optimum is broad (L20–L30)
 542 and the mid-layer probe’s in-distribution area advantage is modest — its real advantage is LOPO
 543 robustness, which an in-mix test cannot show.

544 C Second duplex family: NemotronLabs-VoiceChat-11B

545 The pre-registered prediction test: with the identical recipe (same audio, same judge, 600-query
 546 calibration), the layer sweep on this hybrid-Mamba duplex model (56 blocks: 27 Mamba-2 / 4
 547 attention) peaks mid-band at L30–34 (eot-last AUC 0.714 vs. 0.693/0.682 at the endpoints); stacking
 548 the same three positions lifts OOF 0.714 $\rightarrow$ 0.761 $\rightarrow$ 0.790; the frozen-methodology probe transfers
 549 to the external pools at AUC 0.781/0.793/0.754/0.701 — the deployed-probe band reached with a
 550 quarter of the calibration data. Boundary notes: the backbone is much weaker at knowledge retrieval
 551 (striviaqa local 0.32 vs. MiniCPM 0.62, same judge) and English-only (the cross-lingual pool has no
 552 analog). The live-loop port and the trained-checkpoint ablation remain future work; the port is gated
 553 by the released inference path, not by the architecture. The only supported loading path (the NeMo
 554 Speech nemotron-labs-voicechat branch) runs the duplex loop *cacheless* — every 80 ms frame
 555 re-runs the full prefix, an $O(T^2)$ schedule that cannot hold the frame clock in real time — although
 556 the hybrid backbone’s Mamba-2 state is exactly the $O(1)$ per-frame recurrence a streaming port
 557 would need; wiring stateful inference into the released duplex frame protocol is the engineering gap.
 558 (Offline, the same property is a convenience: the final frame’s forward contains every position, so
 559 one hook captures the whole eot window.) Accuracy-side conclusions do not wait on the port: the
 560 re-mix arithmetic below tracks measured live arms at 0.011 mean deviation on MiniCPM, inside the
 561 ± 0.02 –0.03 replication floor.

562 **Escalation re-mix on every transfer pool.** Completing the grid of Table 1 on this family: an of-
 563 fline re-mix (the same top- r reconstruction arithmetic that, on MiniCPM, tracks the measured live
 564 arms at 0.011 mean deviation) sends the top- r of each pool by NVDA probe score to the measured
 565 GPT-5.5 outcome and keeps NVDA’s own answer for the rest, scored with each pool’s official proto-
 566 col (the OAB judge for the three OpenAudioBench pools — NVDA local floors 0.352/0.376/0.705 —

our judge for SD-QA, VoiceBench 1–5 for AlpacaEval; the expert answers the heard transcript and is judged directly, so the relay-back channel is excluded — including the measured relay outcomes where they exist shifts tier points by ≤ 0.032). **Selective escalation beats matched-rate random on all five pools** (permutation $p \leq 0.0004$ at the 30% and 50% budgets, $p < 0.03$ at 15%; Table 7, Figure 4) — including Web Questions, SD-QA and AlpacaEval, the three pools that are flat or negative on MiniCPM (§7). Selectivity is nearly unchanged (AUC 0.72–0.78 under official labels); what changed is headroom — floors of 0.31–0.38 against MiniCPM’s 0.51–0.66 — and on AlpacaEval the failure species: the terse English-only model fails open-ended queries *hard* (the probe-selected half scores 3.01 vs. 4.08 kept), where MiniCPM’s open-ended failures are soft quality gradations a wrongness detector cannot see. The negatives are properties of the pool $\times$ backbone pairing, not of the signal.

Table 7: The frozen-recipe NVDA probe re-mixed at fixed escalation rates, official judge per pool, each against its own matched-rate random reference (same measured outcomes, escalation set drawn at random; identical at 0% and 100% by construction). AlpacaEval rows are VoiceBench 1–5 scores; others accuracy. Every pool clears random at all three budgets: $p \leq 0.0004$ at 30% and 50%, $p < 0.03$ at 15% (20k permutations, `scripts/18_nvda_random_check.py`).

pool	judge	n	0%	15%	30%	50%	100%
Speech TriviaQA	OAB	247	0.356	0.498	0.615	0.741	0.955
matched random				0.446	0.536	0.656	
Speech Web Questions	OAB	250	0.376	0.480	0.588	0.696	0.828
matched random				0.444	0.512	0.602	
Llama Questions	OAB	234	0.705	0.786	0.838	0.876	0.932
matched random				0.739	0.773	0.818	
SD-QA	ours	200	0.310	0.425	0.545	0.675	0.890
matched random				0.397	0.484	0.600	
AlpacaEval (1–5)	VB	199	3.55	3.83	4.13	4.46	4.92
matched random				3.75	3.96	4.23	




Figure 4: Re-mix curves for the frozen-recipe NVDA probe: selective (top- r by probe score) vs. matched-rate random on the five transfer pools, official judge per pool. Every pool separates, including the three that are flat on MiniCPM (all five pools, $p \leq 0.0004$ at 30/50% and $p < 0.03$ at 15%, permutation).




Figure 5: Second-family test: (a) the usable read again peaks mid-network; (b) frozen-recipe transfer lands in the deployed-probe AUC band.

D What breaks at the output: mechanism audit

§4.2 reports *that* the final-layer last-token read inverts. This section reports what we could and could not establish about *why*. Everything below uses the Phase-5d all-layer prefill dumps (every decoder layer, last-token and mean-pooled, same 600 queries) and the identical probe protocol as the layer sweep: calibration split, L_2 logistic regression, LOPO trained on the four non-math pools and scored on held-out math. The reproduction check recovers the published L22 and final-layer numbers to ≤ 0.012 (scripts/14-17_*.py).

Four mechanisms ruled out. Table 8 lists them. (i) *Speak mode*. Prefilling the same text queries under the TTS/speak-mode template leaves the cliff exactly where it was (final-layer LOPO math 0.362 vs. 0.366 plain), so the readout is not being displaced by a prepare-to-speak state. (ii) *Modality axis*. The mean audio-minus-text direction at the final layer is orthogonal to both the inverting final-layer probe and the transferring L22 probe (cosine 0.013 and 0.027; chance is $1/\sqrt{4096} = 0.016$), so the cliff is not the audio/text split that §4.3 measures. (iii) *More shortcut available late*. Source-pool identity is linearly decodable at 0.88–0.96 at *every* depth, in the duplex model and its raw backbone alike — the query-type shortcut of §4.1 is not something the late layers add; it is what remains once the competence signal stops being readable. (iv) *Wholesale representational drift*. Raw linear CKA between the duplex model and its backbone appears to collapse late (0.80→0.47), but transformer residual streams carry a few massive-activation coordinates that dominate the Gram matrix; with per-feature standardization the same comparison is nearly flat (0.82→0.78) and does not localize the cliff. Projecting out the massive-activation axes, or the pool-mean directions, rescues nothing (0.35–0.40 at every $k \leq 20$, against a random-direction control of 0.37).

Table 8: Mechanism candidates for the final-layer inversion. Final-layer LOPO hard-math AUC on MiniCPM-o 4.5 unless noted; 0.366 is the unmodified baseline and 0.931 the L22 read.

candidate	test	verdict
prepare-to-speak state	speak-mode template prefill	0.362, no change
modality (audio vs. text) axis	cosine with the probe direction	0.013 $\approx$ chance
late layers encode more pool identity	pool decoding by depth	0.88–0.96 at all depths
representation rewritten late	standardized CKA vs. backbone	0.82→0.78, flat
massive-activation coordinates	project out top- k such axes	0.35–0.40, no rescue
query-type directions	project out pool-mean directions	0.34–0.45, no rescue
delivery to the output position	residual component of the read	0.93→0.27 (Fig. 6)

What does separate the models. Figure 6 measures the probe rather than operating on the model. For the probe trained at each depth on the four non-math pools, we split its held-out score into the part carried by that layer’s five dominant principal directions (estimated from training rows only)

and the residual, $s = w^\top Px + w^\top (I - P)x$. The residual component is the distributed read, and it is the one that transfers: it holds 0.85–0.93 through mid-network on MiniCPM-o 4.5 and then decays to **0.27** at the output, with the same decay on MiniCPM-o 2.6 (0.81→0.52) and *no* decay on either raw backbone (Qwen3-8B 0.86, Qwen2.5-7B 0.71 at their final layers) or on the omni-streaming control (Qwen2.5-Omni holds 0.75 to its output with a flat variance share) — the decay tracks duplex fine-tuning specifically, not streaming or audio fine-tuning in general. On MiniCPM-o 4.5 the score also concentrates: the dominant subspace carries 39% of the held-out score variance at the final layer against ≤ 0.18 at every earlier depth and 0.05 for its backbone; on MiniCPM-o 2.6 this concentration is present but weak (0.16 vs. 0.14 for its backbone), so we report it as a property of the stronger duplex model, not a general law.

Direct turn-control tests. Two experiments target the turn-control story head-on rather than by exclusion (`modal_interp.py`, `scripts/19_*.py`). *Control-token logit lens*: applying each model’s own final norm and unembedding to the stored per-layer states, the probability mass on the trained-in control vocabulary (added/special tokens) *rises* over MiniCPM-o 4.5’s last layers (log-mass −23.8 at L32 to −11.8 at the output) while the raw backbone’s falls by fifteen log-units over the same span — the duplex output layer, uniquely, stays ready to emit control tokens. But the fine-grained prediction fails: the inverting probe direction is no more aligned with the control tokens’ unembedding rows than with random vocabulary rows ($\max |\cos|$ 0.069 vs. 0.068), and per-query control mass correlates with the probe score at only $r=0.18$. *Listen/speak intervention*: re-prefilling 60 calibration queries with a still-listening vs. answer-now suffix moves the final-layer state substantially in *both* models (relative displacement 0.27 duplex, 0.59 raw — cue text is just prompt semantics), and the cue direction is orthogonal to the inverting probe ($|\cos|=0.019$, chance 0.016), shifting its score by only +0.23 sd (and the deployed L22 read by +0.08 sd). Together: the duplex output layer is measurably control-biased in its output distribution, but the inversion is not carried by control-token directions or by an explicit listen/speak axis — the simple takeover story fails its own direct tests.

What we do not claim. Destructive tests do not settle the mechanism. Removing the top five principal directions from the final layer does raise LOPO math from 0.378 to 0.758 — but the identical operation costs the raw backbone 0.90→0.14 and costs the intact L22 read 0.931→0.760, so it is a damaging intervention that happens to help a broken readout, not evidence of a duplex-added subspace. We therefore state the finding at the level the evidence supports: duplex fine-tuning does not erase the competence signal, it stops delivering it in distributed form to the final last-token position, and the layer where delivery fails is also the one whose output distribution stays control-biased. Attributing the loss to a specific duplex objective needs training-time access we do not have; the practical consequence — read mid-network — does not depend on the attribution.




Figure 6: Non-destructive decomposition of the probe’s held-out score at every depth. **Left:** share of score variance carried by the layer’s five dominant directions. **Right:** transfer AUC of the residual (distributed) component — intact to the output on both raw backbones and on the omni-streaming control, decaying to inversion over the final layers of both duplex models.

636 E Query-surface router audit

637 A RouteLLM-style TF-IDF→LR router under identical labels reaches OOF AUC 0.669 (below the
 638 pool oracle’s 0.715); a 24-configuration sweep tops out at 0.689, so this is the query surface’s ceiling,
 639 not tuning. Under LOPO it collapses to 0.38–0.57 and scores the 100%-fail trap pool at 0.230 —
 640 it misses the pool a router most needs. The training receipt: at the argmax threshold its out-of-fold
 641 accuracy exactly matches the majority class; the internal probes under the same accounting clear
 642 majority by 9–29 points (Figure 8). Nor is it data-starved: the identical recipe on RouterBench [Hu
 643 et al., 2024] (36,497 prompts) reaches only in-domain AUC 0.710 with leave-one-benchmark-out
 644 at chance; RouteLLM’s released preference-trained routers score 0.523/0.533 zero-shot on our pool
 645 and rank the trap pool below ordinary hard pools; our router transfers to RouterBench below chance
 646 (0.440). On audio labels the router improves (test 0.814) but still trails the mid-layer probe (0.879)
 647 and cannot fire mid-stream. In distribution, the gate’s area over random (+0.054) is statistically
 648 on par with the pool oracle’s (+0.042; paired bootstrap delta CI $[-0.003, +0.027]$, $p=0.13$) — the
 649 gate’s case is transfer, not the in-distribution margin. $p(\text{True})$ tradeoff variants: Figure 9.




Figure 7: Accuracy vs. escalation rate, frozen test split: gated escalation vs. random-escalation and pool-oracle (type-only) baselines between the small-only and big-only endpoints.




Figure 8: Identical training receipt for the external router and the two internal probes: the router never clears the majority-class baseline; both probes do, by 9–29 points.

650 F Gate timing details

651 **Against the model’s own thinking mode.** MiniCPM-o 4.5 ships a built-in reasoning mode —
 652 self-escalation in place. Gated-cloud escalation dominates all-THINK on both accuracy and latency




Figure 9: Offline tradeoff for the $p(\text{True})$ signals (pre-answer and post-draft).

at every escalation rate (gated-cloud@33%: 78.7% at mean 6.5 s vs. all-THINK 63.7% at 22.4 s; Table 9). The reason is mechanistic: the gate flags knowledge and trap failures, which extra compute cannot fix — so external escalation is necessary, not merely better.

Table 9: Deployment policies, frozen test split ($n=240$); latencies measured.

policy	acc (%)	lat. mean (s)	P50	P95
fast-only	58.8	3.0	3.5	4.2
all-THINK (built-in)	63.7	22.4	17.2	60.1
gated-think @33%	61.3	12.1	3.6	47.6
gated-cloud @15%	68.8	5.3	3.6	10.1
gated-cloud @33%	78.7	6.5	3.6	20.8
gated-cloud @50%	85.8	7.0	3.6	24.4

Measurement protocol. Gate latency is wall-clock from prefill start to the L22 score being available (truncated forward + probe dot product), CUDA-synchronized, one H100, 50 queries per modality, three warmup calls excluded, interleaved per query with the reference arms; P50/P95: gate 20/25 ms text and 45/104 ms audio vs. full-prefill TTFT 36/47 and 68/144 ms. Network and expert-RPC initiation are excluded: the cloud call fires asynchronously at the gate, and its user-facing cost is reported as expert-content-audible latency in Appendix G. Absolute prefill wall-times vary with call overhead across benches; the robust cross-bench statistic is the ratio — the L22 state is ready at ~ 57 – 61% of prefill.

Per-layer truncated-forward cost vs. quality (Figure 10): the L22 hidden state is available at ~ 57 – 61% of prefill wall-time while L22 is simultaneously the in-mix quality peak; forking much earlier is both weaker and modality-specific, so the fork belongs at ~ 55 – 60% depth. If the cloud call launches at the gate, the expert result arrives before the local answer finishes for 40–58% of the whole mix but only ~ 20 – 31% of *escalated* queries (they skew to short local answers); the residual wait is P50 2–3 s, P90 ~ 27 s (Figure 11) — the stall-phrase budget the injection protocol must cover.

G Live-loop details

What the signal detects. The pre-answer probe is specifically a *retrieval-failure* detector, strongest on missing-knowledge and long-tail-trap failures where the empty lookup is observable in the model’s state. Execution failures (stuck mid-derivation) surface in *decode-time* signals instead, and metacognitive failures — false-premise questions, where no failure event occurs at all —




Figure 10: Fork-depth Pareto: per-layer wall-time vs. gate quality.




Figure 11: Cloud round-trip vs. local answer duration: probability the expert result arrives before the local answer ends.

675 are invisible to every signal we tested, query-surface routers included (fire rate 18% vs. trap’s 90%;
676 Appendix G), which motivates a decode-time check behind the pre-answer gate.

677 **Full live tables.** Table 10 (both scoring views, speakable subset and full pool) and Table 11 (measured response latency). The conservative tier *shortens* the P95 tail: its few escalations displace
678 exactly the longest local decodes. Figure 13 joins accuracy and latency; Figure 14 replays three real
679 sessions from recorded timers.
680

681 **Time to first audible audio.** The sweep above is text-out; a TTS-on re-run of the balanced arm
682 (same 240 queries, talker head + vocoder on, stall pre-synthesized once in the talker’s own voice)
683 measures speech onset directly. The gate decision is unchanged — the TTS token enters the context
684 only after the end-of-turn read (read p50 21 ms / p99 32 ms), and the same 84/240 queries fire. From
685 gate decision to first audible audio: **local answers p50 0.552 s / p95 0.599 / p99 0.644; escalated**
686 **turns p50 0.022 s / p99 0.029** — the canned stall (3.2–4.1 s of speech) plays the moment the gate
687 fires, so the escalated path reaches first audio 25× faster. Expert *content* becomes audible at p50
688 6.4 s / p95 26.0 / p99 34.0 (cache-corrected expert round-trip + first relay PCM, p50 0.66 s after the
689 answer returns); per-turn dead air between stall end and expert content is p50 2.5 s / p95 22.4 / p99
690 30.7, with 25/84 escalated turns fully covered — the stall covers the head of that wait, and Table 11’s
691 completion profile bounds the tail.

692 **Speculative prefetch, measured and dropped.** Framed as speculative escalation (draft–verify at
693 sentence granularity), its acceptance rate — does an expert answer to the partial transcript answer
694 the full question — is 0.07/0.19/0.51 at 25/50/75% of the audio; at the mid-stream gate’s typical fire




Figure 12: Live dual-view tradeoff ($n=218$ speakable subset). Blue: honest heard accuracy of the live system (paired-bootstrap CIs). Green: the channel-controlled counterfactual — same sessions, escalated rows re-scored with gold-text expert answers. Each view is plotted against *its own* matched-rate random reference (faint dashed, same floor, each view’s own 100% endpoint): selection beats random within both views at every arm, by 0.072–0.095 heard and 0.054–0.084 gold. The vertical view-to-view gap is the speech-channel price; the floor-to-floor gap to offline operation is the audio + streaming-answer tax.

Table 10: Live streaming sweep, paired bootstrap 95% CIs ($B=10^4$). Top: speakable subset ($n=218$); bottom: full pool ($n=240$). Gold-inject = escalated rows re-scored with gold-text expert answers; channel cost = paired difference. * CI excludes 0; † CI lower bound at zero.

tier	esc	heard-acc	Δ vs. floor	gold-inject	channel cost
never	0%	0.422 [0.36,0.49]	—	0.422	—
conservative	14%	0.528 [0.46,0.59]	+0.106*	0.550	+0.023†
balanced	31%	0.592 [0.53,0.66]	+0.170*	0.661	+0.069*
aggressive	57%	0.633 [0.57,0.70]	+0.211*	0.761	+0.128*
always (synth)	100%	—	—	0.922 [0.89,0.95]	—
never (full)	0%	0.383 [0.32,0.45]	—	0.383	—
conservative	16%	0.483 [0.42,0.55]	+0.100*	0.525	+0.042*
balanced	35%	0.554 [0.49,0.62]	+0.171*	0.654	+0.100*
aggressive	61%	0.596 [0.53,0.66]	+0.212*	0.771	+0.175*
always (synth)	100%	—	—	0.917 [0.88,0.95]	—

point 6–8 of 10 calls would be wasted. The trap pool’s floor (0.00/0.10/0.27) shares the back-loaded-core property that breaks mid-stream probe reads.

Conflicting-injection controls (172 sessions). Fabricated conflicting expert answers by within-pool derangement, four framings: worst-case comply is 0.00 neutral / 0.53 deployed authority / 0.79 strong-override, while *seeding* words into the talker’s mouth backfires (comply 0.25, lip-service 0.62 — forced speech is treated as revisable). Math is silently re-derived rather than relayed or disputed. Model-wrong $\times$ wrong-injection swallow rate is 0.64: expert quality is the accuracy ceiling of the cascade.

Uplink diagnostics. On the escalated population the expert answering the talker’s own transcript scores 0.585 vs. 0.871 on gold text. Prompt-level rescues on the self-transcript are no-ops (the damaging errors are internally-coherent entity substitutions); replacing the ears helps: open-weights Whisper recovers little (+4 pp) but a stronger cloud ASR on the same audio recovers 38% of the gap (0.585→0.694, McNemar $p=0.007$) — diagnostic only (its critical-path latency and external-pool survival are unmeasured). A same-model gold-text control with an audio-capable expert shows the

Table 11: Measured response latency (s), $n=240$ per arm (latency is channel-independent). P99 on $n=240$ is indicative only.

tier	mean	P50	P95	P99	Δ mean	Δ P50	Δ P95	Δ P99
never	4.6	2.0	15.8	17.8	—	—	—	—
conservative	4.6	2.6	12.8	23.0	+0.0	+0.6	−3.0	+5.2
balanced	5.8	3.5	15.9	32.7	+1.2	+1.5	+0.1	+14.9
aggressive	6.6	4.4	18.5	33.0	+2.0	+2.4	+2.7	+15.2




Figure 13: Accuracy vs. per-arm P50 latency (speakable subset), both scoring views, with the simulated-P50 random reference.

709 audio channel itself is lossless for speakable content: the leak is the talker’s transcription, not the
710 synthesis or the relay (relay drops: 5 cases of 108; corrupted transcripts dominate the escalated
711 heard-failures, first-generation sweep).

712 **Boundary query classes.** *False premises:* the model fails substantially (0.63 text / 0.47 audio)
713 and the expert would help (0.80), but the gate is blind (fire@balanced 0.18 vs. trap’s 0.90; no score
714 separation) — answering along a false premise does not feel like inability. *Real-time queries* (“what
715 is NVDA trading at”): the gate is not blind, the class was simply absent from every static training
716 family; extending calibration with FreshQA [Vu et al., 2023] (fast-changing labeled escalate a priori,
717 never-changing as in-family controls) raises held-out fast-changing fire rates to 0.80/0.93/1.0 across
718 tiers with frozen-test AUC unchanged (0.877 vs. 0.879). One calibration subtlety: tier budgets must
719 remain quantiles of the deployment-mix scores, or the new capability is spent on its own threshold
720 shift.

721 **Demonstration system.** The pipeline runs as an interactive web demo on one H100: live micro-
722 phone audio, end-of-turn read (~ 54 ms) against the balanced threshold, escalated turns stall, consult
723 the expert, and *speak* the relayed answer through the model’s native TTS head (probe context and
724 thresholds untouched). Two demo-only extensions: a web-search tool for the expert (real-time
725 queries then return live values), and backchannel-aware barge-in — speech over the talker ducks
726 playback, a fast ASR pass returns the floor on backchannels (“ok”, “mm-hm”; an English+Chinese
727 lexicon) or commits the interrupt, aborting generation at chunk granularity and seeding the next turn.
728 All measured arms keep the tool-free expert and turn-based loop; the demo will be made available
729 upon publication.

Three live sessions, timestamped (balanced arm; every event time is a recorded live-loop timer)




Figure 14: Three real sessions from recorded timers: expert beats the stall phrase (seamless); expert outlives it (dead air, the P99 signature); gate stays closed (local answer fully decoded 0.4 s after end of turn). Speech occupancy is estimated at 150 wpm (hatched); the measured arms are text-out, so no vocoder onset is plotted.

H Duplex-validation sweep: the gate under the talker

The headless live sweep (§6) never engages the talker head; this sweep replays its full 240-query pool through the deployed voice loop above to test whether the frozen probe read survives the spoken regime. Two arms, one session per query, deployed v4 gate with scores logged but escalation disabled (the arm isolates gate validity, not routing outcomes). **Clean:** the query streams as mic-shaped PCM (pool TTS audio plus steady noise at RMS 0.008), the energy VAD ends the turn, the talker answers aloud, and the session is cut after the first audible chunk. **Overlap:** an easy warmup question (probe score <0.25 , locally answered, >5 s spoken answer) puts the talker mid-answer, and the target query is then spoken over it; sustained overlapping speech commits a barge-in whose audio seeds the next turn. The barge-in path engaged on 237/239 pairs.

Against the frozen local-failure labels of the same pool, the identical frozen probe reads: headless AUC 0.866 [0.817, 0.909], clean 0.871 [0.824, 0.914], overlap 0.856 [0.803, 0.901]; barged-only subset ($n=237$) 0.854. Paired bootstrap deltas vs. headless are $+0.005$ [-0.025 , $+0.035$] (clean) and -0.009 [-0.036 , $+0.016$] (overlap) — both inside the ± 0.02 – 0.03 live replication floor. Per-query scores correlate across regimes at $r=0.90$ – 0.92 , and the balanced threshold makes the same fire/hold decision as the headless read on 0.887 (clean) and 0.912 (overlap) of queries. The end-of-turn read costs 24–26 ms at p50 ($p95 \leq 30$ ms) in both arms, with the first audible answer chunk arriving ≥ 0.5 s after the gate decision ($p50 \approx 0.58$ s), so the decision precedes any audible commitment. Scope: this loop keeps per-turn session control — the mic stays open while the talker speaks and interruptions seed the next read, but incoming audio is not prefilled *during* generation. The concurrent regime is measured separately below.

Concurrent prefill: interleaved listen/speak in one KV stream. A third arm drives the streaming API in time division: while the talker is mid-answer, each 1 s chunk of the target query is prefilled into the *same* session between generation chunks, so the end-of-turn read fires while generation is still active (238/240 queries, full interleave coverage). The carrier answer is teacher-forced: under sampling the duplex model reliably yields the turn — it emits EOS within ~ 2 s of user audio appearing in its context, the trained turn-taking behavior surfacing under interleave. This is the one shift on the transfer ladder that breaks the frozen read: scores drift up (mean 0.48 $\rightarrow$ 0.71; the frozen balanced threshold’s fire rate jumps 37 $\rightarrow$ 75%), AUC drops 0.869 $\rightarrow$ 0.758 [0.693, 0.816] (paired Δ AUC -0.109 [-0.158 , -0.064]), and label-free requantiling recovers decision agreement only to 75%. The information, however, survives in the state: refitting the *same* 12288-d linear read on 360 calibration queries collected in-regime reaches **0.818** [0.761, 0.871] on the held-out 240 (calibration OOF 0.775) — more than half the gap recovered at a quarter of the deployed probe’s calibration scale. The mid-layer competence direction persists under concurrent listen/speak; unlike every earlier shift (modality, external pools, talker-on), calibration must follow the regime. Run-to-run replication of the concurrent read correlates at 0.895. Caveat: the interleave goes through the public streaming API — turn-role headers enter mid-generation and the carrier text is forced — so this measures the concurrent *state*, not the official duplex serving format.

The full loop in the concurrent regime. With the in-regime gate (thresholds = label-free quantiles of each pool’s own concurrent never-arm scores), the complete escalation loop — carrier speaking, target audio interleaved, EOT read mid-generation, stall $\rightarrow$ expert $\rightarrow$ relay — was re-run end to end on all seven pools (Table 12). The mechanism reproduces: heard accuracy rises monotonically in budget on every pool, realized rates track nominal (10–51%), and *selective escalation beats always-escalate* on the internal pool (66.2 vs. 61.7%; speakable subset 71.1 vs. 66.5%) and on Llama Questions (86.4 vs. 85.6%, our judge) — the turn-based signature result carries over. On the speakable subset the concurrent loop reaches **44.5 $\rightarrow$ 71.1%** at 63% escalation (turn-based reference 42.2 $\rightarrow$ 63.3% at 57%). Floors move by pool — the regime cost story: Llama $\approx$ unchanged, TriviaQA -16 , Reasoning zh -18 (the English carrier hurts the zh pool most), WebQ $+4.8$ under the official judge (within the ± 2 –3 replication floor; the larger our-judge shift is alias-matching style sensitivity). AlpacaEval reverses its turn-based honest-negative: 4.36 $\rightarrow$ 4.60 (1–5 scale, always 4.76). Two honest notes: the in-regime probe, calibrated on 360 internal rows only, transfers externally at AUC .57–.67 against the full-scale turn-based probe’s .68–.81 — a calibration-*scale* gap, not a signal gap (internally, where calibration matches the distribution, it reads .857 vs. the turn-based .879) — and the internal conservative threshold under-fired (2.5% realized vs. 15.8 nominal; the top calib quantile did not transfer, balanced and aggressive did). Latency holds: EOT read p50 23 ms, canned-stall onset 27–34 ms, local first audio p50 .655 s (vs. .552 turn-based — the price of the longer concurrent context).

Table 12: The full escalation loop in the concurrent regime, all pools, judge-matched to Table 1 (OAB official for TriviaQA/WebQ/Llama; ours for SD-QA/Reasoning/internal; VoiceBench 1–5 for AlpacaEval). Gate = in-regime probe (360-row calibration), per-pool label-free quantile thresholds. Accuracies in %; AUC = in-regime probe vs. never-arm failure, decimals.

pool	never	@15%	@30%	@50%	always	AUC
TriviaQA	50.4	54.0	64.0	71.6	91.6	0.636
WebQ	61.6	65.2	68.8	71.2	81.6	0.636
Llama Q.	81.6	83.2	84.8	90.8	91.2	0.672
SD-QA	49.5	55.5	60.5	70.5	89.5	0.636
Reason. zh	40.6	48.5	58.4	64.4	80.7	0.570
internal (240)	40.4	40.0	52.1	66.2	61.7	—
internal speakable (218)	44.5	43.6	56.0	71.1	66.5	0.857
AlpacaEval (1–5)	4.36	4.45	4.53	4.60	4.76	—

A serendipitous recovery behavior. An accidental run surfaced one bonus behavior worth reporting: when a long question contains a >1.25 s internal pause, the endpointer fires early and the talker

starts answering the fragment — but the question’s own continuation then triggers barge-in, and the remainder becomes the next turn, where the gate simply re-scores it. Early end-pointing thus degrades into an extra turn rather than a lost query — the honest duplex-ish recovery available to a turn-based loop.

I Fixed-threshold transfer and the expert-benefit oracle

Both analyses below run in the *gold-inject* view (escalated rows take the expert’s answer to the gold text). The deployed heard view cannot evaluate a freely moving threshold: the most permissive arm escalates 50%, so the other half of each pool has no expert outcome on record, whereas the ceiling runs cover every id. Labels follow each pool’s own judge (Table 1); the two judges disagree by up to 13 points on Web Questions, so they are never mixed within a row. AUC columns are recomputed from the deployed sweep’s stored end-of-turn scores; they differ from Table 1’s AUC column — captured in a separate scoring pass — by up to 0.03 externally and 0.04 on the internal pool (0.840 here vs. 0.879 there): the passes render and prefill the audio independently, and the probe score carries render-to-render jitter. AlpacaEval is excluded — it has no binary correctness label on either side, only a 1–5 score whose expert ceiling is degenerate (195/199 at 5).

One absolute threshold, six pools. The deployed gate adapts per pool, but only through label-free quantiles of its own score distribution. To separate what the probe’s *ranking* carries from what its *scale* carries, we freeze the single absolute threshold the internal system actually deployed ($\tau=0.628$, the balanced tier) and apply that number verbatim everywhere. **Scale does not transfer:** the same τ realizes escalation rates from 2.0% (Reasoning QA) to 48.0% (SD-QA), a 24-fold spread, so a fixed cut cannot hold an escalation budget across pools. **Ranking does:** at whatever rate it fires, fixed- τ escalation beats a random reference at the same realized rate on four of five external pools, paired-bootstrap interval excluding zero (Table 13). The exception is Reasoning QA, where τ fires on 2% of queries and there is almost nothing to measure. This is the concrete case for the label-free quantile step: it buys budget control, not discrimination.

Does the gate find failures the expert can fix? The probe is fit and read against *local failure*, but what a routing system needs is *expert benefit*: $y=1$ when the local model is wrong *and* the expert is right. Re-scoring the same frozen scores against this harder label costs little on the retrieval pools (Speech TriviaQA 0.796→0.782, Llama Questions 0.830→0.813) and more where the expert is itself weak (Web Questions 0.773→0.689 against a 0.864 expert ceiling; internal 0.840→0.732). The gate thus remains a usable benefit predictor out of distribution (0.63–0.81) despite never having seen an expert outcome, and the drop is largest exactly where the paper already reports a flat live curve. On Web Questions the mechanism is direct: 75.0% of local failures are expert-fixable pool-wide, but only 65.5% of the failures inside the gate’s top-scoring 15% are, and that share climbs monotonically as the cut widens (0.655/0.685/0.711 at 15/30/50%) — the gate’s most confident failure calls are the ones the expert also misses. Table 14 places the deployed gate against the matched-budget oracle that escalates true benefit cases first: the gate clears random at every pool and budget, capturing a fifth to three fifths of the random-to-oracle span, and the residual gap is the headroom a decode-time or expert-aware second stage would have to claim.

J Transfer latency shapes

Matched-rate precision. Precision at a fixed escalation rate is capped at base-fail/rate: on Llama Questions (base fail 0.16) the 50% cap is 0.32 and the deployed probe reaches 0.304 — 95% of the cap (recall 0.93). Quoting raw precision would misread the strongest pool as the weakest.

The median left-fold. On easy pools per-arm P50 is non-monotonic (Llama Questions 1.52→1.17→1.60/1.42 s): a verified median-of-a-mixture mechanism, not noise. The probe preferentially escalates the queries whose *local* decode is longest — on retrieval pools wrong answers

Table 13: One absolute threshold ($\tau=0.628$, frozen on the internal pool) applied verbatim to every pool; gold-inject view. Δ is fixed- τ accuracy minus a random reference at the *same realized rate* (paired bootstrap over ids, $B=10^4$, seed 42). AUC columns re-score the same frozen scores against local failure and against expert benefit ($y = \text{local wrong} \wedge \text{expert right}$, base rate in the last column). Rates and accuracies in %; AUC in decimals.

pool	n	rate	local	acc	rand	Δ 95% CI	AUC _{fail}	AUC _{benefit}	base
Speech TriviaQA	250	28.4	66.4	83.2	75.0	[+4.4, +12.1]	0.796	0.782	30.8
Speech Web Questions	250	36.0	56.8	72.0	67.5	[+0.8, +8.4]	0.773	0.689	32.4
Llama Questions	250	8.0	83.6	86.8	84.3	[+0.6, +4.6]	0.830	0.813	11.2
SD-QA	200	48.0	51.5	80.5	71.4	[+4.5, +13.6]	0.800	0.765	43.5
Reasoning QA (zh)	202	2.0	58.9	60.9	59.5	[−0.1, +3.5]	0.679	0.634	30.2
internal frozen test	240	35.0	38.3	65.4	57.0	[+3.8, +13.1]	0.840	0.732	53.3

Table 14: Oracle frontier at matched escalation budgets (gold-inject view): random reference / deployed quantile gate / benefit oracle, which escalates the true $y=1$ cases first. The gate clears random everywhere; the oracle column is the remaining headroom. Accuracies in %.

pool	15% budget			30% budget			50% budget		
	rand	gate	oracle	rand	gate	oracle	rand	gate	oracle
Speech TriviaQA	71.0	76.0	81.6	75.5	83.6	96.4	81.6	91.2	97.2
Speech Web Questions	61.2	64.0	72.0	65.7	70.0	86.8	71.6	76.4	89.2
Llama Questions	85.0	88.8	94.8	86.4	91.2	94.8	88.2	93.2	94.8
SD-QA	57.7	63.5	66.5	63.9	74.5	81.5	72.2	82.0	95.0
Reasoning QA (zh)	63.1	67.8	73.8	67.4	72.8	89.1	73.0	76.7	89.1
internal frozen test	46.3	49.2	53.3	54.3	60.8	68.3	65.0	74.2	88.3

835 hedge at length (decode time and answer length $r=0.97$), so a wrongness-selecting probe strips the
836 slow tail as a side effect (escalated-at-15% ids: never-arm local P50 1.711 s vs. 1.198 s kept; local
837 accuracy 0.32 vs. 0.73). Controls: the probe does not key on length ($\text{corr}(\text{score}, \text{length}) = +0.05$;
838 ≈ 0 conditioned on wrongness), and on SD-QA — where wrong answers are short — the same probe
839 selects wrongness identically and produces no fold. The mean is monotonic (1.65→2.06 s); kept-
840 local P50 falls 1.52→0.43 s across tiers (Figure 15). On Reasoning QA the shape inverts usefully:
841 the local chain-of-thought is *spoken*, so every reasoning token enters the response clock, while the
842 expert’s chain is hidden behind a flat ~ 3.7 s round-trip — escalation truncates the latency tail (P90
843 13.25→10.94 s) rather than fattening it.




Figure 15: Mixture decomposition of Llama Questions latency per arm: kept-local P50 falls 1.52→0.43 s as the probe strips the slowest local decodes while escalated rows pay a flat ~ 3 s expert round-trip; the arm median (dashed) is their mixture.

AlpacaEval scale caveats. The official 4.8 is an offline chat-mode number — our chat-mode control reproduces it at 4.86 — and the chat-vs-live gap (4.86 vs. 3.99) is the duplex loop’s spoken-reply style plus a token cap, so only live arms are comparable to each other. Selected queries score 3.90 on the never arm vs. 3.98 for skipped: no discrimination; escalated rows gain (3.90→4.71) because the expert writes better long-form answers — random picks would harvest the same gain.

K Full-pool and per-family figures

Figures 16 (dual-view) on the full frozen pool, and the complete external-family figure set (Appendix L); the noise audit behind the ± 0.02 – 0.03 floor is Figure 17.




Figure 16: Dual-view live tradeoff, full pool ($n=240$); companion to Figure 12.




Figure 17: Replication-noise audit: same query, same audio, re-run across arms. Kept-local verdicts flip 2.3–18.8% by pool; repeated escalation 0.7–10.6%; implied paired SE 0.009–0.028.

L External-benchmark figure set









Figure 18: Speech TriviaQA, Speech Web Questions, SD-QA: dual-view (left) and accuracy–latency (right).




Figure 19: Llama Questions, Reasoning QA (zh), VoiceBench AlpacaEval: dual-view (left) and accuracy-latency (right).

853 **M Judge and self-evaluation prompts**

854 All adequacy labels use the following fixed system prompt (structured outputs with a strict {reason,
855 adequate} schema; the judge never sees which system produced the answer):

```
856     You are a strict, fair grader of a small language model's answers.  
857     You are given a QUESTION, an optional REFERENCE answer, and a  
858     CANDIDATE answer produced by a small model. Decide whether the  
859     CANDIDATE is ADEQUATE, meaning ALL of: (1) factually correct;  
860     (2) any reasoning shown is valid; (3) it actually and completely  
861     answers the question that was asked. For multiple-choice questions  
862     the candidate must land on the correct option (by letter or by  
863     unambiguously stating the option's content). A partially correct,  
864     hedged, refused, off-topic, or truncated answer is NOT adequate. If  
865     a REFERENCE is provided, grade against it. Be strict but do not  
866     penalize harmless differences in format or extra correct detail.
```

867 The verbalized self-evaluation baselines use one fixed phrasing each; $P(\text{Yes})$ is read from the first
868 generated token's distribution:

```
869     [pre] I am going to show you a question. Do NOT answer it. Judge  
870     honestly whether you yourself would answer it correctly. Question:  
871     {q} Would you answer this question correctly? Reply with exactly one  
872     word: Yes or No.  
873     [post] Here is a question and a proposed answer. Question: {q}  
874     Proposed answer: {a} Is the proposed answer correct? Reply with  
875     exactly one word: Yes or No.
```

876 **N Reproducibility**

877 All sampling and splits use a fixed seed (42); the calibration/test split is frozen once and the test split
878 is touched only by the evaluation scripts. Environment pins: torch 2.8.0; transformers 4.51.0 for
879 MiniCPM-o 4.5, with per-model pins for comparison backbones; all experiments run on single-H100
880 instances. Probe artifacts (weights, feature configuration, per-pool label-free thresholds), the judge
881 prompt above, all trace files, and the per-experiment execution log — including per-run GPU and
882 API spend, on the order of \$500 total — ship in the project repository, released upon publication.

883 **TODO (process only — delete before submission)**

- 884 • [TODO: P1 — DONE 2026-08-25 (external review applied): main text now ends
885 within 8 pages (refs start mid-page 8). Restructure: 3 contributions + hero figure;
886 merged signal section; live 6pp → 1.2pp; router audit 3 lines; title dropped “Zero-
887 Training”; causal claims softened (“consistent with”); no version numbers / artifact
888 filenames in main text; everything cut lives in the appendix (A–L) or git history]
- 889 • [TODO: P2 — peer + meeting feedback (GPT reviewer pass 2026-08-26, borderline
890 4/7). (a) and (b) DONE 2026-08-26; (c) dropped as not computable; (d) unchanged.
891 (a) DONE — three claims narrowed: “routing beats scaling the speech backbone”
892 → “selective cloud escalation outperforms larger standalone backbones”; “standard
893 readouts fail” → “late-layer readouts turn brittle after duplex fine-tuning” (abstract,
894 intro contribution 1, hero + teaser panel-1 titles); and the EOT loop is no longer
895 equated with native full-duplex — intro now carries an explicit scope sentence (“*du-*
896 *plex* names the training regime, not the interaction protocol”), live.tex states the turn-
897 based protocol up front and is retitled “Live End-of-Turn Escalation”, and “duplex
898 session/system” → “spoken session”/“live system” throughout. (b) DONE — new ap-
899 pendix section (app:fixedthr) + two tables: one absolute threshold ($\tau=.628$, frozen
900 internally) applied verbatim to all pools fires at 2–48% yet still beats matched-rate
901 random on 4/5 (ranking transfers, scale does not); expert-benefit oracle (y = local-
902 wrong AND expert-right) with AUC re-scored per pool and a gate/random/oracle
903 frontier at 15/30/50% budgets. Script: scripts/09_fixed_threshold_oracle.py → fig-
904 ures/fixed_threshold_oracle.json. Gold-inject view (heard view only observes the
905 ~50% of ids some arm escalated); labels follow each pool’s own judge (oab_ok vs
906 adequate — they differ by 13pp on Web Questions, never mix); AlpacaEval excluded,
907 it has no binary label on either side. (c) DONE 2026-08-26 (was DROPPED, then
908 re-run with TTS on): bench_live gained a tts flag (init_tts + Token2wav, canned stall,
909 gen_speak port of demo_app); balanced arm re-run on the frozen 240 (~50 min H100,
910 expert cache hits). EOT→first-audible: local p50 .552s/p99 .644s, escalated (canned
911 stall onset) p50 .022s/p99 .029s, expert-content p50 6.4s/p95 26.0/p99 34.0. Same
912 84/240 fire as text-out; probe read unperturbed (p50 21ms). Landed: live.tex sentence
913 + app:livedetail paragraph. Traces data/frozen_v3tts_traces.jsonl.balanced.shard0-3;
914 analysis scripts/10_tts_latency.py. Fixed en route: fig:timeline-live panel (c) was cap-
915 tioned “first token 0.4s” but the plotted value is answer_ms = *full* answer decode;
916 caption + figure title corrected and the figure regenerated. (d) DONE 2026-08-27
917 (was SKIP). 8as closed the last tier: concurrent prefill — target audio interleaved into
918 the SAME KV stream between generation chunks, EOT read fires while generation
919 is active (238/240). Frozen v3 probe degrades AUC 0.869→0.758 (paired Δ −0.109
920 [−0.158, −0.064]) — the ONE shift on the ladder that breaks frozen transfer; label-
921 free requantiling recovers decisions only to 75%. In-regime refit of the same 12288-d
922 read on 360 concurrent calib rows recovers 0.818 [0.761,0.871] at quarter-scale data
923 ⇒ signal intact, calibration must follow the regime. Landed: app:duplexval third
924 arm, Limitations (vi), intro clause. Harness modal_duplex_concurrent.py; scripts
925 11/12; RESULTS 8as. Page rebalance done same pass: Table 2 policy table, dualview
926 figure, second-family block and the serendipity paragraph moved to the appendix;
927 main text ends p8, refs p9. Scale unified (accuracies/rates in %, AUC in decimals).
928 Still open: official duplex serving format, full-scale (2310) in-regime recalibration,
929 human continuity study, second expert. (d-earlier) 8aq duplex-validation sweep ran
930 — the frozen-pool 240 replayed through the DEPLOYED voice loop (talker on, spo-
931 ken answers) in clean + overlap/barge-in arms via _ws_duplex.py against the live
932 demo; frozen v4 probe, no refit: AUC .871/.856 vs .866 headless (paired deltas inside
933 the replication floor), decision agreement .89/.91, barged 237/239. Landed: live.tex
934 closing paragraph, app:duplexval, intro scope sentence updated, Limitations (vi) nar-
935 rowed to “concurrent-with-generation read remains open”. Data: gate-data volume

duplex_sweep/{clean,overlap}.jsonl; analysis_duplex_analyze.py. This also partially recovers (c): the voice loop logs EOT→first-audible per turn (first_audio_client_ms, p50 ≈.58s, both arms) — usable if a reviewer asks, though only under probe-off replay, not the escalating sweep. Still skipped: human continuity study, second expert, model-level concurrent listen/speak read. Frame the paper as what it is: pre-response competence probing in duplex-trained models + turn-based live escalation, now with a deployed-loop validity check.]

- [TODO: P2.2 — meeting feedback round 2 (2026-08-26), four items, all landed same day: (1) ugly zigzag figures: fork_pareto redrawn as depth-colored scatter + Pareto frontier (figures/fork_pareto.py, new); sllama_latency_decomp cut to the clean single decomposition panel (median-vs-mean fold fact stays in prose). (2) main table: tab:transfer upgraded Kitty-style — live tiers + ceiling + AUC + official PLUS gold-inject@30% rand/gate/oracle block (numbers from tab:oracle); judge folded into row names. Views kept separate and labeled — do NOT let anyone compare live heard vs gold-inject across blocks. (3) special-token routing discussed in related work: pause/thought tokens (goyal2024pause, zelikman2024quietstar), fused think modes (yang2025qwen3, guo2025deepseekr1), MiniCPM-o’s own think gate — they decide more compute not which model, need talker fine-tuning, policy baked per checkpoint; ours = frozen + pluggable + beats the built-in gate at matched budget (tab:policy). (4) simple-but-important framing: intro minimalism sentence (“competing signals read the wrong place / arrive too late / must be trained in”). Page budget paid by: fair_dualview .56→.45 linewidth, caption + related + gap compressions.]
- [TODO: P2.3 — storytelling pass (advisor feedback 2026-08-26, selectively adopted): narrative arc sharpened to “looks solved → the obvious readout is exactly what duplex training corrupts → signal sits one band earlier → read it before speech commits”. Abstract rewritten to that arc; intro adds the question line (“not whether one can build a router — whether the model knows it will fail early enough to do anything about it”) and the find-it/read-it-in-time/show-it-matters framing; “before the first token” upgraded to “half a second before the first audible word” (grounded in the TTS-on p50 .552s + gate ~30ms); conclusion closer = “the signal was never gone — only misread”. NOT adopted: restructuring the three contributions, cutting appendix richness. Space paid by trimming the extended-analyses name-drop line.]
- [TODO: P0-R — reviewer P0 pass (2026-08-27), five questions; answers exist in-repo, story needs surfacing: (1) why route pre-first-token: NOT a safety necessity (first token → first audible ≈0.5s exists) — it’s that waiting buys no AUC (measured: entropy@4 0.696, decode-hidden mean 0.776 vs prefill read 0.828/0.887) and every ms of delay is dead air on the escalated path (expert RTT = 85–88% of escalated wall time). Edit: one sentence in gate.tex reframing pre-token as free-and-optimal, citing the measured decode-time rows. (2) final-layer mean-pool WAS run: LOPO math o4.5 ≈0.80 (L35) / o2.6 0.674 vs mid last-token 0.931/0.822; in-dist uniformly worse (−0.058 at L22). Edit: replace signal.tex “are fine” with the numbers; claim = “only uniformly robust readout”, not “mid-layer necessary”. Watch counter-q: audio-native final layer is intact (0.936) — answer = text side still inverts and text→audio calibration only transfers mid-band. (3) label = local-wrong by design (expert-agnostic / pluggable); measured benefit-AUC cost already in app:fixedthr (0.63–0.81; internal 0.840→0.732; WebQ mechanism paragraph). Edit: promote 2 sentences to main text. Optional if labels exist: benefit-trained refit on stored EOT hiddens (CPU, minutes). (4) 20/45 ms = wall-clock prefill-start→L22 score, truncated forward + probe dot, CUDA-synced, H100, 50q, warmup excluded, P50 (P95 25/104); TTFT same bench. Excludes network/RPC — async path, reported separately. Edit: measurement-protocol sentence in app:gatedetail + define in gate.tex. (5) user-facing expert latency IS measured: stall onset p50 22 ms, stall 3.2–4.1 s, expert content audible p50 6.4 s/p95 26/p99 34, dead air after stall p50 ≈2.3–3.2 s, tail P90 ≈27 s. Edit: put content-audible numbers next to

the 22 ms claim in live.tex main text so $25\times$ can't be read as answer latency; optional: exact dead-air p50/95/99 from frozen_v3tts traces (scripts/10_tts_latency.py extension, CPU). STATUS 8-27: all five LANDED (gate.tex $\times 2$, signal.tex, live.tex, appendix $\times 3$); dead air computed from the tts traces — p50 2.5 s / p95 22.4 / p99 30.7, 25/84 turns fully covered (scripts/10_tts_latency.py, dead-air block added). Benefit-trained refit NOT run (only re-scoring exists). $\Rightarrow$ WARNING page budget: same-engine (tectonic) check puts References at p9@78% BEFORE this pass and p10@6% after — main text already exceeded the 8-page target by ~ 0.8 pp before these edits (user build 8-27 01:52 also shows refs on p9). Needs a ~ 1 -page trim decision before upload.]

- [TODO: P2 — 8-27 advisor pass: related.tex gained the full-duplex $\times$ special-token cell (wang2024fsm, zhang2024duplex, veluri2024syncllm, ma2024lslm, fu2024vita, yu2025salmonnomni) + explicit 2×2 gap framing. Cross-check citation list against BrownCat's Feishu duplex list before camera-ready; swap in DuplexPO-style entries if the list has better exemplars.]
- [TODO: P2 — 8-27 mechanism audit (8aw) LANDED: §4.2 no longer asserts turn-control specialization — the speak-mode, modality-axis, pool-identity and drift candidates are all ruled out in D, and the surviving statement is the delivery/distributed-read one. UPDATE same day (8ax): the two direct tests RAN (control-token logit lens + listen/speak intervention, modal_interp.py): output layer is control-biased (log-mass rises to -11.8 while backbone falls to -38.5) but the inverting direction aligns with neither control unembeddings ($\cos .069 = \text{random} .068$) nor the cue axis ($\cos .019 \approx \text{chance}$) — simple takeover story fails its own direct tests; appendix says so. If a reviewer pushes on mechanism, that is the honest answer.]
- [TODO: 8at DONE 2026-08-27 — the FULL loop in the concurrent regime, all 7 pools (36 arms): monotone gains everywhere, selective beats always-escalate on internal (66.2 vs 61.7) and Llama Q.; speakable 44.5 $\rightarrow$ 71.1%; in-regime probe AUC .857 internal / .57–.67 external (calibration-scale gap, 360 rows); latency holds (EOT 23ms, stall onset 27–34ms, local first audio .655s). Landed: app:duplexval concurrent paragraph + tab:conclive, live.tex pointer, Limitations. Page rebalance: tab:signals, tab:policy, the THINK paragraph and “What the signal detects” moved to the appendix; Conclusion is now a paragraph of the Discussion. Main text ends p8, refs p9. Known gaps for camera-ready: full-scale (2310-row) in-regime recalibration; Nemotron under interleave (SATURDAY nice-to-have, after P3 — its NeMo stack is cacheless per-frame so overlap can be constructed offline); official duplex serving format.]
- [TODO: P3 — submission day (8-28): delete this section; final anonymization grep (names, rhe9527, modal URLs, gmail); decide public-repo $\rightarrow$ private; OpenReview form + upload (deadline 8-29 AoE, moderation hard limit Fri 8-28)]