
When Does a Speech Model Know to Hand Off? Reading Competence from Hidden States for Real-Time Escalation

Anonymous Author(s)

Affiliation

Address

email

Abstract

Real-time speech systems answer most queries with a local model, but some queries call for a stronger expert. Selective escalation is well studied for text; speech tightens the constraint: the decision must be made before the model commits to a response, and the signal must survive distribution and modality shift. We ask whether a speech model’s own hidden states already carry such a competence signal. Our key finding is that they do, but not where escalation methods usually look: the conventional final-layer read collapses under distribution shift, while a mid-network read stays reliable, transfers from text-calibrated probes to speech input, and can be taken at a single end-of-turn position before the first output token, in about 30 ms. A linear probe on this representation gates a live escalation loop: it raises deployed answer accuracy from 42% to 63% while beating matched-rate random escalation, and, with weights, features, and layer fixed, improves all five external speech-QA benchmarks over always-local operation. These results identify a mid-network representation, rather than the final-layer read, as the practical signal for deciding when a speech model should hand off.

1 Introduction

Real-time spoken interaction leaves only a few hundred milliseconds between turns [Stivers et al., 2009]. Systems that meet this deadline keep a local speech model as the talker, so that the audio context and the voice stay in one session, and call a stronger remote model on turns the local model cannot answer. In a text interface this is a standard routing problem. In a speech interface the decision must also arrive before the talker commits to a response; otherwise the system must delay the expert call or retract an answer already under way. The central question of this paper is whether a duplex-trained speech model exposes its own likely failure early enough to hand the turn to a stronger model.

Part of the answer exists on the text side. Correctness of a text language model can be estimated from verbalized self-assessment, output uncertainty, or hidden states [Kadavath et al., 2022, Azaria and Mitchell, 2023, Burns et al., 2023, Kossen et al., 2024], and text routers use query-side or pre-generation features to choose a model [Ong et al., 2024, Varshney et al., 2026, Lugoloobi et al., 2026]. On the speech side, duplex models learn interaction control as a generation policy: special or synchronization tokens decide when the model should listen or speak [Wang et al., 2024a, Zhang et al., 2024, Veluri et al., 2024, Fu et al., 2024, Yu et al., 2024], and recent systems extend the same




Figure 1: **Primary turn-based system.** (1) Speech reaches a duplex-trained local model. (2) At the end of the user turn, a mid-network (L22) linear read scores the turn in ~ 30 ms, before the first output token. (3) Lower-risk turns stay local; higher-risk turns are sent to a cloud expert. (4) The local model remains the talker and relays the returned answer. Deployed-channel accuracy rises from 42.2% to 63.3% at 57% escalation in the audio-input/text-output sweep; a TTS-on replay validates the speech path separately.

mechanism to hand work to a reasoning layer or a retrieval service [Huang et al., 2026, Chien et al., 2026].

The trained-control path does not close the gap, for two reasons. First, its cost: the policy is trained into the weights, which takes large amounts of duplex-specific data (hundreds of thousands of hours of synthetic dialogue for SyncLLM, millions of hours of audio for Moshi [Veluri et al., 2024, Défossez et al., 2024]) and risks eroding the backbone’s language ability, which is the stated motivation for frozen-backbone designs [Wang et al., 2024b, Yu et al., 2024]. Second, the interface is wrong for escalation: a speak token decides when the model may talk, not whether its answer will be right, and it arrives as a binary action at the moment of commitment. Even when the token triggers a reasoning step instead of speech [Huang et al., 2026], the decision is still a trained binary action: it offers no calibrated score to threshold at a fixed expert budget, and the policy cannot change without retraining. What remains unknown is whether the escalation decision needs training at all: whether a frozen duplex-trained checkpoint already carries a competence signal that can be read before response commitment. Text-probe results make this plausible, but they do not say where such a read stays reliable after duplex speech training, whether a probe calibrated on inexpensive text data still reads speech-input states, or whether the read arrives early enough to act on.

These three conditions are exactly what a deployment needs: a score that separates individual failures rather than memorizing the query types seen in calibration, that survives the text-to-speech shift, and that improves which turns are escalated at a fixed call budget. If such a signal exists, any frozen duplex model gains selective escalation for the cost of a linear probe; if it does not, speech escalation inherits the training costs above.

We test these conditions on MiniCPM-o 4.5 [Cui et al., 2026], a 9B omni model with full-duplex training, against matched controls: raw backbones, streaming-omni and audio-only fine-tunes, a frozen-backbone adapter, and a second duplex generation. The study uses judged failure labels on 600 public-benchmark queries in five frozen pools, leave-one-pool-out (LOPO) transfer as the primary robustness test, matched text and speech inputs, and a complete live escalation loop. We compare decode-time uncertainty, verbalized self-assessment, and linear hidden-state probes across layers, then freeze the winning read for evaluation on public speech benchmarks. Model checkpoints stay frozen; only linear probes are fit on stored activations. Throughout, duplex-trained refers to the weights; the primary live evaluation is turn-based, with audio input and text output, and the concurrent and native full-duplex serving regimes are validated separately (Appendices H and O).

This study makes three contributions:

1. **Transfer moves the readable layer.** The conventional final-layer probe holds in distribution but inverts on held-out math (AUC 0.372); the same last-token read at a mid-network layer reaches 0.931. Matched comparisons tie the late-layer failure to the duplex checkpoints we test; its mechanism remains open (Appendix D).
2. **The mid-layer read satisfies the deployment conditions.** A text-calibrated probe at the same layer reads speech-input states, reproduces on real human recordings, and is available at the end of the user turn in ~ 30 ms, before the first output token.
3. **The read is sufficient to drive escalation.** In the live loop it raises deployed-channel accuracy from 42.2% to 63.3% at 57% escalation, beats matched-rate random escalation at every tier, and dominates the model’s own built-in reasoning mode on both accuracy and latency (Appendix F). With its layer and features frozen, the gate improves all five external speech-QA benchmarks over always-local operation, although selectivity over random is limited on some pools.

The claim is therefore about readout rather than mechanism: the model does not explain why it will fail, but its likely failure is more reliably readable mid-network than at the conventional output interface, early enough to support a live handoff.

2 Related Work

Pre-answer competence estimation. Text-language-model correctness can be estimated from verbal self-evaluation [Kadavath et al., 2022, Lin et al., 2022], output uncertainty, or hidden representations [Azaria and Mitchell, 2023, Burns et al., 2023, Marks and Tegmark, 2024]. Intermediate layers can carry more useful information than the output distribution [Orgad et al., 2025], and trained hidden-state probes are strong factual-confidence estimators [Mahaut et al., 2024]. Semantic entropy improves uncertainty estimation by comparing sampled answers [Kuhn et al., 2023, Farquhar et al., 2024]; Semantic Entropy Probes amortize this signal into a pre-generation or post-hoc hidden-state read [Kossen et al., 2024]. Attention-map probes and multimodal uncertainty fusion extend the same broad idea [Chuang et al., 2024, Zhang et al., 2026a]. This work establishes that competence is readable in text models. It does not show where the read remains reliable after duplex speech training or whether a probe calibrated on text transfers to speech-input states.

Selective routing. Text routing splits into two families by where the decision lives. In the first, the decision is read from outside the generator: HybridLLM [Ding et al., 2024], RouteLLM [Ong et al., 2024], and query-side routers [Xue et al., 2025, Hu et al., 2024, Li et al., 2026] choose a model for a complete text query from query features or preference data; FrugalGPT organizes cascades that can score an earlier model’s response [Chen et al., 2023], while AutoMix verifies a drafted answer before escalating [Aggarwal et al., 2024]. Closest to our signal, recent methods read pre-generation activations or fit linear success probes on frozen states [Varshney et al., 2026, Lugoloobi and Russell, 2025, Lugoloobi et al., 2026]. In the second family, the decision is trained into the generator: Toolformer and Self-RAG learn tool or reflection tokens [Schick et al., 2023, Asai et al., 2024]; Co-LLM and CITER learn token-level deferral [Shen et al., 2024, Zheng et al., 2025]; and pause, thought, or think/no-think tokens allocate more internal computation [Goyal et al., 2024, Zelikman et al., 2024, Yang et al., 2025, DeepSeek-AI et al., 2025]. Learning-to-defer and early-exit methods make related decisions within a model [Madras et al., 2018, Schuster et al., 2022]. Our contribution is not the linear pre-generation probe by itself. We test whether this kind of read survives duplex training and audio input, and use a deployment-time threshold without modifying the speech checkpoint.

Duplex control and live handoff. End-to-end spoken dialogue includes native duplex models [Nguyen et al., 2023, Défossez et al., 2024, Zeng et al., 2024, Fang et al., 2026], streaming omni models [Xu et al., 2025, Cui et al., 2026], frozen-backbone adapters [Wang et al., 2024b], and audio fine-tunes [Chu et al., 2024]. Most control work in these systems concerns the floor: state, idle,

synchronization, or interruption tokens [Wang et al., 2024a, Zhang et al., 2024, Veluri et al., 2024, Ma et al., 2024, Fu et al., 2024, Yu et al., 2024], structured action channels [Zhang et al., 2026b], or classifier heads over a frozen or lightly adapted backbone [Wang et al., 2024b, Chen et al., 2025, Liao et al., 2025]. Training this control is costly enough that reducing the cost is itself a research goal [Xu et al., 2024], and trained turn-taking remains unreliable across models [Lin et al., 2025]. These mechanisms decide when the model may speak; even the most recent native design adds only a few turn tokens to the vocabulary and answers every turn locally, whatever the model’s competence on that turn [Fang et al., 2026].

The direct neighbors of our gate are trained think triggers. DuplexOmni emits a think-control token at the end of a user utterance to hand work to a separate reasoning layer [Huang et al., 2026]; chronological thinking interleaves reasoning with the speech stream [Wu et al., 2026]; MiniCPM-o 4.5 ships a built-in reasoning mode [Cui et al., 2026]; MoshiRAG triggers asynchronous retrieval from the model’s emerging output [Chien et al., 2026]. These systems show that a live handoff is possible. But the trigger is trained into the weights from purpose-built duplex data, so changing the policy or the downstream resource means retraining; it is a binary action with no calibrated score, so the escalation rate cannot be set to a call budget; and the trained policies are hard to reuse or compare, because released artifacts often include the training recipe but not the trained checkpoint [Huang et al., 2026]. None of these works evaluates a competence score read from a frozen checkpoint before generation. Our probe replaces the trained trigger with such a read: it needs no training of the speech model, exposes a thresholdable score, and on our target model it dominates the built-in reasoning mode on both accuracy and latency (Appendix F).

Positioning. Prior work therefore establishes three ingredients separately: competence can be read from text models, text queries can be routed, and duplex models can coordinate speaking or external computation. To our knowledge, the reviewed literature does not establish their conjunction: whether a frozen duplex-trained speech model exposes a pre-answer competence signal that survives unseen query pools and text-to-speech transfer, and whether that signal arrives early enough to improve live selective escalation. We evaluate these conditions directly. The claims are about readout robustness and system utility; the training mechanism behind the late-layer failure remains open.

3 Reading Competence from Hidden States

This section establishes where and when a competence signal can be read from the frozen checkpoint. The analyses use the target model and matched controls introduced above: 600 public-benchmark queries in five frozen pools with judged failure labels, and logistic probes fit on stored hidden states. Full details of models, data, labels, and probe fitting are in §4.

3.1 Signal comparison under distribution shift

Table 3 compares the candidate signals on MiniCPM-o 4.5. In distribution the final-layer probe looks healthy (AUC 0.822), but most of that accuracy reflects query type rather than instance difficulty: the same hidden state predicts which pool a query came from, a classifier given only the pool identity recovers most of the probe’s margin, and appending true pool labels to the probe’s features adds nothing. The same decomposition holds at scale on the deployed calibration mix (Appendix O). We therefore treat leave-one-pool-out (LOPO) transfer as the primary metric, and under it the picture reverses: with text input the final-layer read falls to AUC 0.372 on held-out math, and when the trap pool is held out the probe scores it on the confident side even though the model fails every query in it. A router built on query text shows the same dependence on pool structure and also degrades under LOPO, with little improvement from substantially more data (Appendices B, E). In-distribution routing accuracy alone therefore cannot establish instance-level transfer.

Verbalized self-assessment transfers differently. Asked before answering, $p(\text{True})$ catches exactly the trap pool the probe misses (0.945). Asked after drafting, the model endorses its own confident-

Table 1: Main results: selective escalation on the internal test set and five public speech-QA benchmarks (post-relay answer accuracy, %, audio-input/text-output live loop). The probe is fixed; tier thresholds are label-free per-pool quantiles (the internal tiers realize 14/31/57% escalation). Matched-rate random mixes the always-local and measured always-escalate arms at each pool’s realized rate. avg and Δ (average gain over always-local) cover the five external pools. The bottom block applies the same recipe to NemotronLabs-VoiceChat-11B offline, re-mixing cached local and expert outcomes (Appendix C). AlpacaEval, reported separately on its 1–5 scale, changes from 3.99 to 4.35 on MiniCPM but does not beat matched-rate random; the second family changes from 3.55 to 4.46 versus 4.23 for random. Per-tier random references and gold-input ceilings: Appendices I–J.

	Internal	TriviaQA	WebQ	Llama Q.	SD-QA	Reason. zh	avg	Δ
<i>live post-relay view (audio input, text output)</i>								
always-local	42.2	66.4	56.8	83.6	51.5	58.9	63.4	—
+ gate @ 15%	52.8	73.2	62.8	90.4	60.0	65.3	70.3	+6.9
+ gate @ 30%	59.2	80.0	68.0	90.4	72.0	71.3	76.3	+12.9
+ gate @ 50%	63.3	86.0	73.2	94.8	79.5	77.2	82.1	+18.7
matched random @ 50%	56.1	77.2	68.8	88.6	70.0	71.3	75.2	+11.8
always-escalate (measured)	66.5	88.0	80.8	93.6	88.5	83.7	86.9	+23.5
always-expert (gold-text ceiling)	92.2	96.8	86.4	92.8	93.0	87.1	91.2	+27.8
<i>second family: NemotronLabs-VoiceChat-11B (offline re-mix, official judges)</i>								
always-local	—	35.6	37.6	70.5	31.0	—	43.7	—
+ gate @ 15%	—	49.8	48.0	78.6	42.5	—	54.7	+11.0
+ gate @ 30%	—	61.5	58.8	83.8	54.5	—	64.6	+21.0
+ gate @ 50%	—	74.1	69.6	87.6	67.5	—	74.7	+31.0
matched random @ 50%	—	65.6	60.2	81.8	60.0	—	66.9	+23.2
always-expert (ceiling)	—	95.5	82.8	93.2	89.0	—	90.1	+46.4

160 wrong answers and the signal inverts: post-draft $p(\text{True})$ is the strongest single readout in distribu-
 161 tion but no longer provides a pre-answer routing signal.

162 3.2 Layer and position sweep

163 Matched comparisons tie the failure to the duplex checkpoints rather than to the backbone alone.
 164 The raw Qwen backbone, subsampled to match MiniCPM-o 4.5’s per-pool failure rates, holds LOPO
 165 math near 0.96 where the duplex checkpoint scores 0.372, and the Qwen2.5 family shows the same
 166 ordering, raw above streaming-omni above duplex (Table 4, appendix). The Qwen2-Audio com-
 167 parison is not attribution evidence because the audio fine-tune also changes task capability and the
 168 failure distribution substantially (Appendix B). A layer $\times$ position sweep localizes the difference
 169 (Figure 2): it is concentrated at the last-token position in late layers on text input, mean-pooled
 170 reads degrade far less, and at L22 the last-token probe reaches **0.931** on the same held-out math
 171 pool, close to the raw backbone.

172 The failure is input-specific rather than universal: with audio calibration and audio input, the final-
 173 layer probe reaches 0.936 on LOPO math. We use L22 because one read point must remain robust
 174 for shifted text queries and support text-to-speech transfer, not because the final-layer signal is absent
 175 on the model’s native audio input.

176 The sweep establishes where the transferable read survives, not why it weakens late. Direct tests do
 177 not support the simple hypothesis that turn-control state takes over the last-token representation: the
 178 layer at which performance drops does not move in speak mode, and the inverting probe direction
 179 aligns with neither the control vocabulary nor an explicit listen/speak direction. A decomposition
 180 further shows that a distributed transferable component weakens over the final layers of the two
 181 duplex checkpoints but not their raw backbones. These observations constrain the explanation but
 182 do not identify a training objective or causal mechanism; the full audit is in Appendix D.

183 The mid-network pattern also appears outside the MiniCPM family: on NemotronLabs-VoiceChat-
 184 11B, a hybrid-Mamba duplex model, the same recipe peaks mid-network and transfers to the external
 185 pools from a 600-query calibration set (Appendix C).




Figure 2: The core representational result: LOPO math AUC by layer and position. On the duplex model the standard (final-layer, last-token) read inverts on text input while mid-network layers hold ≥ 0.9 ; raw backbones show no cliff. The deployed gate reads L22.

3.3 Modality transfer

Probes trained on text hidden states and scored on speech-input hidden states show where the representation becomes shared across modalities: early layers are modality-specific, and from about half depth the transfer plateaus (0.855 at the deployed L22). This replicates on MiniCPM-o 2.6 and cross-family on Qwen2.5-Omni. The decisive control is Freeze-Omni: identical frozen LLM weights with audio attached by an adapter, and cross-modal transfer stays far below the end-to-end models. Text calibration therefore transfers in the end-to-end speech models tested here, but not in the adapter-attached control. Verbalized introspection shows the complementary failure: on audio input, pre-answer $p(\text{True})$ collapses on both MiniCPM duplex generations while both non-duplex controls stay intact. Controlled arms show the verbal self-check depends on the text-token pathway (duplicating the question as text restores it), while a linear probe remains usable on the same forward pass, and the restored self-check is additive with the probe (Appendix B, Table 7). The mid-layer band also replicates on 200 real human recordings from SD-QA, so the result is not limited to the TTS rendering. The live system keeps L22 and adds two pooled positions to form the deployed probe (Appendix B).

3.4 Decision timing and escalation budget

Scoring latency. The gate is a linear read on L22 states the forward pass already computes: a 4096-d dot product per position (the deployed probe adds two pooled positions at no added scoring latency; Appendix B). On an H100, measured from prefill start with CUDA synchronization, the score is available at P50 in 20 ms for text and 45 ms for audio, against a time to first token of 36/68 ms on the same benchmark. The measurement excludes network time, but it establishes the ordering the system needs: the cloud request can start before local decoding commits to an answer. Waiting for decoded tokens does not improve the read in our data and postpones the expert request (Appendix F).

Read position in the audio loop. In the streaming loop, scores taken from intermediate audio chunks are weaker than a read at the end of the turn and can miss information in the last chunk. The system therefore reads once, at the end of the user turn: after the final audio chunk it prefills the assistant turn and reads L22 at its start. In a 360-query replay, this single read reaches AUC 0.887 in about 30 ms, against 0.843 for the offline reference, and arrives before the first output token (Appendix G).

From score to escalation budget. Thresholding the score converts it into a call budget. On the frozen test split, the probe beats matched-rate random escalation at every operating point and recovers over half of the local-to-expert gap with a third of the expert calls (Figure 8, appendix). In distribution it is statistically tied with a pool-identity classifier, as §3.1 predicts; its advantage appears under transfer, where pool identity stops being sufficient. The gate is fit to local failure, not expert benefit: a benefit label would specialize the ranking to the current expert. The frozen scores remain useful under the stricter benefit label, but pools with little expert headroom remain a system-level limit (Appendix I). Comparisons with the model’s built-in reasoning mode are in Appendix F.

4 Experimental Setup

4.1 Models

The target model is MiniCPM-o 4.5 (9B; omni + full-duplex fine-tune of Qwen3-8B), run in bf16 with greedy decoding and seed 42. To attribute effects to the kind of fine-tune rather than the backbone, we assemble ten models spanning raw backbones, a second duplex generation (MiniCPM-o 2.6), a streaming-omni fine-tune (Qwen2.5-Omni-7B), an audio-only fine-tune (Qwen2-Audio-7B), and a frozen-backbone adapter model (Freeze-Omni, LLM weights byte-identical to its raw control); the full roster is Table 2 (appendix). The comparison statistic is always the within-pair difference. All model checkpoints remain frozen throughout; the escalation expert is gpt-5.5 at low reasoning effort.

4.2 Dataset and labels

The internal dataset is 600 queries in five pools, frozen before any gate experiment (360 calibration / 240 test): *trap* = SimpleQA [Wei et al., 2024] (the target fails 100%), *hard-knowledge* = MMLU-Pro [Wang et al., 2024c], *easy-fact* = TriviaQA [Joshi et al., 2017], *easy-chat* = Dolly/Alpaca [Conover et al., 2023, Taori et al., 2023], *hard-math* = GSM8K tail + MATH-500 [Cobbe et al., 2021, Hendrycks et al., 2021, Lightman et al., 2024]; public benchmarks only, no LLM-generated queries. Adequacy labels come from an LLM judge [Zheng et al., 2023] with structured outputs, blind to the gate and validated by a stronger re-judge (agreement 96.2% text / 94.5% audio; every headline number moves ≤ 1 point under the stronger labels). The audio arm renders the same frozen pool to speech with single-voice TTS (content unchanged; only modality varies), following Spoken-SQuAD / VoiceBench practice [Li et al., 2018, Chen et al., 2024]. Judge prompts, pins, and spend: Appendix M–N.

4.3 Probe training and baseline signals

We compare three signal families: (i) decode scalars: max token entropy or margin over the first k decode steps; (ii) linear probes: logistic regression on hidden states, with layer and position selected on calibration data; and (iii) verbalized self-evaluation [Kadavath et al., 2022]: $p(\text{True})$ -pre asks “would you answer this correctly?” before answering, $p(\text{True})$ -post asks about a drafted answer; $P(\text{Yes})$ is read from the first token’s distribution. The representation analyses use a single-position, 4096-d L22 probe fit on the 360-query calibration split. The live system uses the same layer plus two pooled positions, fit on a wider 2,310-query calibration mix that is disjoint from the external benchmarks. Concurrent and native serving change the read point, so those validations refit the same feature set in regime; Appendix B records the full calibration ledger.

4.4 Benchmarks and evaluation protocol

Audio-side findings are validated on real human recordings (SD-QA, Faisal et al., 2021); the frozen gate is evaluated on external public speech benchmarks with their official audio (§5.2). Following §3.1, we use LOPO AUC to test distribution transfer and text-trained/audio-scored AUC to test modality transfer. For the live system, the key comparison is against matched-rate random escalation:




Figure 3: Live escalation on the internal test set (speakable subset, $n=218$). Blue: deployed-channel accuracy, where the expert answers from the talker’s own transcript. Green: the same sessions re-scored with the expert answering the reference transcript, which factors out the relay channel. Each curve has its own matched-rate random reference (dashed, same floor and endpoint). Selection beats random in both, by 7.2–9.4 points deployed and 5.3–8.3 points channel-controlled; the gap between the curves is the cost of the relay channel.

an improvement over always-local operation can come from the expert alone, while an improvement over random at the same call rate measures the value of the gate’s ranking.

4.5 Live loop and escalation path

The live loop is turn-based: audio streams in 1 s chunks, and the talker speaks only after the user’s turn ends; the model’s native listen/speak head is not engaged, which isolates competence handoff from turn-taking (native-head validations are reported in §5.1 and Appendix O). The loop uses the deployed three-position probe of §4.3 (frozen-test AUC 0.879). Escalation runs in four tiers from conservative to aggressive; each tier’s threshold is a quantile of the probe’s score, so a tier sets the call budget without correctness labels. If a turn escalates, the talker first produces a short stall phrase; the model’s own transcription of the query is sent to gpt-5.5, and the returned answer is injected as context for the talker to relay. The expert request runs asynchronously. Partial-query prefetch, expert reasoning effort, and injection controls are in Appendix G.

5 Results

5.1 Live escalation

We ran the 240 frozen test queries through four escalation tiers, each a complete turn-based loop with audio input (§4.5). The primary sweep retains the model’s output as text; a separate run of the balanced tier with speech synthesis enabled leaves all 84/240 gate decisions unchanged and measures speech onset. We score deployed-channel answer accuracy: the correctness of the relayed content after the full system path. Headline numbers use a pre-registered input filter that removes 21 formula-LaTeX queries and one broken recording, because the speech rendering destroys the symbolic input before the model receives it. Accuracy rises monotonically, 42.2% → 52.8% → 59.2% → 63.3%

at 0/14/31/57% escalation, and every tier’s gain is significant (paired bootstrap; Table 1, Figure 3). The gate also exceeds matched-rate random escalation at every tier: by 7.2–9.4 points on deployed-channel accuracy, and by 5.3–8.3 points when the relay channel is factored out (Figure 3).

A measured always-escalate arm reaches 66.5%. The gain from 57% to 100% escalation is small because the expert sees the model’s own transcription: the same sessions reach 92.2% when the expert receives the reference transcript, so the transcription channel, not selection, is the limit at high escalation rates. Live accuracies have a replication floor of ± 2 –3 points; no claim rests on a smaller difference. Full tables, latency profiles, and boundary cases are in Appendix G.

Serving-regime checks. Replaying the loop with the talker enabled leaves the frozen ranking nearly unchanged (AUC 0.871/0.856 versus 0.866 with the talker disabled). A teacher-forced concurrent-prefill variant is less stable and requires an in-regime refit; its external results remain exploratory (Appendix H). Under the model’s native full-duplex head, the read moves to the head’s own listen→speak commit point, and the probe is refit in regime on a widened 5,228-row calibration merge under the official serving configuration, with labels judged on the native answers. The refit reads at AUC 0.876 internally and 0.771 on the four English external QA pools. The Mandarin reasoning pool is lower, at 0.621; a 1,500-question held-out Mandarin set of facts, knowledge, and commonsense reads at 0.811, so the gap is multi-step reasoning rather than language (Appendix O). Re-mixing the native gate’s decisions with cached expert outcomes beats matched-rate random on all five English pools at both deployed tiers; with per-language thresholds the Mandarin pool escalates at the nominal rates and beats matched-rate random at the two lower tiers. A full native live run on the internal 240 queries, under the earlier inherited serving configuration, improves accuracy from 0.371 to 0.537 at 45% escalation, below the prediction of 0.596 from the same re-mix arithmetic, because relaying the returned answer is lossy. These checks support the ranking beyond the primary turn-based loop, but they also show that read points, probe weights, and thresholds must be calibrated for the serving regime (Appendix O).

5.2 Transfer to public speech benchmarks

The shortcut analysis of §3.1 makes external transfer necessary. We freeze the probe weights, feature set, and layer selected on the internal calibration mix, then run the live loop of §4.5 on six public speech benchmarks with audio input and text output: Speech TriviaQA, Speech Web Questions, Llama Questions, and Reasoning QA from OpenAudioBench [Li et al., 2025] (spoken renderings of Joshi et al., 2017, Berant et al., 2013, and the set of Nachmani et al., 2024), SD-QA (real dialectal human speech) [Faisal et al., 2021], and VoiceBench AlpacaEval [Chen et al., 2024, Taori et al., 2023]. All audio is the benchmarks’ official pre-rendered audio; no evaluation query is used to fit the probe. The only per-pool adjustment is label-free: tier thresholds are quantiles of that pool’s score distribution, requiring no correctness labels or judge calls. This controls the call budget but cannot improve ranking. A single absolute threshold also beats matched-rate random on four of five QA pools, but its escalation rate varies from 2% to 48% (Appendix I). We use each benchmark’s official judge when available and verify that the local-only controls are on the published scale.

The frozen scores remain discriminative outside the calibration mix: AUC ranges from 0.683 to 0.806 across the five QA pools, with a mean of 0.771 (Appendix J). The gate also improves deployed-channel accuracy over always-local operation on every QA benchmark and at every budget; at 50% escalation, mean accuracy rises from 63.4% to 82.1%. These gains establish system utility but do not by themselves establish selectivity, because random escalation also adds expert answers.

Llama Questions provides the cleanest fixed-budget test. It contains a single query type, the kept-local subset scores 97.6% while the escalated subset scores 69.6% ($z=6.46$), and expert answers improve the latter to 88.8% ($p<.0001$). At 50% escalation, the measured system reaches 94.8%, compared with 93.6% for always escalating. Matched-rate precision reaches 95% of its base-rate cap (Appendix J). Cross-lingual transfer also improves when the English calibration mix is widened: Chinese Reasoning QA rises from AUC 0.621 to 0.683 despite adding no Chinese calibration data.

The transfer results also expose the boundary between predicting local failure and creating expert benefit. SD-QA is a positive real-speech case and exceeds matched-rate random at all three budgets. Web Questions is less stable, with little gain over random at the low and high budgets; in its top-scoring 15%, many local failures are also missed by the expert (Appendix I). AlpacaEval is a clearer negative: its open-ended failures do not present the retrieval-failure pattern the probe reads. On the second duplex family, which has more expert headroom on these pools, the same recipe beats matched-rate random on all five evaluated pools, including AlpacaEval ($p \leq .0003$). Thus external utility depends on both the quality of the competence ranking and whether the selected failures are repairable by the expert (Appendices C, I, and J).

6 Discussion

Interpretation. The competence signal this system runs on was already in the frozen weights. Nothing was trained into the speech model, and the read costs one dot product on states the forward pass computes anyway, so a deployed duplex model can gain selective escalation without new duplex data, a new control token, or a retrained policy. What matters is where the read is taken. The final layer is the default choice because it is the interface to the output, and on these duplex checkpoints it is the wrong one: the same linear read that inverts at the output holds up in the middle of the network. This also points back at duplex training itself. The matched comparisons show that the fine-tune, not the backbone, is what breaks the late-layer read, so a training recipe that adds interaction control may be degrading the model’s own competence interface while leaving its answers intact. We do not know the mechanism, and a model that reads its own state before answering still cannot explain why it will fail.

Scope of the signal. The probe is strongest on retrieval-like failures, where the state reflects a missing answer before generation. It is weak on false premises and complementary to signals that appear during decoding (Appendix G). This boundary explains why failure AUC and system benefit are not the same quantity. A correct ranking cannot help when the expert fails on the same turns, and an open-ended quality benchmark may not contain the factual-retrieval event the probe detects. The unstable Web Questions curve and the negative AlpacaEval result are part of that boundary rather than exceptions to it; SD-QA, by contrast, confirms selective gains on real human speech. Selective escalation needs both a readable local failure and expert headroom.

Calibration. The model weights stay frozen, but the gate is still a fitted component. Text-only calibration is enough for the cross-modal result on end-to-end speech models, the deployed probe uses a wider benchmark-disjoint mix, and both the concurrent and native read points require refitting in their own regime. The method avoids fine-tuning the talker and moves the operating budget through a threshold, but it does not remove the need for representative calibration data.

Limitations. The internal study uses one phrasing per $p(\text{True})$ variant, and its real-speech validation is scripted read speech in a single dialect. Formula-symbolic content is out of scope for the audio arm, since any speech rendering destroys it. Labels come from one judge family, with no human labels; a stronger re-judge bounds the resulting shifts at about one point. Each query is run once per tier, which puts a replication floor of two to three points on every live accuracy, and no claim here rests on a smaller difference. The second-family replication covers mid-band readability and frozen-recipe transfer only, not matched-pair attribution, the live loop, or the concurrent regime. Layer and features were chosen by cross-validation on the internal calibration split and then confirmed on held-out pools, rather than by nested selection inside each held-out pool. The mechanism behind the late-layer decay on text input remains unresolved (Appendix D).

Serving regime is the sharpest practical limit. Calibration does not carry across regimes: under concurrent prefill the frozen read degrades and needs an in-regime refit, and the native read needs its own refit, wider calibration, and labels judged on native answers before it matches the turn-based loop (Appendices H, O). That native read also appears saturated, in that more labels of the same

kind, larger heads, and a distilled second-signal student all leave it close to where it is. Finally, the evaluation is scripted and single-turn. Disfluent correction, chained tool calls, and multi-turn grounding remain open [Lin et al., 2026], and in native testing a new question asked during the expert wait usually derails the pending relay, so a production system has to cancel or replace an in-flight request when the user moves on.

References

- Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqui, and Mausam. AutoMix: Automatically mixing language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2310.12963.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.11511.
- Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. arXiv:2304.13734.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from question-answer pairs. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2013. WebQuestions.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2212.03827.
- Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.
- Qian Chen, Yafeng Chen, Yanni Chen, et al. MinMo: A multimodal large language model for seamless voice interaction. *arXiv preprint arXiv:2501.06282*, 2025.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench: Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024.
- Chung-Ming Chien, Manu Orsini, Eugene Kharitonov, Neil Zeghidour, Karen Livescu, and Alexandre Défossez. MoshiRAG: Asynchronous knowledge retrieval for full-duplex speech language models. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2604.12928.
- Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, et al. Qwen2-Audio technical report. *arXiv preprint arXiv:2407.10759*, 2024.
- Yung-Sung Chuang, Linlu Qiu, Cheng-Yu Hsieh, Ranjay Krishna, Yoon Kim, and James Glass. Lookback lens: Detecting and mitigating contextual hallucinations in large language models using only attention maps. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. GSM8K.
- Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick Wendell, Matei Zaharia, and Reynold Xin. Free dolly: Introducing the world’s first truly open instruction-tuned LLM. Databricks blog, 2023. databricks-dolly-15k.
- Junbo Cui, Bokai Xu, Chongyi Wang, Tianyu Yu, Weiye Sun, Yingjing Xu, et al. MiniCPM-o 4.5: Towards real-time full-duplex omni-modal interaction. *arXiv preprint arXiv:2604.27393*, 2026.

422 DeepSeek-AI, Daya Guo, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via
423 reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

424 Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou,
425 Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dia-
426 logue. *arXiv preprint arXiv:2410.00037*, 2024.

427 Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks
428 V. S. Lakshmanan, and Ahmed Hassan Awadallah. Hybrid LLM: Cost-efficient and quality-
429 aware query routing. In *International Conference on Learning Representations (ICLR)*, 2024.
430 arXiv:2404.14618.

431 Fahim Faisal, Sharlina Keshava, Md Mahfuz ibn Alam, and Antonios Anastasopoulos. SD-QA:
432 Spoken dialectal question answering for the real world. In *Findings of the Association for Com-
433 putational Linguistics: EMNLP 2021*, 2021. arXiv:2109.12072.

434 Qingkai Fang, Shoutao Guo, and Yang Feng. BayLing-Duplex: Native full-duplex speech dialogue
435 with a single autoregressive LLM. *arXiv preprint arXiv:2606.14528*, 2026.

436 Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large
437 language models using semantic entropy. *Nature*, 630:625–630, 2024.

438 Chaoyou Fu, Haojia Lin, Zuwei Long, Yunhang Shen, Yuhang Dai, Meng Zhao, et al. VITA: To-
439 wards open-source interactive omni multimodal LLM. *arXiv preprint arXiv:2408.05211*, 2024.

440 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
441 Nagarajan. Think before you speak: Training language models with pause tokens. In *ICLR*, 2024.

442 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
443 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
444 *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks*, 2021.
445 arXiv:2103.03874.

446 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
447 Keutzer, and Shriyash Kaustubh Upadhyay. RouterBench: A benchmark for multi-LLM routing
448 systems. *arXiv preprint arXiv:2403.12031*, 2024.

449 Muye Huang, Lingling Zhang, Xingyu Yu, Lei Shi, Zhanyu Ma, Jun Xu, Jiuchong Gao, Jinghua
450 Hao, Renqing He, and Jun Liu. DuplexOmni: Real-time listening, seeing, thinking, and speaking
451 for full-duplex interaction. *arXiv preprint arXiv:2606.09186*, 2026.

452 Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale dis-
453 tantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual
454 Meeting of the Association for Computational Linguistics (ACL)*, 2017.

455 Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez,
456 Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language mod-
457 els (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

458 Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. Se-
459 mantic entropy probes: Robust and cheap hallucination detection in LLMs. *arXiv preprint
460 arXiv:2406.15927*, 2024.

461 Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for
462 uncertainty estimation in natural language generation. In *International Conference on Learning
463 Representations (ICLR)*, 2023. arXiv:2302.09664.

464 Chia-Hsuan Li, Szu-Lin Wu, Chi-Liang Liu, and Hung-yi Lee. Spoken SQuAD: A study of miti-
465 gating the impact of speech recognition errors on listening comprehension. In *Interspeech*, 2018.
466 arXiv:1804.00320.

467 Hao Li, Yiqun Zhang, Zhaoyan Guo, Chenxu Wang, Shengji Tang, Qiaosheng Zhang, Yang Chen,
468 Biqing Qi, Peng Ye, Lei Bai, Zhen Wang, and Shuyue Hu. LLMRouterBench: A massive bench-
469 mark and unified framework for LLM routing. *arXiv preprint arXiv:2601.07206*, 2026.

470 Tianpeng Li, Jun Liu, Tao Zhang, Yuanbo Fang, Da Pan, Mingrui Wang, et al. Baichuan-Audio:
471 A unified framework for end-to-end speech interaction. *arXiv preprint arXiv:2502.17239*, 2025.
472 Introduces OpenAudioBench.

473 Borui Liao, Yulong Xu, Jiao Ou, Kaiyuan Yang, Weihua Jian, Pengfei Wan, and Di Zhang. Flex-
474 Duo: A pluggable system for enabling full-duplex capabilities in speech dialogue systems. *arXiv*
475 *preprint arXiv:2502.13472*, 2025.

476 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
477 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International*
478 *Conference on Learning Representations (ICLR)*, 2024. arXiv:2305.20050; source of the MATH-
479 500 subset.

480 Guan-Ting Lin, Jiachen Lian, Tingle Li, Qirui Wang, Gopala Anumanchipalli, Alexander H. Liu,
481 and Hung-yi Lee. Full-Duplex-Bench: A benchmark to evaluate full-duplex spoken dialogue
482 models on turn-taking capabilities. *arXiv preprint arXiv:2503.04721*, 2025.

483 Guan-Ting Lin, Chen Chen, Zhehuai Chen, and Hung-yi Lee. Full-Duplex-Bench-v3: Bench-
484 marking tool use for full-duplex voice agents under real-world disfluency. *arXiv preprint*
485 *arXiv:2604.04847*, 2026.

486 Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in
487 words. *Transactions on Machine Learning Research*, 2022. arXiv:2205.14334.

488 William Lugoloobi and Chris Russell. LLMs encode how difficult problems are. *arXiv preprint*
489 *arXiv:2510.18147*, 2025.

490 William Lugoloobi, Thomas Foster, William Bankes, and Chris Russell. LLMs encode their failures:
491 Predicting success from pre-generation activations. *arXiv preprint arXiv:2602.09924*, 2026.

492 Ziyang Ma, Yakun Song, Chenpeng Du, Jian Cong, Zhuo Chen, Yuping Wang, Yuxuan Wang, and
493 Xie Chen. Language model can listen while speaking. *arXiv preprint arXiv:2408.02622*, 2024.

494 David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and
495 accuracy by learning to defer. In *Advances in Neural Information Processing Systems (NeurIPS)*,
496 2018.

497 Matéo Mahaut, Laura Aina, Paula Czarnowska, Momchil Hardalov, Thomas Müller, and Lluís
498 Màrquez. Factual confidence of LLMs: on reliability and robustness of current estimators. In
499 *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.

500 Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language
501 model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
502 arXiv:2310.06824.

503 Eliya Nachmani, Alon Levkovitch, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh
504 Mariooryad, Ehud Rivlin, RJ Skerry-Ryan, and Michelle Tadmor Ramanovich. Spoken question
505 answering and speech continuation using spectrogram-powered LLM. In *International Confer-*
506 *ence on Learning Representations (ICLR)*, 2024. arXiv:2305.15255; Spectron, source of the
507 Llama Questions dataset.

508 Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden
509 Tomasello, Robin Algayres, Benoît Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux. Gen-
510 erative spoken dialogue language modeling. *Transactions of the Association for Computational*
511 *Linguistics*, 2023. arXiv:2203.16502.

512 Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez,
513 M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data.
514 *arXiv preprint arXiv:2406.18665*, 2024.

515 Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Kotek, and
516 Yonatan Belinkov. LLMs know more than they show: On the intrinsic representation of
517 LLM hallucinations. In *International Conference on Learning Representations (ICLR)*, 2025.
518 arXiv:2410.02707.

519 Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettle-
520 moyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach them-
521 selves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
522 arXiv:2302.04761.

523 Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and
524 Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information*
525 *Processing Systems (NeurIPS)*, 2022. arXiv:2207.07061.

526 Shannon Zejiang Shen, Hunter Lang, Bailin Wang, Yoon Kim, and David Sontag. Learning to de-
527 code collaboratively with multiple language models. In *Proceedings of the 62nd Annual Meeting*
528 *of the Association for Computational Linguistics (ACL)*, 2024. arXiv:2403.03870; Co-LLM.

529 Tanya Stivers, N. J. Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann,
530 Gertie Hoymann, Federico Rossano, Jan Peter de Ruiter, Kyung-Eun Yoon, and Stephen C. Levin-
531 son. Universals and cultural variation in turn-taking in conversation. *Proceedings of the National*
532 *Academy of Sciences*, 106(26):10587–10592, 2009.

533 Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy
534 Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following LLaMA model.
535 GitHub repository, 2023. zh split: shibing624/alpaca-zh.

536 Tanay Varshney, Annie Surla, Michelle Xu, Gomathy Venkata Krishnan, Maximilian Jeblick, David
537 Austin, Neal Vaidya, and Davide Onofrio. LLM router: Rethinking routing with prefix activations.
538 *arXiv preprint arXiv:2603.20895*, 2026.

539 Bandhav Veluri, Benjamin N. Peloquin, Bokai Yu, Hongyu Gong, and Shyamnath Gollakota. Be-
540 yond turn-based interfaces: Synchronous LLMs as full-duplex dialogue agents. In *Proceedings*
541 *of EMNLP (Main)*, 2024. arXiv:2409.15594.

542 Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan
543 Sung, Denny Zhou, Quoc Le, and Thang Luong. Freshllms: Refreshing large language models
544 with search engine augmentation. *arXiv preprint arXiv:2310.03214*, 2023.

545 Peng Wang, Songshuo Lu, Yaohua Tang, Sijie Yan, Wei Xia, and Yuanjun Xiong. A full-duplex
546 speech dialogue scheme based on large language models. In *Advances in Neural Information*
547 *Processing Systems (NeurIPS)*, 2024a. arXiv:2405.19487.

548 Xiong Wang, Yangze Li, Chaoyou Fu, Lei Xie, Ke Li, Xing Sun, and Long Ma. Freeze-Omni:
549 A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv preprint*
550 *arXiv:2411.00774*, 2024b.

551 Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, et al. MMLU-Pro: A more robust and
552 challenging multi-task language understanding benchmark. In *Advances in Neural Information*
553 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2024c. arXiv:2406.01574.

554 Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese,
555 John Schulman, and William Fedus. Measuring short-form factuality in large language models.
556 *arXiv preprint arXiv:2411.04368*, 2024. SimpleQA.

557 Donghang Wu, Haoyang Zhang, Chen Chen, Tianyu Zhang, Fei Tian, Xuerui Yang, Gang Yu, Hexin
558 Liu, Nana Hou, Yuchen Hu, and Eng Siong Chng. Chronological thinking in full-duplex spoken
559 dialogue language models. In *Proceedings of the 27th Annual Meeting of the Special Interest
560 Group on Discourse and Dialogue (SIGDIAL)*, 2026. arXiv:2510.05150.

561 Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, et al. Qwen2.5-
562 Omni technical report. *arXiv preprint arXiv:2503.20215*, 2025.

563 Wang Xu, Shuo Wang, Weilin Zhao, Xu Han, Yukun Yan, Yudi Zhang, Zhe Tao, Zhiyuan Liu, and
564 Wanxiang Che. Enabling real-time conversations with minimal training costs. *arXiv preprint
565 arXiv:2409.11727*, 2024.

566 Leyang Xue, Yao Fu, Luo Mai, and Mahesh K. Marina. HybridServe: Efficient serving of large AI
567 models with confidence-based cascade routing. *arXiv preprint arXiv:2505.12566*, 2025.

568 An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

569 Wenyi Yu, Siyin Wang, Xiaoyu Yang, Xianzhao Chen, Xiaohai Tian, Jun Zhang, Guangzhi Sun,
570 Lu Lu, Yuxuan Wang, and Chao Zhang. SALMONN-omni: A codec-free LLM for full-duplex
571 speech understanding and generation. *arXiv preprint arXiv:2411.18138*, 2024.

572 Eric Zelikman, Georges Harik, Yijia Shao, Varuna Jayasiri, Nick Haber, and Noah D. Goodman.
573 Quiet-STaR: Language models can teach themselves to think before speaking. *arXiv preprint
574 arXiv:2403.09629*, 2024.

575 Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong,
576 and Jie Tang. GLM-4-Voice: Towards intelligent and human-like end-to-end spoken chatbot.
577 *arXiv preprint arXiv:2412.02612*, 2024.

578 Dengjia Zhang, Alexander Martin, William Jurayj, Kenton Murray, Benjamin Van Durme, and Reno
579 Kriz. Unified multimodal uncertain inference. *arXiv preprint arXiv:2604.08701*, 2026a.

580 Haoyang Zhang, Jun Chen, Donghang Wu, Yuxin Li, Yuxin Zhang, Xiangyu Tony Zhang, Che Liu,
581 Qingjian Lin, Yizhou Peng, Hexin Liu, Eng Siong Chng, Chao Yan, Boyong Wu, Yechang Huang,
582 Xuerui Yang, and Fei Tian. DuplexSLA: A full-duplex spoken language model with synchronized
583 speech, language, and action. *arXiv preprint arXiv:2605.20755*, 2026b.

584 Xinrong Zhang, Yingfa Chen, Shengding Hu, Xu Han, Zihang Xu, Yuanwei Xu, Weilin Zhao,
585 Maosong Sun, and Zhiyuan Liu. Beyond the turn-based game: Enabling real-time conversations
586 with duplex models. In *Proceedings of EMNLP*, 2024. arXiv:2406.15718.

587 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang,
588 Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica.
589 Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information
590 Processing Systems (NeurIPS), Datasets and Benchmarks*, 2023. arXiv:2306.05685.

591 Wenhao Zheng, Yixiao Chen, Weitong Zhang, Souvik Kundu, Yun Li, Zhengzhong Liu, Eric P.
592 Xing, Hongyi Wang, and Huaxiu Yao. CITER: Collaborative inference for efficient large language
593 model decoding with token-level routing. In *Conference on Language Modeling (COLM)*, 2025.
594 arXiv:2502.01976.

595 A Model roster

Table 2: Model roster. “Pair” groups a fine-tune with its raw backbone control; the comparison statistic is always the within-pair difference.

model	type	pair	role
MiniCPM-o 4.5	omni + duplex FT	Qwen3-8B	system target
Qwen3-8B	raw	—	pair control
MiniCPM-o 2.6	duplex FT	Qwen2.5-7B	duplex replication
Qwen2.5-Omni-7B	streaming omni FT	Qwen2.5-7B	omni control
Qwen2.5-7B	raw	—	pair control
Qwen2-Audio-7B	audio FT	Qwen2-7B	audio-FT control
Freeze-Omni	frozen LLM + adapter	Qwen2-7B	adapter control
Qwen2-7B	raw	—	pair control
Mistral-7B-v0.3	raw	—	backbone diversity
GLM4-9B-chat	raw	—	backbone diversity

596 B Extended signal and readout analyses

Table 3: Signals on MiniCPM-o 4.5, judged failure labels. In-dist = calibration-split AUC (probes 5-fold out-of-fold); LOPO math = leave-one-pool-out AUC on held-out math; audio = the same readout under speech input (final/mid-layer probe: audio-input LOPO math / text-only calibration read on audio states). Full per-signal detail: Appendix B.

signal	in-dist	LOPO math	audio input
max entropy@4	0.696	≈chance on raw backbones	—
$p(\text{True})$ -pre	0.807	trap 0.945 (probe: 0.232)	collapses
$p(\text{True})$ -post (full draft)	0.899	—	degrades
final-layer probe	0.822	0.372 (inverts)	0.936 (intact)
mid-layer (L22) probe	0.866	0.931	0.855 (text-calibrated)

597 **The in-distribution shortcut, quantified.** The numbers behind §3.1’s claim that the final-layer
598 probe’s in-distribution AUC of 0.822 is mostly pool identity: a linear read of the same hidden state
599 classifies a query’s source pool at 95.8% accuracy; a pool-identity oracle alone reaches AUC 0.715;
600 and appending true pool labels to the probe’s features moves it 0.822→0.821. In-distribution area
601 over random is statistically on par between the gate and the pool oracle (Appendix E); the two
602 separate only under LOPO, where the probe inverts to 0.372 and the oracle is undefined.

603 **ROC curves.** Figure 4 shows the calibration-split ROC for the signals of Table 3.

604 $p(\text{True})$ **backbone dependence.** Across backbones, post-draft $p(\text{True})$ beats the trained final-
605 layer probe on four of five raw/duplex models tested; the exception (GLM4-9B-chat, 0.685 vs. probe
606 0.798) shows the pattern depends on the backbone’s self-evaluation quality.

607 **Matched pairs.** Table 4: with label coverage matched by subsampling the raw control to the
608 duplex model’s per-pool fail rates, the raw backbone keeps LOPO math at 0.825/0.809 while duplex
609 models sit at chance — label coverage is ruled out; the representation changed. Degradation is
610 ordered raw > omni-streaming > duplex. Table 5 gives the by-depth summary behind Figure 2:
611 the within-pair deficit at the best mid-layer is only −0.03 vs. −0.59 at the final layer; mean-pooled
612 probes survive to the final layer on duplex models (LOPO math ≈0.80 at L35 on o4.5, 0.674 at L27
613 on o2.6 — 0.13–0.15 below each model’s mid-layer last-token peak, and uniformly weaker in-mix,
614 −0.058 at L22); duplex damage is severe but local, omni-streaming damage mild but diffuse.




Figure 4: ROC on the calibration split. $p(\text{True})$ -post leads but needs a full drafted answer; the probes and $p(\text{True})$ -pre are pre-answer.

Table 4: LOPO transfer (AUC, held-out pool) of the final-layer last-token probe, Qwen2.5-7B family.

LOPO pool	Qwen2.5-7B (raw)	Qwen2.5-Omni	MiniCPM-o 2.6	[o4.5 vs Qwen3]
hard-math	0.809	0.745	0.538 $\approx$ chance	0.372 (inverts)
hard-knowledge	0.680	0.613	0.526 $\approx$ chance	0.61

Modality specificity of the cliff. Under audio input the final layer does *not* invert (L35 audio LOPO math 0.936 vs. 0.366 text); a speak-mode template control on text input leaves the cliff unchanged (0.362) — the operative variable is the modality of the input context, not the output mode.

Audio collapse of verbalized introspection. Table 6. Three controlled arms pin the mechanism on MiniCPM-o 4.5: *not perception* (ASR audit WER 0.074; $p(\text{True})$ on the model’s own transcript snaps back to 0.074; the well-heard subset still collapses); *not a context prior* (irrelevant filler audio does not inflate p_{yes}); *binding* (duplicating the question as text alongside the audio fully restores introspection, 0.034). The verbal instance-check runs over text-token pathways that audio embeddings do not feed, while a linear probe reads the same information off the same forward pass at 0.93+. The failure implies its own cheap fix, confirmed twice: re-present the query as text before asking (“repeat-then-judge”).

Probe $\oplus$ repeat-then-judge fusion. The restored self-evaluation is *additive* with the probe. We stack the deployed L22 probe with $p(\text{True})$ -pre asked over the model’s own transcript (one short text prefill + a 1-token decode per turn; logistic stacker fit on the 360 calibration rows, native regime, Table 7). On the internal test split the fusion is the first significant AUC lift over the deployed probe (0.836 $\rightarrow$ 0.862; bootstrap 95% CI on Δ excludes zero), and it transfers to one external pool (TriviaQA 0.789 $\rightarrow$ 0.820); the remaining English pools move positively but within CI, and the Mandarin reasoning pool is untouched — there the self-evaluation itself carries no signal (AUC 0.586). Contrary to our a-priori expectation, the gain tracks where the self-evaluation signal is strong, not where

Table 5: LOPO math transfer by depth (last-token probes).

model	best mid-layer	final layer	shape
Qwen3-8B (raw)	0.964 (L33/36)	0.958	plateau, no cliff
MiniCPM-o 4.5 (duplex)	0.931 (L22/36)	0.366	cliff, inverts
Qwen2.5-7B (raw)	0.893 (L21/28)	0.809	mild dip
Qwen2.5-Omni (streaming)	0.794 (L16/28)	0.746	diffuse depression
MiniCPM-o 2.6 (duplex)	0.822 (L21/28)	0.540	cliff

Table 6: Pre-answer introspection under audio input. Trap p_{yes} = mean stated probability of answering correctly on a 100%-fail pool (lower is better).

model	FT type	trap p_{yes} : text $\rightarrow$ audio	audio $p(\text{True})$ -pre AUC
MiniCPM-o 4.5	duplex	0.055 $\rightarrow$ 0.556 (collapse)	0.786 (trap dead)
MiniCPM-o 2.6	duplex	0.196 $\rightarrow$ 0.345	0.491 $\approx$ chance
Qwen2.5-Omni	omni-streaming	0.279 $\rightarrow$ 0.213 (intact)	0.727
Freeze-Omni	frozen + adapter	0.165 $\rightarrow$ 0.112 (intact)	0.612 (capability caveat)

the probe is weak; a diagnostic arm asking $p(\text{True})$ -pre on the original query text scores within 0.01–0.02 of the transcript arm on every pool, so transcription quality is not the limiter.

Table 7: Probe $\oplus$ repeat-then-judge $p(\text{True})$ -pre fusion (native regime, deployed read point). Δ = fusion – probe, bootstrap 95% CI.

pool	n	probe	rtj solo	fusion	Δ [95% CI]
internal test	240	0.836	0.836	0.862	+0.026 [+0.004, +0.048]
TriviaQA	250	0.789	0.761	0.820	+0.030 [+0.000, +0.062]
WebQ	250	0.745	0.723	0.771	+0.026 [−0.007, +0.059]
Llama-Q	250	0.686	0.674	0.691	+0.006 [−0.020, +0.031]
SD-QA	200	0.803	0.742	0.808	+0.004 [−0.023, +0.032]
Reasoning-zh	202	0.640	0.586	0.618	−0.021 [−0.076, +0.032]

Real-speech validation. The matched-pair design rerun on 200 real human recordings (VoiceBench SD-QA, USA split) replicates all three audio-side findings: modality tax (0.400 $\rightarrow$ 0.450), audio overconfidence (paired p_{yes} shift +0.089; failure subset 0.415 $\rightarrow$ 0.581), and the layer structure (text-calibrated read on real-speech audio holds 0.76–0.80 at L22–L26, peak 0.800 at L25 — transfer across both content and modality).

Synthesis. End-to-end multimodal training *builds* the shared mid-layer core; duplex-style training additionally *damages* the readouts; omni-streaming gets both right; an adapter alone gets neither (probe readout text-fragile, verbal readout audio-fragile; in both cases knowledge survives and a readout breaks, and in both cases the breakage is duplex-fine-tune-specific).

The deployed probe, and which calibration each result uses. The mechanistic probe is a single-position L22 read (4096-d), fit on the 360-query calibration split; every representational result (§3, Appendices B, D) is this probe. The deployed probe adds two zero-latency pooled positions (an 8-token tail mean and the running user-audio mean; 12,288-d), refit on a calibration mix widened to 2,310 public-benchmark queries disjoint from every evaluation benchmark; a pre-registered guard held (frozen internal test AUC 0.860 $\rightarrow$ 0.879) and feature selection never saw the externals. It drives the turn-based live loop (§5.1), the frozen transfer results (§5.2), and the talker-on replay — frozen throughout, calibrated on text and speech-rendered internal data only. The serving-regime validations are the exceptions. In the concurrent regime, the same 12,288-d feature set is refit in-regime on 360 rows for the loop of Table 14 and on the full 2,310-row mix to measure the effect of calibration scale (Appendix H). In the native regime, the same feature set is refit at the native listen $\rightarrow$ speak read point; the deployed native gate is the 5,228-row official-configuration merge of

Table 17 (Appendix O). *Selection history*. L22 and the three-position feature set were chosen on internal calibration-split cross-validation before any external benchmark was scored (the 8-position $\times$ layer sweep and the 19-configuration feature sweep are both calibration-only), and the internal frozen test split was the pre-registered guard. The external pools were then read once with the frozen artifact. Two honest exposures remain: the layer band was picked on internal LOPO rather than by nested selection inside each held-out pool, and the second duplex family’s mid-band peak (L30–34 of 56) was located with its own sweep, so both layer choices are best read as *confirmed on held-out pools*, not as untouched-holdout estimates. In the mid-layer band the in-mix tradeoff optimum is broad (L20–L30) and the mid-layer probe’s in-distribution area advantage is modest — its real advantage is LOPO robustness, which an in-mix test cannot show.

C Second duplex family: NemotronLabs-VoiceChat-11B

The pre-registered prediction test: with the identical recipe (same audio, same judge, 600-query calibration), the layer sweep on this hybrid-Mamba duplex model (56 blocks: 27 Mamba-2 / 4 attention) peaks mid-band at L30–34 (eot-last AUC 0.714 vs. 0.693/0.682 at the endpoints); stacking the same three positions lifts OOF 0.714 $\rightarrow$ 0.761 $\rightarrow$ 0.790; the frozen-methodology probe transfers to the external pools at AUC 0.781/0.793/0.754/0.701 — the deployed-probe band reached from a 600-query calibration set. Boundary notes: the backbone is much weaker at knowledge retrieval (striviaqa local 0.32 vs. MiniCPM 0.62, same judge) and English-only (the cross-lingual pool has no analog). The live-loop port and the trained-checkpoint ablation remain future work; the port is gated by the released inference path, not by the architecture. The only supported loading path (the NeMo Speech nemotron-labs-voicechat branch) runs the duplex loop *cacheless* — every 80 ms frame re-runs the full prefix, an $O(T^2)$ schedule that cannot hold the frame clock in real time — although the hybrid backbone’s Mamba-2 state is exactly the $O(1)$ per-frame recurrence a streaming port would need; wiring stateful inference into the released duplex frame protocol is the engineering gap. (Offline, the same property is a convenience: the final frame’s forward contains every position, so one hook captures the whole eot window.) Accuracy-side conclusions do not wait on the port: the re-mix arithmetic below tracks measured live arms at 0.011 mean deviation on MiniCPM, inside the ± 0.02 –0.03 replication floor.

Escalation re-mix on every transfer pool. Completing the grid of Table 1 on this family: an offline re-mix (the same top- r reconstruction arithmetic that, on MiniCPM, tracks the measured live arms at 0.011 mean deviation) sends the top- r of each pool by NVDA probe score to the measured GPT-5.5 outcome and keeps NVDA’s own answer for the rest, scored with each pool’s official protocol (the OAB judge for the three OpenAudioBench pools — NVDA local floors 0.352/0.376/0.705 — our judge for SD-QA, VoiceBench 1–5 for AlpacaEval; the expert answers the heard transcript and is judged directly, so the relay-back channel is excluded — including the measured relay outcomes where they exist shifts tier points by ≤ 0.032). **Selective escalation beats matched-rate random on all five pools** (permutation $p \leq 0.0004$ at the 30% and 50% budgets, $p < 0.03$ at 15%; Table 8, Figure 5) — including Web Questions, SD-QA and AlpacaEval, the three pools that are flat or negative on MiniCPM (§5.2). Selectivity is nearly unchanged (AUC 0.72–0.78 under official labels); what changed is headroom — floors of 0.31–0.38 against MiniCPM’s 0.51–0.66 — and on AlpacaEval the failure species: the terse English-only model fails open-ended queries *hard* (the probe-selected half scores 3.01 vs. 4.08 kept), where MiniCPM’s open-ended failures are soft quality gradations a wrongness detector cannot see. The negatives are properties of the pool $\times$ backbone pairing, not of the signal.

D What breaks at the output: mechanism audit

§3.2 reports *that* the final-layer last-token read inverts. This section reports what we could and could not establish about *why*. Everything below uses the Phase-5d all-layer prefill dumps (every decoder layer, last-token and mean-pooled, same 600 queries) and the identical probe protocol as the layer

Table 8: The frozen-recipe NVDA probe re-mixed at fixed escalation rates, official judge per pool, each against its own matched-rate random reference (same measured outcomes, escalation set drawn at random; identical at 0% and 100% by construction). AlpacaEval rows are VoiceBench 1–5 scores; others accuracy. Every pool clears random at all three budgets: $p \leq 0.0004$ at 30% and 50%, $p < 0.03$ at 15% (20k permutations, `scripts/18_nvda_random_check.py`).

pool	judge	n	0%	15%	30%	50%	100%
Speech TriviaQA matched random	OAB	247	0.356	0.498 0.446	0.615 0.536	0.741 0.656	0.955
Speech Web Questions matched random	OAB	250	0.376	0.480 0.444	0.588 0.512	0.696 0.602	0.828
Llama Questions matched random	OAB	234	0.705	0.786 0.739	0.838 0.773	0.876 0.818	0.932
SD-QA matched random	ours	200	0.310	0.425 0.397	0.545 0.484	0.675 0.600	0.890
AlpacaEval (1–5) matched random	VB	199	3.55	3.83 3.75	4.13 3.96	4.46 4.23	4.92




Figure 5: Re-mix curves for the frozen-recipe NVDA probe: selective (top- r by probe score) vs. matched-rate random on the five transfer pools, official judge per pool. Every pool separates, including the three that are flat on MiniCPM (all five pools, $p \leq 0.0004$ at 30/50% and $p < 0.03$ at 15%, permutation).

706 sweep: calibration split, L_2 logistic regression, LOPO trained on the four non-math pools and scored
707 on held-out math. The reproduction check recovers the published L22 and final-layer numbers to
708 ≤ 0.012 (`scripts/14-17_*.py`).

709 **Four mechanisms ruled out.** Table 9 lists them. (i) *Speak mode*. Prefilling the same text queries
710 under the TTS/speak-mode template leaves the cliff exactly where it was (final-layer LOPO math
711 0.362 vs. 0.366 plain), so the readout is not being displaced by a prepare-to-speak state. (ii) *Modality*
712 *axis*. The mean audio-minus-text direction at the final layer is orthogonal to both the inverting final-
713 layer probe and the transferring L22 probe (cosine 0.013 and 0.027; chance is $1/\sqrt{4096} = 0.016$), so
714 the cliff is not the audio/text split that §3.3 measures. (iii) *More shortcut available late*. Source-pool
715 identity is linearly decodable at 0.88–0.96 at *every* depth, in the duplex model and its raw backbone
716 alike — the query-type shortcut of §3.1 is not something the late layers add; it is what remains once




Figure 6: Second-family test: (a) the usable read again peaks mid-network; (b) frozen-recipe transfer lands in the deployed-probe AUC band.

717 the competence signal stops being readable. (iv) *Wholesale representational drift*. Raw linear CKA
 718 between the duplex model and its backbone appears to collapse late (0.80→0.47), but transformer
 719 residual streams carry a few massive-activation coordinates that dominate the Gram matrix; with
 720 per-feature standardization the same comparison is nearly flat (0.82→0.78) and does not localize
 721 the cliff. Projecting out the massive-activation axes, or the pool-mean directions, rescues nothing
 722 (0.35–0.40 at every $k \leq 20$, against a random-direction control of 0.37).

Table 9: Mechanism candidates for the final-layer inversion. Final-layer LOPO hard-math AUC on MiniCPM-o 4.5 unless noted; 0.366 is the unmodified baseline and 0.931 the L22 read.

candidate	test	verdict
prepare-to-speak state	speech-mode template prefill	0.362, no change
modality (audio vs. text) axis	cosine with the probe direction	0.013 $\approx$ chance
late layers encode more pool identity	pool decoding by depth	0.88–0.96 at all depths
representation rewritten late	standardized CKA vs. backbone	0.82→0.78, flat
massive-activation coordinates	project out top- k such axes	0.35–0.40, no rescue
query-type directions	project out pool-mean directions	0.34–0.45, no rescue
delivery to the output position	residual component of the read	0.93→0.27 (Fig. 7)

723 **What does separate the models.** Figure 7 measures the probe rather than operating on the model.
 724 For the probe trained at each depth on the four non-math pools, we split its held-out score into the
 725 part carried by that layer’s five dominant principal directions (estimated from training rows only)
 726 and the residual, $s = w^\top Px + w^\top (I - P)x$. The residual component is the distributed read, and
 727 it is the one that transfers: it holds 0.85–0.93 through mid-network on MiniCPM-o 4.5 and then
 728 decays to **0.27** at the output, with the same decay on MiniCPM-o 2.6 (0.81→0.52) and *no* decay
 729 on either raw backbone (Qwen3-8B 0.86, Qwen2.5-7B 0.71 at their final layers) or on the omni-
 730 streaming control (Qwen2.5-Omni holds 0.75 to its output with a flat variance share) — the decay
 731 tracks duplex fine-tuning specifically, not streaming or audio fine-tuning in general. On MiniCPM-
 732 o 4.5 the score also concentrates: the dominant subspace carries 39% of the held-out score variance
 733 at the final layer against ≤ 0.18 at every earlier depth and 0.05 for its backbone; on MiniCPM-o 2.6
 734 this concentration is present but weak (0.16 vs. 0.14 for its backbone), so we report it as a property
 735 of the stronger duplex model, not a general law.

736 **Direct turn-control tests.** Two experiments target the turn-control story head-on rather than by ex-
 737 clusion (modal_interp.py, scripts/19_*.py). *Control-token logit lens*: applying each model’s
 738 own final norm and unembedding to the stored per-layer states, the probability mass on the trained-in
 739 control vocabulary (added/special tokens) *rises* over MiniCPM-o 4.5’s last layers (log-mass −23.8
 740 at L32 to −11.8 at the output) while the raw backbone’s falls by fifteen log-units over the same span
 741 — the duplex output layer, uniquely, stays ready to emit control tokens. But the fine-grained predic-

tion fails: the inverting probe direction is no more aligned with the control tokens' unembedding rows than with random vocabulary rows (max $|\cos|$ 0.069 vs. 0.068), and per-query control mass correlates with the probe score at only $r=0.18$. *Listen/speak intervention*: re-prefilling 60 calibration queries with a still-listening vs. answer-now suffix moves the final-layer state substantially in *both* models (relative displacement 0.27 duplex, 0.59 raw — cue text is just prompt semantics), and the cue direction is orthogonal to the inverting probe ($|\cos|=0.019$, chance 0.016), shifting its score by only $+0.23$ sd (and the deployed L22 read by $+0.08$ sd). Together: the duplex output layer is measurably control-biased in its output distribution, but the inversion is not carried by control-token directions or by an explicit listen/speak axis — the simple takeover story fails its own direct tests.

What we do not claim. Destructive tests do not settle the mechanism. Removing the top five principal directions from the final layer does raise LOPO math from 0.378 to 0.758 — but the identical operation costs the raw backbone $0.90 \rightarrow 0.14$ and costs the intact L22 read $0.931 \rightarrow 0.760$, so it is a damaging intervention that happens to help a broken readout, not evidence of a duplex-added subspace. We therefore state the finding at the level the evidence supports: duplex fine-tuning does not erase the competence signal, it stops delivering it in distributed form to the final last-token position, and the layer where delivery fails is also the one whose output distribution stays control-biased. Attributing the loss to a specific duplex objective needs training-time access we do not have; the practical consequence — read mid-network — does not depend on the attribution.




Figure 7: Non-destructive decomposition of the probe's held-out score at every depth. **Left:** share of score variance carried by the layer's five dominant directions. **Right:** transfer AUC of the residual (distributed) component — intact to the output on both raw backbones and on the omni-streaming control, decaying to inversion over the final layers of both duplex models.

760 E Query-surface router audit

A RouteLLM-style TF-IDF $\rightarrow$ LR router under identical labels reaches OOF AUC 0.669 (below the pool oracle's 0.715); a 24-configuration sweep tops out at 0.689, so this is the query surface's ceiling, not tuning. Under LOPO it collapses to 0.38–0.57 and scores the 100%-fail trap pool at 0.230 — it misses the pool a router most needs. The training receipt: at the argmax threshold its out-of-fold accuracy exactly matches the majority class; the internal probes under the same accounting clear majority by 9–29 points (Figure 9). Nor is it data-starved: the identical recipe on RouterBench [Hu et al., 2024] (36,497 prompts) reaches only in-domain AUC 0.710 with leave-one-benchmark-out at chance; RouteLLM's released preference-trained routers score 0.523/0.533 zero-shot on our pool and rank the trap pool below ordinary hard pools; our router transfers to RouterBench below chance (0.440). On audio labels the router improves (test 0.814) but still trails the mid-layer probe (0.879) and cannot fire mid-stream. In distribution, the gate's area over random ($+0.054$) is statistically on par with the pool oracle's ($+0.042$; paired bootstrap delta CI $[-0.003, +0.027]$, $p=0.13$) — the gate's case is transfer, not the in-distribution margin. $p(\text{True})$ tradeoff variants: Figure 10.




Figure 8: Accuracy vs. escalation rate, frozen test split: gated escalation vs. random-escalation and pool-oracle (type-only) baselines between the small-only and big-only endpoints.




Figure 9: Identical training receipt for the external router and the two internal probes: the router never clears the majority-class baseline; both probes do, by 9–29 points.

774 F Gate timing details

775 **Against the model’s own thinking mode.** MiniCPM-o 4.5 ships a built-in reasoning mode —
 776 self-escalation in place. Gated-cloud escalation dominates all-THINK on both accuracy and latency
 777 at every escalation rate (gated-cloud@33%: 78.7% at mean 6.5 s vs. all-THINK 63.7% at 22.4 s; Ta-
 778 ble 10). The reason is mechanistic: the gate flags knowledge and trap failures, which extra compute
 779 cannot fix — so external escalation is necessary, not merely better.

Table 10: Deployment policies, frozen test split ($n=240$); latencies measured.

policy	acc (%)	lat. mean (s)	P50	P95
fast-only	58.8	3.0	3.5	4.2
all-THINK (built-in)	63.7	22.4	17.2	60.1
gated-think @33%	61.3	12.1	3.6	47.6
gated-cloud @15%	68.8	5.3	3.6	10.1
gated-cloud @33%	78.7	6.5	3.6	20.8
gated-cloud @50%	85.8	7.0	3.6	24.4

780 **Measurement protocol.** Gate latency is wall-clock from prefill start to the L22 score being avail-
 781 able (truncated forward + probe dot product), CUDA-synchronized, one H100, 50 queries per




Figure 10: Offline tradeoff for the $p(\text{True})$ signals (pre-answer and post-draft).

modality, three warmup calls excluded, interleaved per query with the reference arms; P50/P95: gate 20/25 ms text and 45/104 ms audio vs. full-prefill TTFT 36/47 and 68/144 ms. Network and expert-RPC initiation are excluded: the cloud call fires asynchronously at the gate, and its user-facing cost is reported as expert-content-audible latency in Appendix G. Absolute prefill wall-times vary with call overhead across benches; the robust cross-bench statistic is the ratio — the L22 state is ready at $\sim 57\text{--}61\%$ of prefill.

Per-layer truncated-forward cost vs. quality (Figure 11): the L22 hidden state is available at $\sim 57\text{--}61\%$ of prefill wall-time while L22 is simultaneously the in-mix quality peak; forking much earlier is both weaker and modality-specific, so the fork belongs at $\sim 55\text{--}60\%$ depth. If the cloud call launches at the gate, the expert result arrives before the local answer finishes for 40–58% of the whole mix but only $\sim 20\text{--}31\%$ of *escalated* queries (they skew to short local answers); the residual wait is P50 2–3 s, P90 ~ 27 s (Figure 12) — the stall-phrase budget the injection protocol must cover.




Figure 11: Fork-depth Pareto: per-layer wall-time vs. gate quality.

G Live-loop details

What the signal detects. The pre-answer probe is specifically a *retrieval-failure* detector, strongest on missing-knowledge and long-tail-trap failures where the empty lookup is observable in the model’s state. Execution failures (stuck mid-derivation) surface in *decode-time* signals instead, and metacognitive failures — false-premise questions, where no failure event occurs at all — are invisible to every signal we tested, query-surface routers included (fire rate 18% vs. trap’s 90%; Appendix G), which motivates a decode-time check behind the pre-answer gate.




Figure 12: Cloud round-trip vs. local answer duration: probability the expert result arrives before the local answer ends.

801 **Full live tables.** Table 11 (both scoring views, speakable subset and full pool) and Table 12 (measured response latency). The conservative tier *shortens* the P95 tail: its few escalations displace exactly the longest local decodes. Figure 13 joins accuracy and latency; Figure 14 replays three real sessions from recorded timers.

805 **Time to first audible audio.** The sweep above is text-out; a TTS-on re-run of the balanced arm (same 240 queries, talker head + vocoder on, stall pre-synthesized once in the talker’s own voice) measures speech onset directly. The gate decision is unchanged — the TTS token enters the context only after the end-of-turn read (read p50 21 ms / p99 32 ms), and the same 84/240 queries fire. From gate decision to first audible audio: **local answers p50 0.552 s / p95 0.599 / p99 0.644; escalated turns p50 0.022 s / p99 0.029** — the canned stall (3.2–4.1 s of speech) plays the moment the gate fires, so the escalated path reaches first audio 25× faster. Expert *content* becomes audible at p50 6.4 s / p95 26.0 / p99 34.0 (cache-corrected expert round-trip + first relay PCM, p50 0.66 s after the answer returns); per-turn dead air between stall end and expert content is p50 2.5 s / p95 22.4 / p99 30.7, with 25/84 escalated turns fully covered — the stall covers the head of that wait, and Table 12’s completion profile bounds the tail.

Table 11: Live streaming sweep, paired bootstrap 95% CIs ($B=10^4$). Top: speakable subset ($n=218$); bottom: full pool ($n=240$). Gold-inject = escalated rows re-scored with gold-text expert answers; channel cost = paired difference. * CI excludes 0; † CI lower bound at zero.

tier	esc	deployed	Δ vs. floor	gold-inject	channel cost
never	0%	0.422 [0.36,0.49]	—	0.422	—
conservative	14%	0.528 [0.46,0.59]	+0.106*	0.550	+0.023†
balanced	31%	0.592 [0.53,0.66]	+0.170*	0.661	+0.069*
aggressive	57%	0.633 [0.57,0.70]	+0.211*	0.761	+0.128*
always (measured)	100%	0.665 [0.60,0.73]	+0.243	0.922 [0.89,0.95]	+0.257
never (full)	0%	0.383 [0.32,0.45]	—	0.383	—
conservative	16%	0.483 [0.42,0.55]	+0.100*	0.525	+0.042*
balanced	35%	0.554 [0.49,0.62]	+0.171*	0.654	+0.100*
aggressive	61%	0.596 [0.53,0.66]	+0.212*	0.771	+0.175*
always (measured)	100%	0.617	+0.234	0.917 [0.88,0.95]	+0.300

816 **Speculative prefetch, measured and dropped.** Framed as speculative escalation (draft–verify at sentence granularity), its acceptance rate — does an expert answer to the partial transcript answer the full question — is 0.07/0.19/0.51 at 25/50/75% of the audio; at the mid-stream gate’s typical fire

Table 12: Measured response latency (s), $n=240$ per arm (latency is channel-independent). P99 on $n=240$ is indicative only.

tier	mean	P50	P95	P99	Δ mean	Δ P50	Δ P95	Δ P99
never	4.6	2.0	15.8	17.8	—	—	—	—
conservative	4.6	2.6	12.8	23.0	+0.0	+0.6	-3.0	+5.2
balanced	5.8	3.5	15.9	32.7	+1.2	+1.5	+0.1	+14.9
aggressive	6.6	4.4	18.5	33.0	+2.0	+2.4	+2.7	+15.2




Figure 13: Accuracy vs. per-arm P50 latency (speakable subset), both scoring views, with the simulated-P50 random reference.

point 6–8 of 10 calls would be wasted. The trap pool’s floor (0.00/0.10/0.27) shares the back-loaded-core property that breaks mid-stream probe reads.

Conflicting-injection controls (172 sessions). Fabricated conflicting expert answers by within-pool derangement, four framings: worst-case comply is 0.00 neutral / 0.53 deployed authority / 0.79 strong-override, while *seeding* words into the talker’s mouth backfires (comply 0.25, lip-service 0.62 — forced speech is treated as revisable). Math is silently re-derived rather than relayed or disputed. Model-wrong $\times$ wrong-injection swallow rate is 0.64: expert quality is the accuracy ceiling of the cascade.

Uplink diagnostics. On the escalated population the expert answering the talker’s own transcript scores 0.585 vs. 0.871 on gold text. Prompt-level rescues on the self-transcript are no-ops (the damaging errors are internally-coherent entity substitutions); replacing the ears helps: open-weights Whisper recovers little (+4 pp) but a stronger cloud ASR on the same audio recovers 38% of the gap (0.585→0.694, McNemar $p=0.007$) — diagnostic only (its critical-path latency and external-pool survival are unmeasured). A same-model gold-text control with an audio-capable expert shows the audio channel itself is lossless for speakable content: the leak is the talker’s transcription, not the synthesis or the relay (relay drops: 5 cases of 108; corrupted transcripts dominate the escalated failures, first-generation sweep).

Boundary query classes. *False premises:* the model fails substantially (0.63 text / 0.47 audio) and the expert would help (0.80), but the gate is blind (fire@balanced 0.18 vs. trap’s 0.90; no score separation) — answering along a false premise does not feel like inability. *Real-time queries* (“what is NVDA trading at”): the gate is not blind, the class was simply absent from every static training

Three live sessions, timestamped (balanced arm; every event time is a recorded live-loop timer)




Figure 14: Three real sessions from recorded timers: expert beats the stall phrase (seamless); expert outlives it (dead air, the P99 signature); gate stays closed (local answer fully decoded 0.4 s after end of turn). Speech occupancy is estimated at 150 wpm (hatched); the measured arms are text-out, so no vocoder onset is plotted.

family; extending calibration with FreshQA [Vu et al., 2023] (fast-changing labeled escalate a priori, never-changing as in-family controls) raises held-out fast-changing fire rates to 0.80/0.93/1.0 across tiers with frozen-test AUC unchanged (0.877 vs. 0.879). One calibration subtlety: tier budgets must remain quantiles of the deployment-mix scores, or the new capability is spent on its own threshold shift.

Demonstration system. The pipeline runs as an interactive web demo on one H100, on the model’s *native full-duplex* head (Appendix O): every second of microphone audio is prefilled into the context the talker is generating from, the head itself decides listen/speak per chunk, and the gate reads L22 at the listen→speak transition against the in-regime thresholds. Escalated turns play a canned stall (teacher-forced once, the talker’s own voice), consult the expert (with a demo-only web-search tool for real-time queries), and the verified answer re-enters the stream as a text unit that the talker voices — an ordinary duplex turn, hence interruptible like any other. There is no VAD, no ASR barge-in classifier, and no abort machinery: floor control is entirely the duplex head’s own behavior (Appendix O). An earlier demo build approximated barge-in with an energy-VAD + burst-ASR harness over the turn-based loop; it is retired, and survives only as the harness-overlap arm measured above. The demo will be made available upon publication.

H Duplex-validation sweep: the gate under the talker

The headless live sweep (§5.1) never engages the talker head; this sweep replays its full 240-query pool through the deployed voice loop above to test whether the frozen probe read survives the spoken regime. Two arms, one session per query, deployed v4 gate with scores logged but escalation disabled (the arm isolates gate validity, not routing outcomes). **Clean:** the query streams as mic-

shaped PCM (pool TTS audio plus steady noise at RMS 0.008), the energy VAD ends the turn, the talker answers aloud, and the session is cut after the first audible chunk. **Overlap:** an easy warmup question (probe score <0.25 , locally answered, >5 s spoken answer) puts the talker mid-answer, and the target query is then spoken over it; sustained overlapping speech commits a barge-in whose audio seeds the next turn. The barge-in path engaged on 237/239 pairs.

Against the frozen local-failure labels of the same pool, the identical frozen probe reads: headless AUC 0.866 [0.817, 0.909], clean 0.871 [0.824, 0.914], overlap 0.856 [0.803, 0.901]; barged-only subset ($n=237$) 0.854. Paired bootstrap deltas vs. headless are $+0.005$ [-0.025 , $+0.035$] (clean) and -0.009 [-0.036 , $+0.016$] (overlap) — both inside the ± 0.02 – 0.03 live replication floor. Per-query scores correlate across regimes at $r=0.90$ – 0.92 , and the balanced threshold makes the same fire/hold decision as the headless read on 0.887 (clean) and 0.912 (overlap) of queries. The end-of-turn read costs 24–26 ms at p50 ($p95 \leq 30$ ms) in both arms, with the first audible answer chunk arriving ≥ 0.5 s after the gate decision ($p50 \approx 0.58$ s), so the decision precedes any audible commitment. Scope: this loop keeps per-turn session control — the mic stays open while the talker speaks and interruptions seed the next read, but incoming audio is not prefilled *during* generation. In hindsight this sweep and the concurrent arm are best read as *read-point ablations*: the deployed native full-duplex regime (Appendix O) sits strictly between the turn-based ideal and the harness-concurrent worst case on every pool measured. The concurrent regime is measured separately below.

Floor control under escalation: barge-in vs. backchannel. The demonstration loop’s floor policy (speech over the talker ducks playback; a fast ASR pass returns the floor on backchannels or commits the interrupt) raises a distinct worry: did adding the gate change the talker’s overlap behavior? A dedicated sweep injects overlap stimuli into the deployed loop at real-time pacing — four classes: short backchannels (0.40–0.46 s, the English+Chinese lexicon), long lexicon-only continuers (time-stretched to 1.33–2.08 s, probing the sustained-speech commit rule), out-of-lexicon stop commands, and pool queries spoken over the answer — in matched pairs with the gate on and off under an identical floor policy (non-firing queries, $n=40$ per class per arm). The answer is exact orthogonality: all 160 pairs make the identical duck/resume/interrupt decision in both arms (every paired rate delta 0.000 [0.000, 0.000]; paired latency medians within 31 ms), down to the single seeding failure, which occurs on the same query in both arms. By construction the gate reads once at end of turn and never enters the floor state machine; this measures it. The three phases that exist only under escalation — the stall sentence, the expert wait, the relay speech — are the actual new risk surface and behave (Table 13): a genuine barge-in aborts the escalated turn and seeds the next turn 40/40, and short backchannels hold the floor at the same rate as during local answers. One asymmetry is disclosed: a false interrupt during the expert wait cancels the pending expert call, so the policy’s fail-toward-interrupt bias is recoverable but not free there. That residual backchannel false-interrupt rate (0.40 short, 0.50 long; identical in both arms) is a property of the harness policy and is per-stimulus deterministic: non-lexical hums (“mm-hm” and its Chinese counterpart) die by the ASR-empty fail-closed rule (16/16 each; worded tokens such as “okay” 0/16), and continuers sustaining loud speech past the ≥ 1.2 s commit rule die without any ASR check (20/20; pause-bearing variants 0/20). We keep the policy as deployed for measurement integrity and note both rules as the obvious fixes. For scope: the checkpoint’s *native* duplex profile, measured separately on Full-Duplex-Bench v1.0 [Lin et al., 2025] (727 samples, official scripts), shows best-in-class pause handling (turn-over rate 0.125/0.117 Candor/synthetic), competitive user-interruption response (0.915 at 0.90 s), and *zero* spontaneous backchannels — the lexicon floor is a harness-level approximation added precisely because the talker head has no native backchannel behavior to preserve.

Concurrent prefill: interleaved listen/speak in one KV stream. A third arm drives the streaming API in time division: while the talker is mid-answer, each 1 s chunk of the target query is prefilled into the *same* session between generation chunks, so the end-of-turn read fires while generation is still active (238/240 queries, full interleave coverage). The carrier answer is teacher-forced: under sampling the duplex model reliably yields the turn — it emits EOS within ~ 2 s of user audio appearing in its context, the trained turn-taking behavior surfacing under interleave. This is the one shift on the transfer ladder that breaks the frozen read: scores drift up (mean 0.48 $\rightarrow$ 0.71;

overlap injected during	backchannel holds floor	barge-in aborts	seeds next turn
local answer (gate off)	0.60 [0.45, 0.75]	40/40	39/40
local answer (gate on)	0.60 [0.45, 0.75]	40/40	39/40
stall sentence	0.63 [0.38, 0.88]	16/16	16/16
expert wait	0.60 [0.33, 0.87]	14/14	14/14
relay speech	0.55 [0.27, 0.82]	10/10	10/10

Table 13: Floor-control sweep over the deployed voice loop. Short backchannels (lexicon bursts) vs. genuine barge-ins (pool queries spoken over the talker), by the phase the overlap lands in; the last three phases exist only under escalation. Interrupt latency p50 1.39–1.40 s from overlap onset in every phase; escalated-turn aborts discard the pending expert answer and the barge-in speech becomes the next turn.

the frozen balanced threshold’s fire rate jumps 37→75%, AUC drops 0.869→0.758 [0.693, 0.816] (paired Δ AUC -0.109 [$-0.158, -0.064$]), and label-free requantiling recovers decision agreement only to 75%. The information, however, survives in the state: refitting the *same* 12288-d linear read on 360 calibration queries collected in-regime reaches **0.818** [0.761, 0.871] on the held-out 240 (calibration OOF 0.775) — more than half the gap recovered from those 360 in-regime rows alone. The mid-layer competence direction persists under concurrent listen/speak; unlike every earlier shift (modality, external pools, talker-on), calibration must follow the regime. Run-to-run replication of the concurrent read correlates at 0.895. Caveat: the interleave goes through the public streaming API — turn-role headers enter mid-generation and the carrier text is forced — so this measures the concurrent *state*, not the official duplex serving format.

The full loop in the concurrent regime. With the in-regime gate (thresholds = label-free quantiles of each pool’s own concurrent never-arm scores), the complete escalation loop — carrier speaking, target audio interleaved, EOT read mid-generation, stall → expert → relay — was re-run end to end on all seven pools (Table 14). We report this sweep as *exploratory*: one session per query per arm, an in-regime probe calibrated on 360 internal rows only, teacher-forced carrier (the caveat above). The loop’s mechanics reproduce — realized escalation rates track nominal on the external pools (10.8–51.2%), and latency holds (EOT read p50 23 ms, canned-stall onset 27–34 ms, local first audio p50 .655 s vs. .552 turn-based, the price of the longer concurrent context) — and on the internal pool, where calibration matches the regime, so does *selection*: accuracy climbs 40.4→52.1→66.2% (speakable **44.5→56.0→71.1%**) over the balanced and aggressive tiers, clearing the matched-rate random remix of the same never/always outcomes ($p=.017$ and 5×10^{-5} ; speakable .009 and 5×10^{-5}), and *selective escalation beats always-escalate* (66.2 vs. 61.7% at two thirds of the calls; speakable 71.1 vs. 66.5) — the turn-based signature result carries over. The matched-random rows of Table 14 bound any stronger reading, in three ways. (i) Gains are not monotone everywhere: the internal conservative tier under-fires (2.5% realized vs. 15.8% nominal — the top calibration quantile did not transfer to the regime; balanced and aggressive did) and lands flat (40.4→40.0, 44.5→43.6, inside the ± 2 –3 floor). (ii) Externally the 360-row in-regime probe separates from matched-rate random only in spots (Llama @50%: 90.8 vs. 86.5, $p=3\times 10^{-4}$; Reasoning zh @30/50%: $p \leq .0024$; AlpacaEval @15%: $p=.031$); at the other external operating points the budget gains are what random escalation would buy — consistent with this probe’s external AUC of .57–.67 against the turn-based probe’s .68–.81 (internally, where calibration matches the distribution, it reads .857 vs. the turn-based .879). We then tested the calibration-scale attribution directly: the full 2,310-row calibration mix was re-collected in the concurrent state and the same read refit at full scale. Scale is real — every English pool lifts (paired against the 360-row probe on the same features: TriviaQA .645→.699, WebQ .641→.713, Llama Q .674→.755, SD-QA .639→.701; mean Δ AUC $+.068$ [.036, .099]) along a monotone data-scaling curve (external mean .625→.689 over 360→2,310 rows; internal .818 unchanged) — but it buys only about a third of the distance to the turn-based .785–.806, and Reasoning zh, absent from the English calibration mix, moves within noise (.615→.577 [$-.117, .037$]). The residue is the regime itself, not calibration scale. Table 14’s loop ran with the 360-row probe and its numbers are unchanged. (iii) On Llama Questions selective-

953 vs-always is judge- and noise-sensitive (official judge 90.8 vs. 91.2; ours 86.4 vs. 85.6 — both deltas
 954 inside the floor), so we claim that comparison only on the internal pool. Floors move by pool — the
 955 regime cost story: Llama $\approx$ unchanged, TriviaQA -16 , Reasoning zh -18 (the English carrier hurts
 956 the zh pool most), WebQ $+4.8$ under the official judge (within the $\pm 2-3$ replication floor; the larger
 957 our-judge shift is alias-matching style sensitivity). AlpacaEval rises $4.36 \rightarrow 4.60$ (1–5 scale, always
 958 4.76), but only its conservative tier clears matched random ($p=.031$; 4.40–4.55 across tiers), so the
 959 turn-based honest-negative stands.

Table 14: The full escalation loop in the concurrent regime (*exploratory*: one session per arm, 360-row in-regime calibration, teacher-forced carrier), all pools, judge-matched to Table 1 (OAB official for TriviaQA/WebQ/Llama; ours for SD-QA/Reasoning/internal; VoiceBench 1–5 for AlpacaEval). Gate = in-regime probe, per-pool label-free quantile thresholds; tier columns are headed by each pool’s *realized* rate, and the *rand* row under each pool is the matched-rate random remix of the same never/always outcomes (20k draws; $p = P(\text{random} \geq \text{gate})$, marked * $p < .05$, ** $p < .01$). Accuracies in %; AUC = in-regime probe vs. never-arm failure, decimals. Script: scripts/20_conclive_random_check.py.

pool	never	cons.	bal.	aggr.	always	AUC
TriviaQA	50.4	54.0	64.0	71.6	91.6	0.636
<i>rand</i> (10.8/30.8/46.8%)		54.9	63.1	69.6		
WebQ	61.6	65.2	68.8	71.2	81.6	0.636
<i>rand</i> (15.6/30.4/48.8%)		64.7	67.7	71.4		
Llama Q.	81.6	83.2	84.8	90.8**	91.2	0.672
<i>rand</i> (14.8/29.2/51.2%)		83.0	84.4	86.5		
SD-QA	49.5	55.5	60.5	70.5	89.5	0.636
<i>rand</i> (15.0/29.5/50.5%)		55.5	61.3	69.7		
Reason. zh	40.6	48.5	58.4**	64.4**	80.7	0.570
<i>rand</i> (15.3/30.2/45.0%)		46.8	52.7	58.7		
internal (240)	40.4	40.0	52.1*	66.2**	61.7	—
<i>rand</i> (2.5/36.2/66.2%)		41.0	48.1	54.5		
internal speakable (218)	44.5	43.6	56.0**	71.1**	66.5	0.857
<i>rand</i> (2.3/31.2/62.8%)		45.0	51.4	58.3		
AlpacaEval (1–5)	4.36	4.45*	4.53	4.60	4.76	—
<i>rand</i> (10.1/30.7/47.7%)		4.40	4.48	4.55		

960 **A serendipitous recovery behavior.** An accidental run surfaced one bonus behavior worth report-
 961 ing: when a long question contains a >1.25 s internal pause, the endpointer fires early and the talker
 962 starts answering the fragment — but the question’s own continuation then triggers barge-in, and
 963 the remainder becomes the next turn, where the gate simply re-scores it. Early end-pointing thus
 964 degrades into an extra turn rather than a lost query — the honest duplex-ish recovery available to a
 965 turn-based loop.

966 I Fixed-threshold transfer and the expert-benefit oracle

967 Both analyses below run in the *gold-inject* view (escalated rows take the expert’s answer to the gold
 968 text). The deployed-channel view cannot evaluate a freely moving threshold: the most permissive
 969 arm escalates 50%, so the other half of each pool has no expert outcome on record, whereas the
 970 ceiling runs cover every id. Labels follow each pool’s own judge (Table 1); the two judges disagree
 971 by up to 13 points on Web Questions, so they are never mixed within a row. AUC columns are
 972 recomputed from the deployed sweep’s stored end-of-turn scores; they differ from Table 1’s AUC
 973 column — captured in a separate scoring pass — by up to 0.03 externally and 0.04 on the internal
 974 pool (0.840 here vs. 0.879 there): the passes render and prefill the audio independently, and the
 975 probe score carries render-to-render jitter. AlpacaEval is excluded — it has no binary correctness
 976 label on either side, only a 1–5 score whose expert ceiling is degenerate (195/199 at 5).

977 **One absolute threshold, six pools.** The deployed gate adapts per pool, but only through label-free
 978 quantiles of its own score distribution. To separate what the probe’s *ranking* carries from what its

979 *scale* carries, we freeze the single absolute threshold the internal system actually deployed ($\tau=0.628$,
 980 the balanced tier) and apply that number verbatim everywhere. **Scale does not transfer:** the same
 981 τ realizes escalation rates from 2.0% (Reasoning QA) to 48.0% (SD-QA), a 24-fold spread, so a
 982 fixed cut cannot hold an escalation budget across pools. **Ranking does:** at whatever rate it fires,
 983 fixed- τ escalation beats a random reference at the same realized rate on four of five external pools,
 984 paired-bootstrap interval excluding zero (Table 15). The exception is Reasoning QA, where τ fires
 985 on 2% of queries and there is almost nothing to measure. This is the concrete case for the label-free
 986 quantile step: it buys budget control, not discrimination.

987 **Does the gate find failures the expert can fix?** The probe is fit and read against *local failure*, but
 988 what a routing system needs is *expert benefit*: $y=1$ when the local model is wrong *and* the expert is
 989 right. Re-scoring the same frozen scores against this harder label costs little on the retrieval pools
 990 (Speech TriviaQA 0.796 $\rightarrow$ 0.782, Llama Questions 0.830 $\rightarrow$ 0.813) and more where the expert is it-
 991 self weak (Web Questions 0.773 $\rightarrow$ 0.689 against a 0.864 expert ceiling; internal 0.840 $\rightarrow$ 0.732). The
 992 gate thus remains a usable benefit predictor out of distribution (0.63–0.81) despite never having seen
 993 an expert outcome, and the drop is largest exactly where the paper already reports a flat live curve.
 994 On Web Questions the mechanism is direct: 75.0% of local failures are expert-fixable pool-wide, but
 995 only 65.5% of the failures inside the gate’s top-scoring 15% are, and that share climbs monotonically
 996 as the cut widens (0.655/0.685/0.711 at 15/30/50%) — the gate’s most confident failure calls are the
 997 ones the expert also misses. A benefit-trained refit closes none of this: with expert outcomes col-
 998 lected for all 2,310 calibration queries (gold-text expert, expert-right rate .853), retraining the same
 999 12,288-d read directly on the benefit label moves internal benefit-AUC .732 $\rightarrow$.759 (paired +.027
 1000 [.004, .052]) and leaves every external pool inside noise (Δ $-.029$ to $+.037$, all CIs spanning zero),
 1001 at no internal cost on the fail label (.840 $\rightarrow$.839). The benefit gap reflects the label’s dependence
 1002 on the expert’s own failure surface, not a trainable objective mismatch — the expert-agnostic label
 1003 leaves nothing on the table. Table 16 places the deployed gate against the matched-budget oracle
 1004 that escalates true benefit cases first: the gate clears random at every pool and budget, capturing a
 1005 fifth to three fifths of the random-to-oracle span, and the residual gap is the headroom a decode-time
 1006 or expert-aware second stage would have to claim.

Table 15: One absolute threshold ($\tau=0.628$, frozen on the internal pool) applied verbatim to every pool; gold-inject view. Δ is fixed- τ accuracy minus a random reference at the *same realized rate* (paired bootstrap over ids, $B=10^4$, seed 42). AUC columns re-score the same frozen scores against local failure and against expert benefit ($y=$ local wrong $\wedge$ expert right, base rate in the last column). Rates and accuracies in %; AUC in decimals.

pool	n	rate	local	acc	rand	Δ 95% CI	AUC _{fail}	AUC _{benefit}	base
Speech TriviaQA	250	28.4	66.4	83.2	75.0	[+4.4, +12.1]	0.796	0.782	30.8
Speech Web Questions	250	36.0	56.8	72.0	67.5	[+0.8, +8.4]	0.773	0.689	32.4
Llama Questions	250	8.0	83.6	86.8	84.3	[+0.6, +4.6]	0.830	0.813	11.2
SD-QA	200	48.0	51.5	80.5	71.4	[+4.5, +13.6]	0.800	0.765	43.5
Reasoning QA (zh)	202	2.0	58.9	60.9	59.5	[−0.1, +3.5]	0.679	0.634	30.2
internal frozen test	240	35.0	38.3	65.4	57.0	[+3.8, +13.1]	0.840	0.732	53.3

1007 J Transfer latency shapes

1008 **Matched-rate precision.** Precision at a fixed escalation rate is capped at base-fail/rate: on Llama
 1009 Questions (base fail 0.16) the 50% cap is 0.32 and the deployed probe reaches 0.304 — 95% of the
 1010 cap (recall 0.93). Quoting raw precision would misread the strongest pool as the weakest.

1011 **The median left-fold.** On easy pools per-arm P50 is non-monotonic (Llama Questions
 1012 1.52 $\rightarrow$ 1.17 $\rightarrow$ 1.60/1.42 s): a verified median-of-a-mixture mechanism, not noise. The probe pref-
 1013 erentially escalates the queries whose *local* decode is longest — on retrieval pools wrong answers
 1014 hedge at length (decode time and answer length $r=0.97$), so a wrongness-selecting probe strips the

Table 16: Oracle frontier at matched escalation budgets (gold-inject view): random reference / deployed quantile gate / benefit oracle, which escalates the true $y=1$ cases first. The gate clears random everywhere; the oracle column is the remaining headroom. Accuracies in %.

pool	15% budget			30% budget			50% budget		
	rand	gate	oracle	rand	gate	oracle	rand	gate	oracle
Speech TriviaQA	71.0	76.0	81.6	75.5	83.6	96.4	81.6	91.2	97.2
Speech Web Questions	61.2	64.0	72.0	65.7	70.0	86.8	71.6	76.4	89.2
Llama Questions	85.0	88.8	94.8	86.4	91.2	94.8	88.2	93.2	94.8
SD-QA	57.7	63.5	66.5	63.9	74.5	81.5	72.2	82.0	95.0
Reasoning QA (zh)	63.1	67.8	73.8	67.4	72.8	89.1	73.0	76.7	89.1
internal frozen test	46.3	49.2	53.3	54.3	60.8	68.3	65.0	74.2	88.3

slow tail as a side effect (escalated-at-15% ids: never-arm local P50 1.711 s vs. 1.198 s kept; local accuracy 0.32 vs. 0.73). Controls: the probe does not key on length ($\text{corr}(\text{score}, \text{length}) = +0.05$; ≈ 0 conditioned on wrongness), and on SD-QA — where wrong answers are short — the same probe selects wrongness identically and produces no fold. The mean is monotonic (1.65→2.06 s); kept-local P50 falls 1.52→0.43 s across tiers (Figure 15). On Reasoning QA the shape inverts usefully: the local chain-of-thought is *spoken*, so every reasoning token enters the response clock, while the expert’s chain is hidden behind a flat ~ 3.7 s round-trip — escalation truncates the latency tail (P90 13.25→10.94 s) rather than fattening it.




Figure 15: Mixture decomposition of Llama Questions latency per arm: kept-local P50 falls 1.52→0.43 s as the probe strips the slowest local decodes while escalated rows pay a flat ~ 3 s expert round-trip; the arm median (dashed) is their mixture.

AlpacaEval scale caveats. The official 4.8 is an offline chat-mode number — our chat-mode control reproduces it at 4.86 — and the chat-vs-live gap (4.86 vs. 3.99) is the duplex loop’s spoken-reply style plus a token cap, so only live arms are comparable to each other. Selected queries score 3.90 on the never arm vs. 3.98 for skipped: no discrimination; escalated rows gain (3.90→4.71) because the expert writes better long-form answers — random picks would harvest the same gain.

K Full-pool and per-family figures

Figures 16 (dual-view) on the full frozen pool, and the complete external-family figure set (Appendix L); the noise audit behind the ± 0.02 –0.03 floor is Figure 17.

L External-benchmark figure set




Figure 16: Dual-view live tradeoff, full pool ($n=240$); companion to Figure 3.




Figure 17: Replication-noise audit: same query, same audio, re-run across arms. Kept-local verdicts flip 2.3–18.8% by pool; repeated escalation 0.7–10.6%; implied paired SE 0.009–0.028.









Figure 18: Speech TriviaQA, Speech Web Questions, SD-QA: dual-view (left) and accuracy–latency (right).




Figure 19: Llama Questions, Reasoning QA (zh), VoiceBench AlpacaEval: dual-view (left) and accuracy–latency (right).

1032 M Judge and self-evaluation prompts

1033 All adequacy labels use the following fixed system prompt (structured outputs with a strict {reason,
1034 adequate} schema; the judge never sees which system produced the answer):

```
1035     You are a strict, fair grader of a small language model’s answers.  
1036     You are given a QUESTION, an optional REFERENCE answer, and a  
1037     CANDIDATE answer produced by a small model. Decide whether the  
1038     CANDIDATE is ADEQUATE, meaning ALL of: (1) factually correct;  
1039     (2) any reasoning shown is valid; (3) it actually and completely  
1040     answers the question that was asked. For multiple-choice questions  
1041     the candidate must land on the correct option (by letter or by  
1042     unambiguously stating the option’s content). A partially correct,  
1043     hedged, refused, off-topic, or truncated answer is NOT adequate. If  
1044     a REFERENCE is provided, grade against it. Be strict but do not  
1045     penalize harmless differences in format or extra correct detail.
```

1046 The verbalized self-evaluation baselines use one fixed phrasing each; $P(\text{Yes})$ is read from the first
1047 generated token’s distribution:

```
1048     [pre] I am going to show you a question. Do NOT answer it. Judge  
1049     honestly whether you yourself would answer it correctly. Question:  
1050     {q} Would you answer this question correctly? Reply with exactly one  
1051     word: Yes or No.  
1052     [post] Here is a question and a proposed answer. Question: {q}  
1053     Proposed answer: {a} Is the proposed answer correct? Reply with  
1054     exactly one word: Yes or No.
```

1055 N Reproducibility

1056 All sampling and splits use a fixed seed (42); the calibration/test split is frozen once and the test
1057 split is touched only by the evaluation scripts. Environment pins: torch 2.8.0; transformers 4.51.0
1058 for MiniCPM-o 4.5, with per-model pins for comparison backbones; all experiments run on single-
1059 H100 instances. Probe artifacts (weights, feature configuration, per-pool label-free thresholds), the
1060 judge prompt above, all trace files, and the per-experiment execution log — including per-run GPU
1061 and API spend, on the order of \$800 total — ship in the project repository, released upon publication.

1062 O The native full-duplex regime

1063 Every result so far reads the probe under a *turn-based* serving loop (harness-controlled end-of-turn)
1064 or its harness-interleaved concurrent variant (Appendix H). MiniCPM-o 4.5 also ships a native full-
1065 duplex mode: audio prefills in 1 s units into the same context the model generates from, and the
1066 head itself emits ⟨listen⟩ or speaks, yielding the floor via its turn-end token. This is the deployed
1067 regime of the demonstration system, and this section recalibrates and re-validates the gate inside
1068 it. The read point needs no external end-of-turn detector: the chunk where the head first switches
1069 listen→speak *is* the commit-to-answer moment, and the probe reads there (the tail- k features then
1070 cover the model’s first ~ 8 answer tokens — a generation-side read the head hands us for free). For
1071 87% of test queries this commit lands within ± 1 chunk (± 1 s) of the true question end.

1072 **Recalibration.** The full 2,310-row calibration mix was re-collected in this regime (fresh duplex
1073 session per query, features at the deployed read point, labels unchanged) and the same 12,288-d
1074 linear read refit. (The final deployed gate goes one step further and also re-labels: each training
1075 row’s label is the judged outcome of the native answer produced in the same duplex session the fea-
1076 tures were read from, under the deployed sampling, rather than a separate deterministic turn-based
1077 answer. The two label sets agree on 80.8% of the 5,228 rows; the native failure rate is higher, .541
1078 vs. .463, mostly turn-based-correct questions the sampled native answer gets wrong. Ranking is
1079 unchanged within noise, Table 17.) On native test-240 features the in-regime probe reaches AUC

Table 17: Native-regime probe AUC at the deployed listen→speak read point, under the OFFICIAL serving configuration (final deployed system; the chronological narrative above measured under inherited wrapper defaults). The merged refit adds the 2,700-row public-benchmark expansion (incl. 400 Mandarin rows) to the official core. En-4 mean .770 vs. turn-based .771 — the English regime cost is eliminated; the Mandarin slice restores Reasoning-zh from the all-English chance collapse. The last column is the deployed gate: same 5,228 features, but every training label is the judged outcome of the *same* native-duplex answer the features were read from (instead of a separate deterministic turn-based answer). Label source is a wash for ranking (En-4 mean .771 either way; the paired bootstrap crosses zero on six of seven pools, TriviaQA +.035 [+0.006, +.065] being the exception), so the deployed gate uses the in-regime definition. Internal test is scored against both the earlier harness labels and the native-judged labels of the same 240 queries.

pool	<i>n</i>	old cfg	official (2,542)	+exp3 (5,228)	+native lbl. (deployed)
internal (harness lbl.)	239	.830	.846	.844	.843
internal (native lbl.)	239	—	—	.878	.876
TriviaQA	250	.711	.759	.767	.802
WebQ	241	.736	.754	.772	.777
Llama-QA	245	.757	.789	.815	.807
SD-QA	200	.736	.737	.728	.700
Reasoning-zh	202	.606	.495	.605	.621

Table 18: Native-regime gated accuracy under the official serving config and the deployed merged (official+exp3) gate (8ad remix: native scores select, native local outcomes + cached always-arm expert outcomes mix). *p* = permutation vs. matched-rate random; escalation rate in parentheses. Thresholds are the deployed per-language operating points: tier quantiles of the training-set out-of-fold scores, computed separately for the English and the Mandarin training rows (no evaluation pool is consulted); the Mandarin pool uses the zh thresholds. On frozen test the aggressive tier reaches the always-escalate ceiling at 49% of the expert cost. Reasoning-zh, which the global English thresholds never reached (1–2% fire), now fires at its nominal rates and beats matched-rate random at the two lower tiers. AlpacaEval (measured under the 2,542-row official gate) stays a clean honest negative.

pool	local	balanced	aggressive	always
frozen test	.431	.565 (27%, <i>p</i> =.0005)	.669 (49%, <i>p</i> <.0001)	.669
TriviaQA	.568	.684 (20%, <i>p</i> =.0005)	.852 (54%, <i>p</i> <.0001)	.916
WebQ	.411	.527 (22%, <i>p</i> =.045)	.680 (58%, <i>p</i> =.020)	.809
Llama-QA	.767	.792 (4%, <i>p</i> =.002)	.841 (16%, <i>p</i> <.0001)	.910
SD-QA	.455	.605 (24%, <i>p</i> =.0035)	.775 (61%, <i>p</i> =.003)	.895
Reasoning-zh (zh thr.)	.525	.658 (34%, <i>p</i> =.023)	.713 (55%, <i>p</i> =.066)	.807
AlpacaEval (VB 1–5)	3.65	3.70 (6%, n.s.)	4.07 (37%, <i>p</i> =.24)	4.76

1080 **0.830** [0.777, 0.878] (transferred concurrent probes: 0.807–0.818; turn-based reference 0.877). Ex-
1081 ternally the in-regime refit beats both transferred probes on all five pools — TriviaQA .711 (Δ+.080
1082 [+0.030, +.129]), WebQ .736, Llama-QA .757 (Δ+.118), SD-QA .736, Reasoning-zh .606 (Δ+.115);
1083 external mean .709. The regime ordering is monotone everywhere: turn-based (.877/.771 internal/external mean) > native (.830/.709) > concurrent (.689 external) — i.e. a substantial share of
1084 the concurrent regime’s cost was the harness read point, not liveness itself. Thresholds again do *not*
1085 transfer (balanced quantile 0.386→0.645); scores are decoding-temperature invariant (0.7-trained
1086 probe on temp-0.1 features: AUC 0.843).
1087

1088 **Validity.** Re-mixing native local outcomes (gpt-judged, deployed decoding) with cached always-
1089 arm expert outcomes under the native scores (Table 18): under the deployed gate, gated accuracy
1090 clears matched-rate random on all five English pools at both deployed tiers — frozen .431→.669 at
1091 49% escalation (*p*<.0001; ceiling .669), TriviaQA .568→.852 at 54%, WebQ .411→.680, Llama-
1092 QA .767→.841, SD-QA .455→.775 at 61%. In-regime training renormalizes the score distribution,
1093 so realized fire rates sit near their nominal quantiles on the English pools (the earlier 64–73% drift
1094 was a transferred-threshold artifact). Reasoning-zh is an operating-point problem before it is a rank-
1095 ing problem: a single global threshold, quantiled on a 92%-English training mix, places the zh score
1096 distribution below every tier (1–2% fire at “balanced”), the language axis of Appendix H expressed

1097 as a fire-rate failure. The deployed fix needs no evaluation data: tier thresholds are quantiled on the
1098 training-set out-of-fold scores *per language* (en/zh), and the session selects the language’s thresh-
1099 olds. The zh thresholds derived from the 400 Mandarin training rows alone land Reasoning-zh at
1100 17/34/55% fire (nominal 15/30/50) and lift it from .525 to .609/.658/.713, beating matched-rate ran-
1101 dom at the two lower tiers ($p=.0065$, $p=.023$). This static per-language point is only as good as
1102 the match between the language slice’s family mix and the deployment stream: when the Mandarin
1103 training slice is widened from 400 to 1,900 rows with higher-failure families (below), its balanced
1104 quantile moves from .39 to .75 and fires 5% on Reasoning-zh; a family-balanced quantile recovers
1105 13/24/38%. The composition-free alternative — each pool’s own score quantile, or its online form,
1106 a 100-turn sliding-window tracker with no labels and no pool identity — realizes nominal rates on
1107 every pool and converges within $\pm .06$ of nominal on the English pools.

1108 **Is the Mandarin gap a language gap?** No: it is a task-type gap. A targeted 1,500-row Mandarin
1109 expansion (Chinese SimpleQA long-tail facts, C-Eval and CMMLU knowledge MC, the remain-
1110 ing XCOPA-zh; public benchmarks, deduplicated against every pool, judged on the native answers;
1111 native failure .578) does not move Reasoning-zh when added to the calibration mix (.621→.627,
1112 paired $+ .006$ [$-.026$, $+ .040$]; internal test $-.010$, English external $+ .006$), and Mandarin-only train-
1113 ing is worse everywhere (Reasoning-zh .591, internal .758). Read the other way, the deployed gate
1114 scores these 1,500 unseen Mandarin questions at AUC **.811** [.789, .833] (long-tail .719, C-Eval .730,
1115 CMMLU .642, causal .693) — the same band as the English pools. What the probe does not read
1116 is multi-step reasoning, which is what Reasoning-QA is; the “.62 on Mandarin” number is that gap
1117 wearing a language label. The expansion is therefore not deployed and serves as a second, larger
1118 Mandarin held-out pool. (A side observation from the same dump: on the causal-commonsense
1119 family 34% of the native answers switch to English while remaining judged adequate — a language-
1120 consistency defect the adequacy label does not see.)

1121 **Reasoning coverage and in-context rows: the read point saturates.** Two further tar-
1122 getted sets test the two remaining hypotheses. A bilingual multi-step reasoning pool
1123 (1,319 rows: LogiQA and LogiQA-zh, MATH levels 4–5 without LaTeX, StrategyQA,
1124 BBH date/deduction/causal/temporal/navigation, CMATH grades 4–6, Ape210K; native failure
1125 .595) leaves Reasoning-zh unchanged when added to calibration (.621→.620, paired $-.002$
1126 [$-.039$, $+ .038$]) and moves nothing else beyond noise. The reason is visible inside the pool itself:
1127 cross-validated AUC on the reasoning families is only .61–.76 (overall .736 [.708, .762]), against
1128 .77–.82 for fact, knowledge-MC, and causal families. Whether a multi-step chain will come out
1129 right is weakly encoded in the listen→speak commit state, which precedes the chain; the .62 on
1130 Reasoning-zh is that ceiling, not a coverage gap. Second, the calibration questions re-collected
1131 as the *second* turn of a session (a spoken carrier question and its answer first) fail more often
1132 (.517→.620; 26% of labels flip) while the deployed gate scores them *lower* (mean shift $-.097$):
1133 at the balanced tier the in-context fire rate halves (.252→.115) although ranking holds (AUC .861
1134 [.815, .896]). Adding these rows to training does not change AUC either; the defect is the operating
1135 point, and the deployed demo therefore thresholds on a sliding-window quantile of the session’s own
1136 recent onset scores (100-score window, static per-language threshold until 20 scores exist), which
1137 follows such shifts without labels or pool identity. After a label-definition fix and three targeted
1138 expansions ($\approx 4,000$ rows), external ranking sits where the 5,228-row gate left it. The head is not
1139 the bottleneck either: on the same features, per-dimension standardization costs .07 of English ex-
1140 ternal AUC (the raw activation scale is informative), an elastic-net head is a wash (internal $+ .008$,
1141 external $-.007$), and non-linear heads are clearly worse (256-unit MLP .702, PCA-256 + gradient-
1142 boosted trees .715, vs. .771). Nor are several heads: a three-way task-type router trained on the
1143 same features is 97% accurate, yet no family-specialized head beats the single head on its own
1144 family (facts .846 vs. .847, reasoning .761 vs. .766, chat .672 vs. .720), and the routed mixture is
1145 flat-to-negative externally (SD-QA $-.033$ [$-.060$, $-.007$]). What the commit-chunk state encodes
1146 about the coming answer is extracted by one linear read; the remaining headroom lies in *what* is read
1147 — later into the answer, other layers, or a second signal family such as verbalized self-evaluation
1148 (Table 7) — not in how. The talker’s own local floor is lower in-regime (frozen .483→.431 under

the official serving configuration; the inherited-defaults harness measured .371, and the $\approx$ one-third-temperature / $\approx$ 8-point decomposition belongs to that older figure), which *raises* the gate’s marginal value in deployment. The in-regime calibration scaling curve (measured under the default serving configuration) mirrors the concurrent one and is not saturated: internal .717 $\rightarrow$.830, external mean .643 $\rightarrow$.709 from 360 $\rightarrow$ 2,310, and a further public-benchmark expansion to 5,009 rows (2,300 English across the same seven families + 400 Mandarin: MGSM-zh and XCOPA-zh through the same TTS pipeline, source-disjoint from Reasoning-zh) continues the $\approx$ +.02-per-doubling external trend to **.736** while internal stays flat (.834) — and the Mandarin slice moves the axis English data never touched, Reasoning-zh .606 $\rightarrow$.659 (Δ vs. conc-2310 +.169 [+ .085,+.251]; the deployed official-configuration merge reads the same pool at .605 with turn-based labels and .621 with the deployed native labels, Table 17). Two caveats on reading this curve: its six points are the project’s chronological accumulation, each step adding new query families, coverage, or a serving-configuration change, not repeated random subsamples of the final set, so it evidences that added *coverage* transfers, not that another random doubling of the same families would return +.02; and the internal/English-external plateau against the turn-based reference (.771) is an empirical baseline at $n\approx$ 200–250 per pool, not a ceiling.

Ceiling, stability, and the pre-answer read. Three diagnostics on the deployed 5,228-row gate. (i) *Ceiling*. A pool-prior scorer (each query scored by its training pool’s out-of-fold failure rate) reaches AUC .786 against the probe’s .848 on the training out-of-fold scores, and .767 against .876 on frozen test; AUC restricted to same-pool pairs is .743 (train) and .820 (test). Per-question discrimination is thus +.06–.11 of AUC on top of the type prior, the same decomposition as the $n=600$ audit of Section 3.1; at a 30% budget the prior alone recovers .475 failure recall against the probe’s .506 (random .299). The external pools, each a single source, measure the within-pool component directly (.70–.81, Table 17). (ii) *Stability*. An independent second native pass over frozen test (deployed sampling, $T=0.7$, its own judged label) changes 98% of answer texts and flips 14% of native labels, yet the read at the commit chunk barely moves: feature cosine .95, score correlation .95, and the tier decisions agree on 91–94% of queries. Each run’s features predict the *other* run’s outcome about as well as their own (.848/.887 vs. .876/.872), so the probe reads a property of the question in context rather than the fate of one sampled answer; against the 204 queries whose two outcomes agree, AUC is .91–.93, which bounds the single-sample label noise in the reported numbers. (iii) *Pre-answer read*. The same layer-22 read taken before the onset chunk’s generate call (no answer tokens in the tail), scored by the deployed head without refitting, gives AUC .861/.837 against the two label sets, .011–.015 below the deployed read; the first $\approx$ 8 answer tokens are worth about one AUC point, and the pre-answer variant needs its own operating points (the deployed quantiles fire 0/1/20% on it).

Live validation. Full live runs on frozen test-240 (real raw-audio ASR uplink, real expert, wait paced at one chunk per second) validate the remix arithmetic where the channel is thin and expose where it is not: balanced delivers .504 at 23% fire (remix predicted .492 at 24%); aggressive delivers .537 at 45% against a .596 remix prediction, and the gap decomposes entirely onto fired turns — live fired-turn accuracy .495–.589 vs. the .667 cached-expert ceiling. The native *relay channel* (the answer re-entering as a text unit and being voiced) is the regime’s lossy element — 89% of relay units need the nudge retry — playing the role the transcription uplink played in the turn-based loop; the uplink itself is now clean (raw-audio ASR, p50 0.9 s). Latency, expert-cache warm: expert p50 4.4–5.7 s, wait p50 5–6 chunks, fire $\rightarrow$ relay first audio p50 $\sim$ 7–8 s. Selection is visible live: unfired subsets score .478/.573 against the .371 unconditional local floor.

Dialogue-act gating. Full duplex creates commits that are not answers: after a barge-in or a backchannel the head speaks to *manage the floor* (“Okay.”), and the failure probe is out-of-distribution there — on 194 TTS’d floor-management stims (stop commands, backchannels, acknowledgments, fillers; en+zh) the unmodified gate false-fires on 37% at the aggressive tier (stop commands 45%), i.e. it would escalate “stop” to the expert. The fix is a second linear head on the *same* L22 read — info-seeking vs. floor-management — which separates the two acts perfectly

1200 (OOF AUC 1.000); thresholded at the question-side 0.5th percentile it costs 0.5% of true escalations
1201 and admits 0% of floor stims. Escalation requires *both* heads: $\text{act} \geq \tau_a \wedge P(\text{fail}) \geq \tau$. The de-
1202 ployed demo answers floor turns naturally and reports the act score per commit; the failure probe’s
1203 calibration is untouched.

1204 Live use then sharpened both heads within minutes — scripted pools miss what ten minutes of real
1205 conversation finds. (a) Live mic speech shifts both act distributions into the calibration gap (real
1206 request-phrased questions score $\sim .93$, live floor turns $\sim .37$ vs. a standalone-stim maximum of $.05$):
1207 a percentile-anchored τ_a rejected a real question and the talker delivered an empty promise. The fix
1208 is a gap-center threshold plus *in-context* calibration stims — each utterance dumped as the second
1209 turn of a carrier conversation, matching how floor acts actually arrive; the act read stays perfectly
1210 separable (AUC 1.000) under both conditions, and in-context stop commands false-fire the *failure*
1211 probe on 57% at the aggressive tier, so the act term is not optional. (b) The native refit had silently
1212 dropped the FreshQA real-time extension (“the stock price of NVDA today” scored $P(\text{fail})=.29$
1213 and stayed local); replaying that recipe in-regime restores fast-changing heldout fire to $.45/.75/1.00$
1214 across tiers at zero internal-AUC cost ($.830 \rightarrow .833$), with budgets still quantiled on the core mix. (c)
1215 A relay’s own speak-onset must not re-enter the gate (re-fire with an empty uplink); deliveries are
1216 guarded to their end of turn. The pattern across all three: the mid-layer signals separate cleanly
1217 every time, and what re-opens at each regime or domain shift is calibration coverage and thresholds
1218 — interactive use is itself an evaluation modality.

1219 **Multi-turn grounding.** A stateless expert uplink breaks on follow-ups: “what is NVDA trading
1220 at” then “what about Apple” sends only the second utterance, unresolvable in isolation. The probe
1221 itself is unaffected — it reads L22 with the prior turns live in the model’s KV cache — so the fix is
1222 in the uplink, not the gate: the deployed demo threads a rolling resolved-dialogue history into the
1223 expert input and asks it to resolve references, after which the follow-up correctly returns Apple’s
1224 share price. The measured eval pools are single-turn, so this is a deployment capability rather than
1225 a table; systematic multi-turn evaluation is future work.

1226 **Floor control, natively.** With the harness gone, 8ba’s question re-poses on the head itself (108
1227 scripted overlap trials, stims spoken over the talker at fixed offsets). During answer speech,
1228 backchannels almost never stop the floor (false-stop 3/24 within a 6-chunk window) with *no* ASR
1229 or lexicon anywhere in the loop; sustained question speech yields the floor (6/12 within 6 chunks,
1230 median 8.5 post-stim); short burst commands (“Stop!”) mostly do not (3/12) — the native head
1231 reads sustained speech, not commands, a naturalness-for-reliability trade the old energy-VAD did
1232 not make. The escalation machinery is floor-transparent where it should be: relays complete under
1233 overlapping stims (26/30), and backchannels during the thinker wait leave the pending relay
1234 intact (8/10). The weak stop-command sensitivity decomposed into two causes, found in sequence.
1235 *Mechanism:* instrumenting the serving loop shows the head never samples $\langle \text{listen} \rangle$ mid-turn — zero
1236 attempts across 36 trials — so its only stop channel is the turn-end token. *Configuration:* a user-
1237 demanded vanilla control (the bare official serving stack, no gate machinery) revealed that our loop
1238 had inherited the checkpoint wrapper’s sampling defaults ($k=100$) rather than the official demo’s
1239 serving config ($k=20$, a 3-chunk startup listen guard, an assistant-style system prompt). Under the
1240 official config the turn-end channel is *responsive*: a spoken “stop” two seconds into an enumeration
1241 answer ends the turn in 2.0–2.2 s in 3/3 scripted sessions, versus 1–13 s (once running the answer to
1242 completion) under the inherited defaults — wide- k sampling dilutes away the rising turn-end proba-
1243 bility. The deployed demo now serves the official config, and the earlier floor-control stop numbers
1244 measured under the inherited defaults should be read as a lower bound on native responsiveness. Two
1245 durable lessons: serving-parameter parity with the reference stack is part of “native”, and a vanilla
1246 control arm belongs in the harness from day one. The one true interaction: a *new question* during
1247 the wait derails the pending relay 7/10 times — the head simply moves on — which is the empirical
1248 case for aborting the in-flight expert call on user speech, left as deployment future work. Latency:
1249 listen chunks cost 0.04 s and speak chunks ~ 0.5 s on one H100 (realtime holds); fire $\rightarrow$ stall-audio
1250 0.2–1 s; fire $\rightarrow$ relay-first-audio is the expert wall clock (13–20 s observed) plus ~ 2.5 s.