
When Does a Small Model Know to Hand Off? Reading Competence from Full-Duplex Speech Models for Real-Time Escalation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Full-duplex speech models answer in real time with a small on-device backbone,
2 yet many queries exceed its competence. We ask whether the model’s own hid-
3 den state signals, *before the first output token*, that a query is beyond it — reli-
4 ably enough to hand off to a large cloud model without breaking the live session.
5 Standard competence readouts fail on duplex models: final-layer probes collapse
6 under distribution shift on text input (leave-one-pool-out math AUC .93→.37),
7 consistent with late-layer specialization for duplex turn control, while verbalized
8 self-assessment collapses on audio input. A mid-network linear probe survives
9 both: leave-one-pool-out math .93, text-calibrated speech deployment .86, and an
10 end-of-turn read in the streaming loop that reaches AUC .887 in ~30 ms — be-
11 fore decoding starts. Run live end-to-end, the gated duplex system climbs from
12 .422 to .633 heard accuracy at 57% escalation, with the remaining gap dominated
13 by the transcription uplink rather than the gate. The gate then transfers *frozen*
14 to six public speech benchmarks using only label-free threshold quantiles, lifting
15 Speech TriviaQA to .860 live against the model’s official offline .755. A compe-
16 tence signal sufficient for real-time escalation is already present mid-network in
17 duplex speech models; the standard readouts simply miss it.

18 1 Introduction

19 A full-duplex speech model listens and speaks in the same forward pass, under a conversational
20 deadline of a few hundred milliseconds [Stivers et al., 2009]. That deadline forces a small on-device
21 backbone — and a small backbone routinely faces queries beyond its competence: long-tail facts,
22 expert knowledge, multi-step reasoning. The natural architecture is selective escalation to a large
23 cloud model, but the duplex setting imposes two constraints that existing routing work does not
24 meet: the decision must come *before the first output token* (once the model starts speaking, a wrong
25 answer is already on the air), and the session — audio context, voice, cache — must stay with the
26 local model, which will speak the expert’s answer itself.

27 This paper asks whether the small model already contains the signal this decision needs. We study
28 MiniCPM-o 4.5 [Cui et al., 2026], a 9B omni/duplex fine-tune of Qwen3-8B, against matched-
29 pair controls spanning raw backbones, streaming-omni, audio-only, and frozen-adaptor fine-tunes,
30 and evaluate every standard zero-shot competence readout — decode-time entropy, verbalized self-
31 assessment ($p(\text{True})$), and linear probes on hidden states — with judged failure labels on a frozen
32 pool of public-benchmark queries, in matched text and speech conditions.




Figure 1: **Left:** standard competence readouts break on duplex models. **Middle:** a mid-network (L22) linear read decides in ~ 30 ms — before the first output token — so the expert call is already in flight while local decoding would just be starting. **Right:** the live system hands off mid-conversation and the *same talker speaks* the expert’s answer, improving real duplex sessions ($.422 \rightarrow .633$) and, frozen, public speech benchmarks (Speech TriviaQA $.664 \rightarrow .860$).

Everything runs on a frozen checkpoint — no fine-tuning; the only fitted component is a logistic regression on stored hidden states, calibrated once on text. Three findings, in increasing order of system-level consequence:

- Representation: duplex fine-tuning breaks the standard readouts; a mid-network signal survives.** The conventional final-layer probe looks strong in distribution (AUC .82) but is mostly a query-type shortcut: under leave-one-pool-out transfer it *inverts* on held-out math ($.93 \rightarrow .37$). The damage is specific to the late-layer last-token position on text input — absent in raw backbones and non-duplex fine-tunes — consistent with late-layer specialization for duplex turn control. Mid-network (L22 of 36) the same linear read transfers at .93, and the mid-network optimum generalizes to a second, architecturally disjoint duplex family (hybrid-Mamba NemotronLabs-VoiceChat-11B).
- Modality: the mid-layer signal transfers text $\rightarrow$ speech in end-to-end models.** A probe calibrated on cheap text data reads speech-input hidden states at .86; the control that isolates the mechanism is Freeze-Omni, an adapter on frozen identical weights, where cross-modal transfer is dead ($\leq .60$). Verbalized self-assessment shows the mirror-image failure — it collapses on audio input — so the mid-layer probe is the one readout that survives both modalities. The transfer holds on real human recordings, not just TTS.
- System: the signal is available before the first token and works live, in and out of distribution.** The L22 read costs 20 ms (text) / 45 ms (audio) against time-to-first-token of 36/68 ms; in the streaming loop an end-of-turn read reaches AUC .887 in ~ 30 ms. Run live end-to-end — expert consulted mid-conversation, answer injected, talker speaks it — heard accuracy climbs $.422 \rightarrow .633$ at 57% escalation, and the frozen gate transfers to six public speech benchmarks with label-free thresholds, lifting Speech TriviaQA to .860 against the model’s official offline .755.

Extended analyses — readout-damage controls, router baselines, escalation-channel diagnostics, latency mechanics — are in the appendix.

2 Related Work

Full-duplex real-time dialogue. End-to-end spoken dialogue spans native duplex designs — dGSLM [Nguyen et al., 2023], Moshi [Défossez et al., 2024], GLM-4-Voice [Zeng et al., 2024], single-LLM native full-duplex [Fang et al., 2026] — streaming omni models [Xu et al., 2025a, Cui et al., 2026], frozen-backbone adapters [Wang et al., 2024a], and audio fine-tunes [Chu et al., 2024]. A parallel line adds explicit turn-taking control: FlexDuo [Liao et al., 2025], MinMo [Chen et al., 2025], and Freeze-Omni’s state heads decide *when the local model should speak*, at 120–250 ms cadences that show a live loop tolerates far more decision latency than our ~ 30 ms read [cf. Riera et al., 2026, Wu et al., 2026]. Our gate decides the orthogonal question — *whether the local model*

68 *should be the one answering* — and the head-based controllers’ shared choice of a last-layer readout
69 is exactly the point our results show duplex fine-tuning damages.

70 **Does the model know what it knows?** $p(\text{True})$ self-evaluation [Kadavath et al., 2022], verbalized
71 uncertainty [Lin et al., 2022], and truthfulness probes on hidden states [Azaria and Mitchell, 2023,
72 Burns et al., 2023, Marks and Tegmark, 2024] established that correctness is readable internally;
73 intermediate layers carry more of it than the output distribution [Orgad et al., 2025]. Semantic
74 entropy [Kuhn et al., 2023, Farquhar et al., 2024] needs multiple sampled generations; Semantic
75 Entropy Probes [Kossen et al., 2024] amortize it into a hidden-state probe — the closest signal to
76 ours, but read post-hoc on text-only half-duplex models. Multimodal uncertainty fusion [Zhang
77 et al., 2026] asks a downstream question; we ask *where* in a duplex model a modality-robust read
78 survives at all. Our contribution is a controlled account of when the standard readouts break under
79 duplex fine-tuning — and where the surviving signal sits.

80 **Selective escalation and routing.** FrugalGPT [Chen et al., 2023], HybridLLM [Ding et al., 2024],
81 RouteLLM [Ong et al., 2024], and confidence-based cascade serving [Xue et al., 2025] route from
82 query-side features or preference data; RouterBench [Hu et al., 2024] and LLMRouterBench [Li
83 et al., 2026] standardize evaluation and report that many routers barely beat simple baselines. Au-
84 toMix [Aggarwal et al., 2024] is the closest system design — a small model assesses its own answer
85 and escalates — but the decision is *post-answer* self-verification; a voice deadline requires *pre-*
86 *answer*. Both routing assumptions (query-side features; the large model answers the user directly)
87 fail in full-duplex voice, where the session lives in the small model. Learning-to-defer [Madras et al.,
88 2018], early exit [Schuster et al., 2022], speculative decoding [Leviathan et al., 2023], and test-time
89 compute [DeepSeek-AI et al., 2025] make related decisions inside one model; our audit quantifies
90 the query-side ceiling directly (§4.1).

91 **The gap.** Duplex work decides when to speak; routing work decides which model, from the query
92 surface or after a first answer; uncertainty work asks whether the model is confident, almost always
93 in text and offline. No prior work we know of asks whether a full-duplex speech model exposes a
94 *pre-answer, modality-robust, internal* competence signal early enough to drive a live small-to-large
95 handoff while preserving the session — the regime this paper measures end-to-end.

96 3 Experimental Setup

97 **Target model and matched-pair controls.** The system target is **MiniCPM-o 4.5** (9B; omni +
98 full-duplex fine-tune of Qwen3-8B), bf16, greedy decoding, seed 42. To attribute effects to the *kind*
99 of fine-tune rather than the backbone, we assemble ten models spanning raw backbones, a second du-
100 plex generation (MiniCPM-o 2.6), a streaming-omni fine-tune (Qwen2.5-Omni-7B), an audio-only
101 fine-tune (Qwen2-Audio-7B), and a frozen-backbone adapter model (Freeze-Omni, LLM weights
102 byte-identical to its raw control); the full roster is Table 4 (appendix). The comparison statistic is
103 always the within-pair difference.

104 **Query pool, labels, and audio arm.** 600 queries in five pools, frozen before any gate experi-
105 ment (360 calibration / 240 test): *trap* = SimpleQA [Wei et al., 2024] (the target fails 100%), *hard-*
106 *knowledge* = MMLU-Pro [Wang et al., 2024b], *easy-fact* = TriviaQA [Joshi et al., 2017], *easy-chat*
107 = Dolly/Alpaca [Conover et al., 2023, Taori et al., 2023], *hard-math* = GSM8K tail + MATH-500
108 [Cobbe et al., 2021, Hendrycks et al., 2021, Lightman et al., 2024]; public benchmarks only, no
109 LLM-generated queries. Adequacy labels come from an LLM judge [Zheng et al., 2023] with struc-
110 tured outputs, blind to the gate and validated by a stronger re-judge (agreement .962 text / .945 audio;
111 every headline number moves $\leq .010$ under the stronger labels); the escalation expert is gpt-5.5 at
112 low reasoning effort. The audio arm renders the same frozen pool to speech (single-voice TTS, con-
113 tent unchanged — only modality varies), following Spoken-SQuAD / VoiceBench practice [Li et al.,
114 2018, Chen et al., 2024]; audio-side findings are additionally validated on real human recordings

Table 1: Signals on MiniCPM-o 4.5, judged failure labels. In-dist = calibration-split AUC (probes 5-fold out-of-fold); LOPO math = leave-one-pool-out AUC on held-out math; audio = the same readout under speech input (final/mid-layer probe: audio-input LOPO math / text-only calibration read on audio states). Full per-signal detail: Appendix B.

signal	in-dist	LOPO math	audio input
max entropy@4	.696	$\approx$ chance on raw backbones	—
$p(\text{True})$ -pre	.807	trap .945 (probe: .232)	collapses
$p(\text{True})$ -post (full draft)	.899	—	degrades
final-layer probe	.822	.372 (inverts)	.936 (intact)
mid-layer (L22) probe	.866	.931	.855 (text-calibrated)

(SD-QA, Faisal et al., 2021) and on the external benchmarks’ official audio (§7). Judge prompts, pins, and spend: Appendix J–K.

Gate signals (no model training). (i) decode scalars: max token entropy / margin over the first k decode steps; (ii) linear probes: logistic regression on a single hidden state, position $\times$ layer swept, fit on the calibration split only; (iii) verbalized self-evaluation [Kadavath et al., 2022]: $p(\text{True})$ -pre asks “would you answer this correctly?” before answering, $p(\text{True})$ -post asks about a drafted answer; $P(\text{Yes})$ is read from the first token’s distribution.

4 Reading Competence from Duplex Models

4.1 In distribution every readout looks fine; transfer reveals the shortcut

Table 1 compares the candidate signals on MiniCPM-o 4.5. In distribution the final-layer probe reaches .822 — but that number is mostly a *query-type* shortcut: the same hidden state classifies a query’s source pool at 95.8%, a pool-identity oracle alone reaches AUC .715, and appending true pool labels to the probe adds nothing (.822 $\rightarrow$.821). The decisive test is leave-one-pool-out (LOPO) transfer, which we treat as the primary metric throughout: held-out math *inverts* to .372, and the mean escalation score on a 100%-fail trap pool is .232. An external query-surface router makes the same point from the other side: under identical labels a TF-IDF router reaches only OOF AUC .669 (below the pool oracle), collapses to .38–.57 under LOPO, and 100 $\times$ the data on a public routing benchmark buys back only .04 of AUC (Appendix D). Instance-level competence information lives in the model’s internal state, not on the query surface.

Verbalized self-assessment transfers differently: *pre-answer* $p(\text{True})$ catches the trap pool the probe misses (escalation score .945 vs. the probe’s .232) — but timing is essential: asked *after* drafting, the model endorses its own confident-wrong answer (.360). Post-draft $p(\text{True})$ is the strongest single signal (.899) but requires paying for a full draft, which a voice deadline cannot afford.

4.2 The damage is duplex-specific; the signal survives mid-network

Matched pairs attribute the probe failure to the *kind* of fine-tune, not the backbone: the raw Qwen backbone, label-matched to the duplex model’s fail rates, holds LOPO math at .825 where the duplex model sits at .372; streaming-omni and audio-only fine-tunes are undamaged; a second duplex generation (MiniCPM-o 2.6) reproduces the collapse (Appendix B). A layer $\times$ position sweep localizes it (Figure 2): the collapse is specific to (late layer $\times$ last token) on text input — mean-pooled reads at the same depth are fine, and mid-network the last-token read recovers to **.931 at L22**, near the raw backbone. This is the position a streaming head must re-purpose to encode turn-control state, and the pattern is consistent with late-layer specialization for duplex turn control overwriting the readout while the underlying information survives. The mid-network optimum is not a MiniCPM quirk: on NemotronLabs-VoiceChat-11B — a hybrid-Mamba duplex model, architecturally disjoint — the same recipe peaks mid-band at L30–34 of 56 and transfers to the external pools at AUC .70–.79 with a quarter of the calibration data (Appendix C).

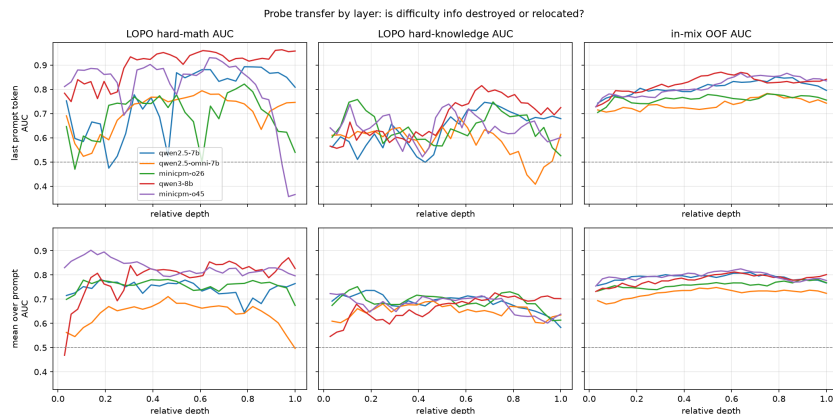


Figure 2: The core representational result: LOPO math AUC by layer and position. On the duplex model the standard (final-layer, last-token) read inverts on text input while mid-network layers hold $\geq .9$; raw backbones show no cliff. The deployed gate reads L22.

151 4.3 Text-calibrated, speech-deployed — in end-to-end models only

152 Probes trained on text hidden states and scored on speech-input hidden states locate where the
 153 representation becomes modality-shared: early layers are modality-specific (.54–.60), and from
 154 $\sim 50\%$ depth text $\rightarrow$ audio transfer plateaus at .82–.87 (.855 at the deployed L22). This replicates
 155 on MiniCPM-o 2.6 and cross-family on Qwen2.5-Omni. The decisive control is **Freeze-Omni**:
 156 identical frozen LLM weights with audio attached by an adapter — and cross-modal transfer is dead
 157 (.52–.60). The modality-shared mid-layer core is *created by training the backbone on the modality*;
 158 the cheap recipe “calibrate on text, deploy on speech” is valid for end-to-end models and invalid
 159 for adapter-attached ones. Verbalized introspection shows the mirror-image failure: on audio input,
 160 pre-answer $p(\text{True})$ collapses on both MiniCPM duplex generations while both non-duplex controls
 161 stay intact; controlled arms show the verbal self-check runs over text-token pathways that audio em-
 162 beddings do not feed — duplicating the question as text restores it — while a linear probe reads the
 163 same information off the same forward pass at .93 (Appendix B). All audio-side findings replicate
 164 on 200 real human recordings (SD-QA), ruling out TTS artifacts. **The mid-layer probe is the one**
 165 **readout that survives every training regime that can converse**; the live system deploys the same
 166 L22 read plus two zero-latency pooled positions (the *deployed* probe; Appendix B).

167 5 From Signal to Gate

168 **The decision exists before decoding starts.** The gate is one 4096-d dot product on a hidden state
 169 the forward pass already computes. Measured on H100, the L22 decision lands in **20 ms (text) /**
 170 **45 ms (audio)** against time-to-first-token of 36/68 ms: $T_{\text{gate}} < T_{\text{first token}}$, far inside the 200–300 ms
 171 turn-taking budget [Stivers et al., 2009], and the cloud call can launch while the remaining $\sim 40\%$
 172 of prefill and the whole decode still run ($p(\text{True})$ -pre costs an extra prefill, $p(\text{True})$ -post a full draft:
 173 1.9/3.5 s; per-layer cost curves in Appendix E).

174 **Accuracy–cost tradeoff.** On the frozen test split ($n=240$): small-only .588, big-only (gpt-5.5)
 175 .917. Gated escalation beats random escalation at every operating point (Figure 6, appendix); the
 176 balanced tier at a 33% budget reaches **.779–.787**, recovering $\sim 58\%$ of the small $\rightarrow$ big gap at one
 177 third of the big-only cost. In distribution the gate is statistically on par with a pool-oracle (type-
 178 only) router — as the type-shortcut audit predicts — so its case over type routing rests on transfer,
 179 where the type shortcut collapses (§4.1) and the mid-layer signal does not.

180 **Against the model’s own thinking mode.** MiniCPM-o 4.5 ships a built-in reasoning mode —
 181 self-escalation in place. Gated-cloud escalation dominates all-THINK on both accuracy and latency

Table 2: Deployment policies, frozen test split ($n=240$); latencies measured.

policy	acc	lat. mean (s)	P50	P95
fast-only	.588	3.0	3.5	4.2
all-THINK (built-in)	.637	22.4	17.2	60.1
gated-think @.33	.613	12.1	3.6	47.6
gated-cloud @.15	.688	5.3	3.6	10.1
gated-cloud @.33	.787	6.5	3.6	20.8
gated-cloud @.50	.858	7.0	3.6	24.4

at every escalation rate (gated-cloud@.33: .787 at mean 6.5 s vs. all-THINK .637 at 22.4 s; Table 2). The reason is mechanistic: the gate predominantly flags knowledge and trap failures, which extra compute cannot fix — thinking helps math and hurts chat — so external escalation is necessary, not merely better.

6 Live Duplex Escalation

The streaming loop. Audio streams in 1 s chunks; per-chunk probe scores are strictly weaker than the offline read (best chunk statistic AUC .764) and miss short utterances whose critical entity lands in the final chunk. The fix is an **end-of-turn read**: after the last audio chunk, prefill an assistant turn containing a single space and read L22 at the assistant-start position. This read reaches AUC **.887** — above the offline reference (.843) — in ~ 30 ms, still before the first output token, and is the loop’s only decision point; escalation-rate thresholds are quantiles of this single score. Speculative mid-stream prefetch was measured and dropped (an expert answer to a partial transcript answers the full question only .07–.51 of the time; Appendix F). On escalation the talker speaks a stall phrase, its *own* transcription of the audio goes to the expert (gpt-5.5, reasoning effort fixed low — higher effort buys zero accuracy on the escalated population at double the latency), and the returned answer is injected as context for the talker to relay. On the deployed channel the relay is clean: model wrong, correct answer injected, comply = 1.00 ($n=24$); more broadly the competence signal also predicts relay compliance (comply .42→.75 across gate-score quartiles) — a model that knows it doesn’t know is willing to listen. Conflicting-injection and framing controls: Appendix F.

The live result. We ran 240 frozen test queries $\times$ four escalation tiers, each a complete duplex session, scored on *heard accuracy* — what a user would actually have heard. Headline numbers use a pre-registered input-side filter (21 formula-LaTeX ids + 1 broken recording removed; spoken formulas are destroyed by any TTS render before the system hears them). Heard accuracy climbs monotonically, .422 → .528 → .592 → .633 at 0/14/31/57% escalation, every tier’s gain significant (paired bootstrap; Figure 3). The remaining gap to the channel-controlled counterfactual (escalated rows re-scored with gold-text expert answers) is dominated by *transcription-uplink* errors, not gate failure: the channel costs +.023/+0.069/+1.28 across tiers, and in the channel-controlled view the gate clears its paired random reference at every arm. Live arm accuracies on this pool carry a replication-noise floor of $\pm .02$ –.03, so no claim here rests on a sub-3-point live delta. Full tables (both views, both pools), latency profiles, uplink diagnostics, and two boundary query classes (false-premise questions, real-time questions) are in Appendix F.

7 Frozen Transfer to Public Speech Benchmarks

The out-of-distribution test the type-shortcut audit demands: we take the deployed gate *frozen* — weights, features, and layer choice fit once on the internal calibration mix — and run the complete live loop of §6 on six public speech benchmarks: Speech TriviaQA, Speech Web Questions, Llama Questions, and Reasoning QA from OpenAudioBench [Li et al., 2025] (spoken renderings of Joshi et al., 2017, Berant et al., 2013, and the set of Nachmani et al., 2024), SD-QA (real dialectal human speech) [Faisal et al., 2021], and VoiceBench AlpacaEval [Chen et al., 2024, Taori et al., 2023]. All audio is the benchmarks’ official pre-rendered audio; no query ever touches probe calibration.

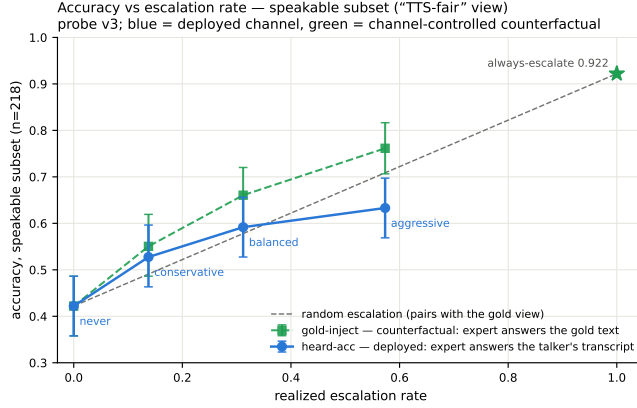


Figure 3: Live dual-view tradeoff ($n=218$ speakable subset). Blue: honest heard accuracy of the live duplex system (paired-bootstrap CIs). Green: the channel-controlled counterfactual (same sessions, escalated rows re-scored with gold-text expert answers) with its paired random reference (dashed). The view-to-view gap is the speech-channel price; the floor-to-floor gap to offline operation is the audio + streaming-answer tax.

Table 3: The frozen gate on six public speech benchmarks (live loop; realized escalation 15/30/50%). Judge: OAB / VB = official; ours = reference-anchored. Ceiling = GPT-5.5 on gold text, same judge; official = MiniCPM-o 4.5’s reported offline number. Live accuracies carry a $\pm .02$ –.03 replication floor.

pool	judge	n	escalation tier				ceiling	AUC	official
			0%	15%	30%	50%			
Speech TriviaQA	OAB	250	.664	.732	.800	.860	.972	.789	.755
Speech Web Questions	OAB	250	.57	.628	.680	.732	.844	.785	.702
Llama Questions	OAB	250	.84	.904	.904	.948	.928	.806	—
SD-QA	ours	200	.51	.600	.720	.795	.930	.792	—
Reasoning QA (zh)	ours	202	.59	.653	.713	.772	.871	.683	—
AlpacaEval (1–5)	VB	199	3.99	3.96	4.23	4.35	4.96	—	4.8

The only per-pool adaptation is *label-free*: tier thresholds are quantiles of the probe’s own score distribution on each pool, consuming no labels and no judge calls. Where an official protocol exists we score with it (the OpenAudioBench judge; VoiceBench’s 1–5 judge), and our chat-mode control reproduces the official anchors on-scale.

The frozen gate improves all five QA benchmarks (Table 3): thresholds fire at sane rates, every curve rises monotonically, and 50% escalation lifts the live floor by +.11 to +.29. On Speech TriviaQA the 9B system reaches **.860 live** — against its own official offline .755 and .629 for Qwen3-Omni-30B [Xu et al., 2025b] at $3.3\times$ the parameters: routing beats scaling the speech backbone. Transfer AUC holds a .68–.81 band on every pool against .879 internally, and the cleanest evidence that the signal is difficulty rather than surface is cross-lingual: an English-only calibration expansion moved *Chinese* Reasoning QA from AUC .621 to .683, the largest per-pool gain. On Llama Questions — single-type, so no shortcut to exploit — **selective escalation matches the expert at half the expert calls**: the probe’s split is decisive (kept-local .976 vs. escalated .696, $z=6.46$), the expert’s whole advantage lands where the probe sends it (.696→.888, $p<.0001$), and always-escalate buys nothing measurable over selective; matched-rate precision reaches 95% of the base-rate cap (Appendix G).

Two honest negatives. **AlpacaEval**: the gate does not beat random — open-ended instruction following has no “missing fact” event for a retrieval-failure detector to read; escalated rows still gain, but random picks would harvest the same gain. **Web Questions and SD-QA are flat** under the deployed-probe refit: on Web Questions strict alias-based labels decouple “wrong” from “uncertain”

(even the gold-text ceiling is .844). The external channel cost is small (heard-vs-gold gap .035–.056 at aggressive vs. .175 internally), and SD-QA’s escalated subset converts at .82–.84 on real dialectal speech. As a side effect of routing, selective escalation can also *reduce* median and tail latency by replacing the slowest wrong local decodes with short expert round-trips (Reasoning QA P90 13.3→10.9 s); Appendix G analyzes the latency shapes.

8 Discussion

What the signal detects. The pre-answer probe is specifically a *retrieval-failure* detector: it is strongest on missing-knowledge and long-tail-trap failures, where the empty lookup is an observable event in the model’s state. Execution failures (getting stuck mid-derivation) surface in *decode-time* signals instead, and metacognitive failures — questions with false premises, where no failure event occurs in the model’s experience — are invisible to every signal we tested, query-surface routers included (fire rate .18 vs. trap’s .90; Appendix F). This is the honest scope statement for pre-answer gating, and it motivates adding a decode-time check behind the pre-answer gate.

Limitations. (i) One phrasing per $p(\text{True})$ variant. (ii) Real-speech validation covers scripted read speech in one dialect. (iii) Formula-symbolic content is out of scope for the audio arm (destroyed by any TTS render — an input-side limit distinct from the transcription-uplink loss). (iv) The second-family replication covers mid-band readability and frozen-recipe transfer, not the full matched-pair attribution or the live loop. (v) Labels come from one judge family, bounded by a stronger re-judge ($\text{deltas} \leq .010$), without human labels. (vi) The live sweep is one session per query per tier ($\pm .02$ –.03 replication floor); the native duplex listen/speak head remains unused — probe validity under the true duplex loop is an open experiment.

9 Conclusion

Duplex speech models carry a competence signal their standard readouts miss: a mid-network linear read that survives the readout damage, transfers from text calibration to speech, and lands tens of milliseconds before the first output token. Deployed live, it turns a 9B talker into a system that knows when to hand off; frozen, it lifts public speech benchmarks past models three times its size. Reading the right layer makes real-time escalation possible.

References

- Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqui, and Mausam. AutoMix: Automatically mixing language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2310.12963.
- Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. arXiv:2304.13734.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from question-answer pairs. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2013. WebQuestions.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2212.03827.
- Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.
- Qian Chen, Yafeng Chen, Yanni Chen, et al. MinMo: A multimodal large language model for seamless voice interaction. *arXiv preprint arXiv:2501.06282*, 2025.

285 Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench:
286 Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024.

287 Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, et al. Qwen2-Audio technical
288 report. *arXiv preprint arXiv:2407.10759*, 2024.

289 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al.
290 Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. GSM8K.

291 Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick
292 Wendell, Matei Zaharia, and Reynold Xin. Free dolly: Introducing the world’s first truly open
293 instruction-tuned LLM. Databricks blog, 2023. databricks-dolly-15k.

294 Junbo Cui, Bokai Xu, Chongyi Wang, Tianyu Yu, Weiye Sun, Yingjing Xu, et al. MiniCPM-o 4.5:
295 Towards real-time full-duplex omni-modal interaction. *arXiv preprint arXiv:2604.27393*, 2026.

296 DeepSeek-AI, Daya Guo, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via
297 reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

298 Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou,
299 Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dia-
300 logue. *arXiv preprint arXiv:2410.00037*, 2024.

301 Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks
302 V. S. Lakshmanan, and Ahmed Hassan Awadallah. Hybrid LLM: Cost-efficient and quality-
303 aware query routing. In *International Conference on Learning Representations (ICLR)*, 2024.
304 arXiv:2404.14618.

305 Fahim Faisal, Sharlina Keshava, Md Mahfuz ibn Alam, and Antonios Anastasopoulos. SD-QA:
306 Spoken dialectal question answering for the real world. In *Findings of the Association for Com-
307 putational Linguistics: EMNLP 2021*, 2021. arXiv:2109.12072.

308 Qingkai Fang, Shoutao Guo, and Yang Feng. BayLing-Duplex: Native full-duplex speech dialogue
309 with a single autoregressive LLM. *arXiv preprint arXiv:2606.14528*, 2026.

310 Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large
311 language models using semantic entropy. *Nature*, 630:625–630, 2024.

312 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
313 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
314 *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks*, 2021.
315 arXiv:2103.03874.

316 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
317 Keutzer, and Shriyash Kaustubh Upadhyay. RouterBench: A benchmark for multi-LLM routing
318 systems. *arXiv preprint arXiv:2403.12031*, 2024.

319 Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale dis-
320 tantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual
321 Meeting of the Association for Computational Linguistics (ACL)*, 2017.

322 Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez,
323 Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language mod-
324 els (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

325 Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. Se-
326 mantic entropy probes: Robust and cheap hallucination detection in LLMs. *arXiv preprint
327 arXiv:2406.15927*, 2024.

328 Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for
329 uncertainty estimation in natural language generation. In *International Conference on Learning*
330 *Representations (ICLR)*, 2023. arXiv:2302.09664.

331 Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative
332 decoding. In *International Conference on Machine Learning (ICML)*, 2023. arXiv:2211.17192.

333 Chia-Hsuan Li, Szu-Lin Wu, Chi-Liang Liu, and Hung-yi Lee. Spoken SQuAD: A study of miti-
334 gating the impact of speech recognition errors on listening comprehension. In *Interspeech*, 2018.
335 arXiv:1804.00320.

336 Hao Li, Yiqun Zhang, Zhaoyan Guo, Chenxu Wang, Shengji Tang, Qiaosheng Zhang, Yang Chen,
337 Biqing Qi, Peng Ye, Lei Bai, Zhen Wang, and Shuyue Hu. LLMRouterBench: A massive bench-
338 mark and unified framework for LLM routing. *arXiv preprint arXiv:2601.07206*, 2026.

339 Tianpeng Li, Jun Liu, Tao Zhang, Yuanbo Fang, Da Pan, Mingrui Wang, et al. Baichuan-Audio:
340 A unified framework for end-to-end speech interaction. *arXiv preprint arXiv:2502.17239*, 2025.
341 Introduces OpenAudioBench.

342 Borui Liao, Yulong Xu, Jiao Ou, Kaiyuan Yang, Weihua Jian, Pengfei Wan, and Di Zhang. Flex-
343 Duo: A pluggable system for enabling full-duplex capabilities in speech dialogue systems. *arXiv*
344 *preprint arXiv:2502.13472*, 2025.

345 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
346 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International*
347 *Conference on Learning Representations (ICLR)*, 2024. arXiv:2305.20050; source of the MATH-
348 500 subset.

349 Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in
350 words. *Transactions on Machine Learning Research*, 2022. arXiv:2205.14334.

351 David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and
352 accuracy by learning to defer. In *Advances in Neural Information Processing Systems (NeurIPS)*,
353 2018.

354 Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language
355 model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
356 arXiv:2310.06824.

357 Eliya Nachmani, Alon Levkovitch, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh
358 Mariooryad, Ehud Rivlin, RJ Skerry-Ryan, and Michelle Tadmor Ramanovich. Spoken question
359 answering and speech continuation using spectrogram-powered LLM. In *International Confer-*
360 *ence on Learning Representations (ICLR)*, 2024. arXiv:2305.15255; Spectron, source of the
361 Llama Questions dataset.

362 Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden
363 Tomasello, Robin Algayres, Benoît Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux. Gen-
364 erative spoken dialogue language modeling. *Transactions of the Association for Computational*
365 *Linguistics*, 2023. arXiv:2203.16502.

366 Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez,
367 M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data.
368 *arXiv preprint arXiv:2406.18665*, 2024.

369 Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, and Yonatan Belinkov.
370 LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In
371 *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2410.02707. TODO
372 verify author list.

373 Pablo Riera, Pablo Brusco, Cristina Kuo, Marcelo Sancinetti, and S. R. K. Branavan. Synchroniza-
374 tion and turn-taking in full-duplex speech dialogue models. *arXiv preprint arXiv:2605.20356*,
375 2026.

376 Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and
377 Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information*
378 *Processing Systems (NeurIPS)*, 2022. arXiv:2207.07061.

379 Tanya Stivers, N. J. Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann,
380 Gertie Hoymann, Federico Rossano, Jan Peter de Ruiter, Kyung-Eun Yoon, and Stephen C. Levin-
381 son. Universals and cultural variation in turn-taking in conversation. *Proceedings of the National*
382 *Academy of Sciences*, 106(26):10587–10592, 2009.

383 Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy
384 Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following LLaMA model.
385 GitHub repository, 2023. zh split: shibing624/alpaca-zh.

386 Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan
387 Sung, Denny Zhou, Quoc Le, and Thang Luong. Freshllms: Refreshing large language models
388 with search engine augmentation. *arXiv preprint arXiv:2310.03214*, 2023.

389 Xiong Wang, Yangze Li, Chaoyou Fu, Lei Xie, Ke Li, Xing Sun, and Long Ma. Freeze-Omni:
390 A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv preprint*
391 *arXiv:2411.00774*, 2024a.

392 Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, et al. MMLU-Pro: A more robust and
393 challenging multi-task language understanding benchmark. In *Advances in Neural Information*
394 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2024b. arXiv:2406.01574.

395 Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese,
396 John Schulman, and William Fedus. Measuring short-form factuality in large language models.
397 *arXiv preprint arXiv:2411.04368*, 2024. SimpleQA.

398 Donghang Wu, Haoyang Zhang, Chen Chen, Tianyu Zhang, Fei Tian, Xuerui Yang, Gang Yu, Hexin
399 Liu, Nana Hou, Yuchen Hu, and Eng Siong Chng. Chronological thinking in full-duplex spoken
400 dialogue language models. In *Proceedings of the 27th Annual Meeting of the Special Interest*
401 *Group on Discourse and Dialogue (SIGDIAL)*, 2026. arXiv:2510.05150.

402 Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, et al. Qwen2.5-
403 Omni technical report. *arXiv preprint arXiv:2503.20215*, 2025a.

404 Jin Xu, Zhifang Guo, Hangrui Hu, Yunfei Chu, Xiong Wang, et al. Qwen3-Omni technical report.
405 *arXiv preprint arXiv:2509.17765*, 2025b.

406 Leyang Xue, Yao Fu, Luo Mai, and Mahesh K. Marina. HybridServe: Efficient serving of large AI
407 models with confidence-based cascade routing. *arXiv preprint arXiv:2505.12566*, 2025.

408 Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong,
409 and Jie Tang. GLM-4-Voice: Towards intelligent and human-like end-to-end spoken chatbot.
410 *arXiv preprint arXiv:2412.02612*, 2024.

411 Dengjia Zhang, Alexander Martin, William Jurayj, Kenton Murray, Benjamin Van Durme, and Reno
412 Kriz. Unified multimodal uncertain inference. *arXiv preprint arXiv:2604.08701*, 2026.

413 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang,
414 Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica.
415 Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information*
416 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2023. arXiv:2306.05685.

417 A Model roster

Table 4: Model roster. “Pair” groups a fine-tune with its raw backbone control; the comparison statistic is always the within-pair difference.

model	type	pair	role
MiniCPM-o 4.5	omni + duplex FT	Qwen3-8B	system target
Qwen3-8B	raw	—	pair control
MiniCPM-o 2.6	duplex FT	Qwen2.5-7B	duplex replication
Qwen2.5-Omni-7B	streaming omni FT	Qwen2.5-7B	omni control
Qwen2.5-7B	raw	—	pair control
Qwen2-Audio-7B	audio FT	Qwen2-7B	audio-FT control
Freeze-Omni	frozen LLM + adapter	Qwen2-7B	adapter control
Qwen2-7B	raw	—	pair control
Mistral-7B-v0.3	raw	—	backbone diversity
GLM4-9B-chat	raw	—	backbone diversity

418 B Extended signal and readout analyses

419 **ROC curves.** Figure 4 shows the calibration-split ROC for the signals of Table 1.

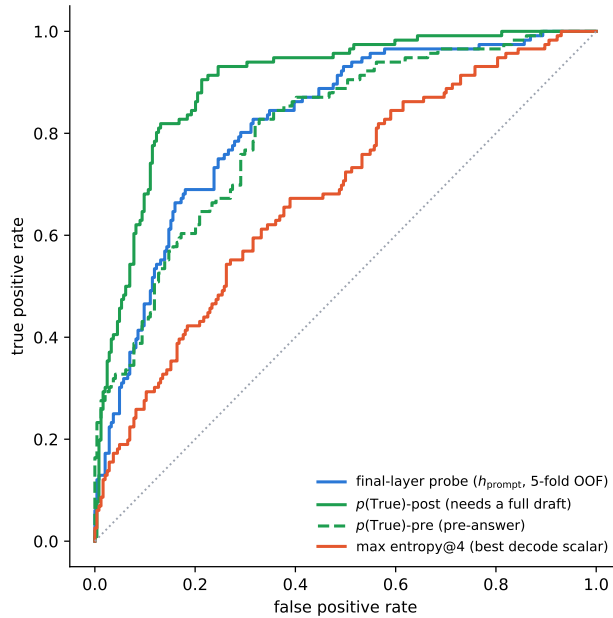


Figure 4: ROC on the calibration split. $p(\text{True})$ -post leads but needs a full drafted answer; the probes and $p(\text{True})$ -pre are pre-answer.

420 **$p(\text{True})$ backbone dependence.** Across backbones, post-draft $p(\text{True})$ beats the trained final-
 421 layer probe on four of five raw/duplex models tested; the exception (GLM4-9B-chat, .685 vs. probe
 422 .798) shows the pattern depends on the backbone’s self-evaluation quality.

423 **Matched pairs.** Table 5: with label coverage matched by subsampling the raw control to the
 424 duplex model’s per-pool fail rates, the raw backbone keeps LOPO math at .825/.809 while duplex
 425 models sit at chance — label coverage is ruled out; the representation changed. Degradation is
 426 ordered raw > omni-streaming > duplex. Table 6 gives the by-depth summary behind Figure 2: the
 427 within-pair deficit at the best mid-layer is only $-.03$ vs. $-.59$ at the final layer; mean-pooled probes

survive to the final layer on duplex models; duplex damage is severe but local, omni-streaming damage mild but diffuse.

Table 5: LOPO transfer (AUC, held-out pool) of the final-layer last-token probe, Qwen2.5-7B family, label-coverage matched.

LOPO pool	Qwen2.5-7B (raw)	Qwen2.5-Omni	MiniCPM-o 2.6	[o4.5 vs Qwen3]
hard-math	0.809	0.745	0.538 $\approx$ chance	0.372 (inverts)
hard-knowledge	0.680	0.613	0.526 $\approx$ chance	0.61

Table 6: LOPO math transfer by depth (last-token probes).

model	best mid-layer	final layer	shape
Qwen3-8B (raw)	0.964 (L33/36)	0.958	plateau, no cliff
MiniCPM-o 4.5 (duplex)	0.931 (L22/36)	0.366	cliff, inverts
Qwen2.5-7B (raw)	0.893 (L21/28)	0.809	mild dip
Qwen2.5-Omni (streaming)	0.794 (L16/28)	0.746	diffuse depression
MiniCPM-o 2.6 (duplex)	0.822 (L21/28)	0.540	cliff

Modality specificity of the cliff. Under audio input the final layer does *not* invert (L35 audio LOPO math .936 vs. .366 text); a speak-mode template control on text input leaves the cliff unchanged (.362) — the operative variable is the modality of the input context, not the output mode.

Audio collapse of verbalized introspection. Table 7. Three controlled arms pin the mechanism on MiniCPM-o 4.5: *not perception* (ASR audit WER .074; $p(\text{True})$ on the model’s own transcript snaps back to .074; the well-heard subset still collapses); *not a context prior* (irrelevant filler audio does not inflate p_{yes}); *binding* (duplicating the question as text alongside the audio fully restores introspection, .034). The verbal instance-check runs over text-token pathways that audio embeddings do not feed, while a linear probe reads the same information off the same forward pass at .93+. The failure implies its own cheap fix, confirmed twice: re-present the query as text before asking (“repeat-then-judge”).

Table 7: Pre-answer introspection under audio input. Trap p_{yes} = mean stated probability of answering correctly on a 100%-fail pool (lower is better).

model	FT type	trap p_{yes} : text $\rightarrow$ audio	audio $p(\text{True})$ -pre AUC
MiniCPM-o 4.5	duplex	.055 $\rightarrow$.556 (collapse)	.786 (trap dead)
MiniCPM-o 2.6	duplex	.196 $\rightarrow$.345	.491 $\approx$ chance
Qwen2.5-Omni	omni-streaming	.279 $\rightarrow$.213 (intact)	.727
Freeze-Omni	frozen + adapter	.165 $\rightarrow$.112 (intact)	.612 (capability caveat)

Real-speech validation. The matched-pair design rerun on 200 real human recordings (VoiceBench SD-QA, USA split) replicates all three audio-side findings: modality tax (.400 $\rightarrow$.450), audio overconfidence (paired p_{yes} shift +.089; failure subset .415 $\rightarrow$.581), and the layer structure (text-calibrated read on real-speech audio holds .76–.80 at L22–L26, peak .800 at L25 — transfer across both content and modality).

Synthesis. End-to-end multimodal training *builds* the shared mid-layer core; duplex-style training additionally *damages* the readouts; omni-streaming gets both right; an adapter alone gets neither (probe readout text-fragile, verbal readout audio-fragile; in both cases knowledge survives and a readout breaks, and in both cases the breakage is duplex-fine-tune-specific).

The deployed probe. The mechanistic probe is a single-position L22 read. The deployed probe adds two zero-latency pooled positions (an 8-token tail mean and the running user-audio mean),

refit on a calibration mix widened to 2,310 public-benchmark queries disjoint from every evaluation benchmark; a pre-registered guard held (frozen internal test AUC .860→.879) and feature selection never saw the externals. In the mid-layer band the in-mix tradeoff optimum is broad (L20–L30) and the mid-layer probe’s in-distribution area advantage is modest — its real advantage is LOPO robustness, which an in-mix test cannot show.

C Second duplex family: NemotronLabs-VoiceChat-11B

The pre-registered prediction test: with the identical recipe (same audio, same judge, 600-query calibration), the layer sweep on this hybrid-Mamba duplex model (56 blocks: 27 Mamba-2 / 4 attention) peaks mid-band at L30–34 (eot-last AUC .714 vs. .693/.682 at the endpoints); stacking the same three positions lifts OOF .714→.761→.790; the frozen-methodology probe transfers to the external pools at AUC .781/.793/.754/.701 — the deployed-probe band reached with a quarter of the calibration data. Boundary notes: the backbone is much weaker at knowledge retrieval (striviaqa local .32 vs. MiniCPM .62, same judge) and English-only (the cross-lingual pool has no analog). The live-loop port and the trained-checkpoint ablation remain future work.

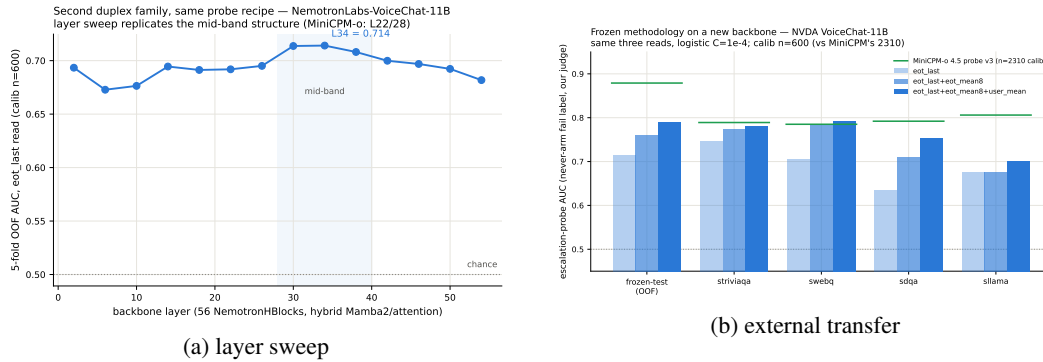


Figure 5: Second-family test: (a) the usable read again peaks mid-network; (b) frozen-recipe transfer lands in the deployed-probe AUC band.

D Query-surface router audit

A RouteLLM-style TF-IDF→LR router under identical labels reaches OOF AUC .669 (below the pool oracle’s .715); a 24-configuration sweep tops out at .689, so this is the query surface’s ceiling, not tuning. Under LOPO it collapses to .38–.57 and scores the 100%-fail trap pool at .230 — it misses the pool a router most needs. The training receipt: at the argmax threshold its out-of-fold accuracy exactly matches the majority class; the internal probes under the same accounting clear majority by 9–29 points (Figure 7). Nor is it data-starved: the identical recipe on RouterBench [Hu et al., 2024] (36,497 prompts) reaches only in-domain AUC .710 with leave-one-benchmark-out at chance; RouteLLM’s released preference-trained routers score .523/.533 zero-shot on our pool and rank the trap pool below ordinary hard pools; our router transfers to RouterBench below chance (.440). On audio labels the router improves (test .814) but still trails the mid-layer probe (.879) and cannot fire mid-stream. In distribution, the gate’s area over random (+.054) is statistically on par with the pool oracle’s (+.042; paired bootstrap delta CI [−.003, +.027], $p=.13$) — the gate’s case is transfer, not the in-distribution margin. $p(\text{True})$ tradeoff variants: Figure 8.

E Gate timing details

Per-layer truncated-forward cost vs. quality (Figure 9): the L22 hidden state is available at ~57–61% of prefill wall-time while L22 is simultaneously the in-mix quality peak; forking much earlier is both weaker and modality-specific, so the fork belongs at ~55–60% depth. If the cloud call launches at

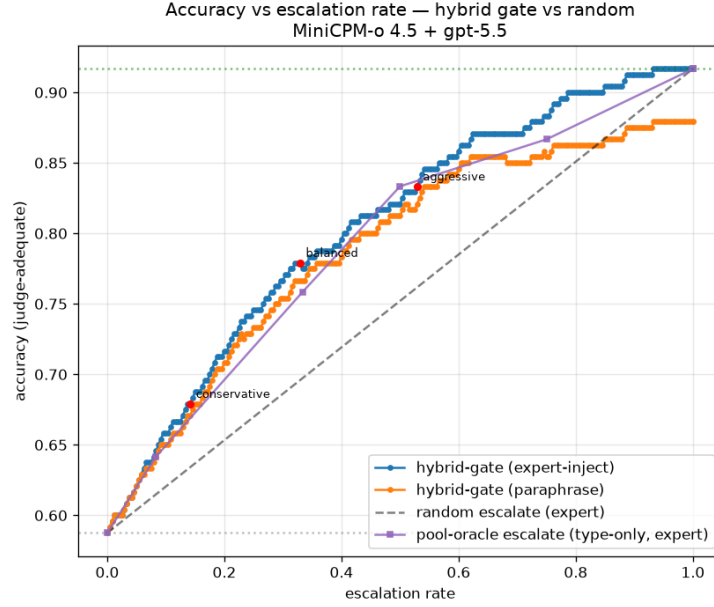


Figure 6: Accuracy vs. escalation rate, frozen test split: gated escalation vs. random-escalation and pool-oracle (type-only) baselines between the small-only and big-only endpoints.

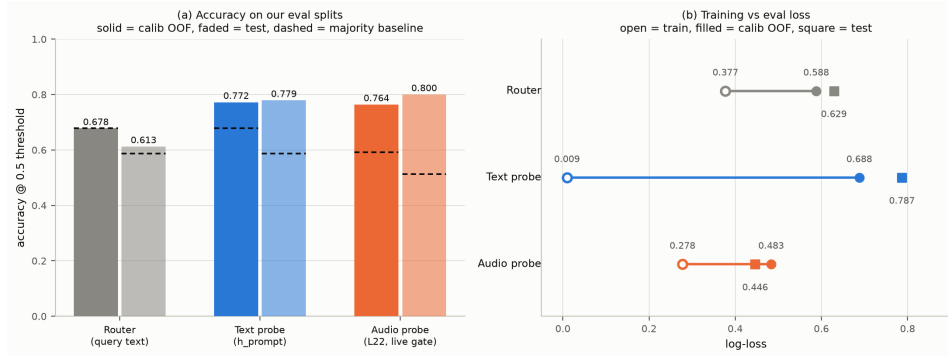


Figure 7: Identical training receipt for the external router and the two internal probes: the router never clears the majority-class baseline; both probes do, by 9–29 points.

484 the gate, the expert result arrives before the local answer finishes for 40–58% of the whole mix but
 485 only ~20–31% of *escalated* queries (they skew to short local answers); the residual wait is P50
 486 2–3 s, P90 ~27 s (Figure 10) — the stall-phrase budget the injection protocol must cover.

487 F Live-loop details

488 **Full live tables.** Table 8 (both scoring views, speakable subset and full pool) and Table 9 (measured response latency). The conservative tier *shortens* the P95 tail: its few escalations displace
 489 exactly the longest local decodes. Figure 11 joins accuracy and latency; Figure 12 replays three real
 490 sessions from recorded timers.
 491

492 **Speculative prefetch, measured and dropped.** Framed as speculative escalation (draft–verify at
 493 sentence granularity), its acceptance rate — does an expert answer to the partial transcript answer the
 494 full question — is .07/.19/.51 at 25/50/75% of the audio; at the mid-stream gate’s typical fire point




Figure 8: Offline tradeoff for the $p(\text{True})$ signals (pre-answer and post-draft).

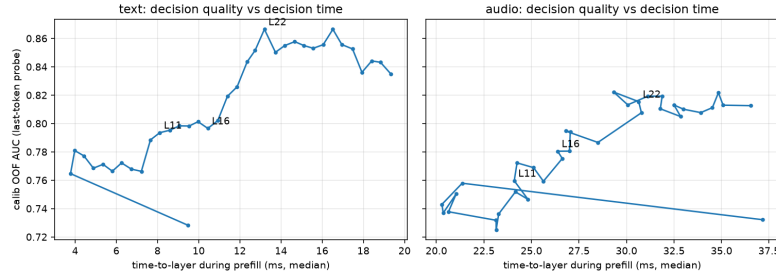


Figure 9: Fork-depth Pareto: per-layer wall-time vs. gate quality.

495 6–8 of 10 calls would be wasted. The trap pool’s floor (.00/.10/.27) shares the back-loaded-core
 496 property that breaks mid-stream probe reads.

497 **Conflicting-injection controls (172 sessions).** Fabricated conflicting expert answers by within-
 498 pool derangement, four framings: worst-case comply is .00 neutral / .53 deployed authority / .79
 499 strong-override, while *seeding* words into the talker’s mouth backfires (comply .25, lip-service .62
 500 — forced speech is treated as revisable). Math is silently re-derived rather than relayed or disputed.
 501 Model-wrong $\times$ wrong-injection swallow rate is .64: expert quality is the accuracy ceiling of the
 502 cascade.

503 **Uplink diagnostics.** On the escalated population the expert answering the talker’s own transcript
 504 scores .585 vs. .871 on gold text. Prompt-level rescues on the self-transcript are no-ops (the dam-
 505 aging errors are internally-coherent entity substitutions); replacing the ears helps: open-weights
 506 Whisper recovers little (+4 pp) but a stronger cloud ASR on the same audio recovers 38% of the
 507 gap (.585 $\rightarrow$.694, McNemar $p=.007$) — diagnostic only (its critical-path latency and external-pool
 508 survival are unmeasured). A same-model gold-text control with an audio-capable expert shows the
 509 audio channel itself is lossless for speakable content: the leak is the talker’s transcription, not the

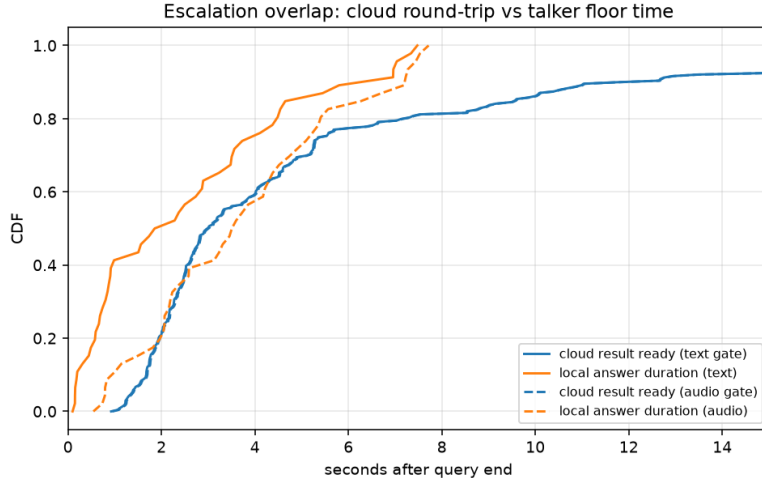


Figure 10: Cloud round-trip vs. local answer duration: probability the expert result arrives before the local answer ends.

Table 8: Live streaming sweep, paired bootstrap 95% CIs ($B=10^4$). Top: speakable subset ($n=218$); bottom: full pool ($n=240$). Gold-inject = escalated rows re-scored with gold-text expert answers; channel cost = paired difference. * CI excludes 0; † CI lower bound at zero.

tier	esc	heard-acc	Δ vs. floor	gold-inject	channel cost
never	0%	.422 [.36,.49]	—	.422	—
conservative	14%	.528 [.46,.59]	+106*	.550	+.023†
balanced	31%	.592 [.53,.66]	+170*	.661	+.069*
aggressive	57%	.633 [.57,.70]	+211*	.761	+.128*
always (synth)	100%	—	—	.922 [.89,.95]	—
never (full)	0%	.383 [.32,.45]	—	.383	—
conservative	16%	.483 [.42,.55]	+100*	.525	+.042*
balanced	35%	.554 [.49,.62]	+171*	.654	+.100*
aggressive	61%	.596 [.53,.66]	+212*	.771	+.175*
always (synth)	100%	—	—	.917 [.88,.95]	—

510 synthesis or the relay (relay drops: 5 cases of 108; corrupted transcripts dominate the escalated
511 heard-failures, first-generation sweep).

512 **Boundary query classes.** *False premises:* the model fails substantially (.63 text / .47 audio) and
513 the expert would help (.80), but the gate is blind (fire@balanced .18 vs. trap’s .90; no score separation)
514 — answering along a false premise does not feel like inability. *Real-time queries* (“what is
515 NVDA trading at”): the gate is not blind, the class was simply absent from every static training fam-
516 ily; extending calibration with FreshQA [Vu et al., 2023] (fast-changing labeled escalate a priori,
517 never-changing as in-family controls) raises held-out fast-changing fire rates to .80/.93/1.0 across
518 tiers with frozen-test AUC unchanged (.877 vs. .879). One calibration subtlety: tier budgets must
519 remain quantiles of the deployment-mix scores, or the new capability is spent on its own threshold
520 shift.

521 **Demonstration system.** The pipeline runs as an interactive web demo on one H100: live micro-
522 phone audio, end-of-turn read (~ 54 ms) against the balanced threshold, escalated turns stall, consult
523 the expert, and *speak* the relayed answer through the model’s native TTS head (probe context and
524 thresholds untouched). Two demo-only extensions: a web-search tool for the expert (real-time
525 queries then return live values), and backchannel-aware barge-in — speech over the talker ducks
526 playback, a fast ASR pass returns the floor on backchannels (“ok”, “mm-hm”; an English+Chinese

Table 9: Measured response latency (s), $n=240$ per arm (latency is channel-independent). P99 on $n=240$ is indicative only.

tier	mean	P50	P95	P99	Δ mean	Δ P50	Δ P95	Δ P99
never	4.6	2.0	15.8	17.8	—	—	—	—
conservative	4.6	2.6	12.8	23.0	+0.0	+0.6	-3.0	+5.2
balanced	5.8	3.5	15.9	32.7	+1.2	+1.5	+0.1	+14.9
aggressive	6.6	4.4	18.5	33.0	+2.0	+2.4	+2.7	+15.2

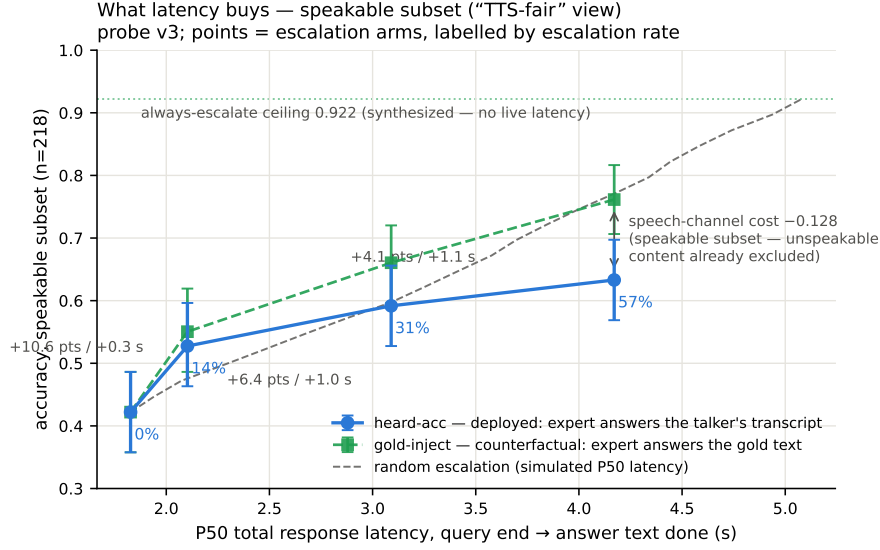


Figure 11: Accuracy vs. per-arm P50 latency (speakable subset), both scoring views, with the simulated-P50 random reference.

lexicon) or commits the interrupt, aborting generation at chunk granularity and seeding the next turn. All measured arms keep the tool-free expert and turn-based loop; the demo will be made available upon publication.

G Transfer latency shapes

Matched-rate precision. Precision at a fixed escalation rate is capped at base-fail/rate: on Llama Questions (base fail .16) the 50% cap is .32 and the deployed probe reaches .304 — 95% of the cap (recall .93). Quoting raw precision would misread the strongest pool as the weakest.

The median left-fold. On easy pools per-arm P50 is non-monotonic (Llama Questions 1.52→1.17→1.60/1.42 s): a verified median-of-a-mixture mechanism, not noise. The probe preferentially escalates the queries whose *local* decode is longest — on retrieval pools wrong answers hedge at length (decode time and answer length $r=.97$), so a wrongness-selecting probe strips the slow tail as a side effect (escalated-at-15% ids: never-arm local P50 1.711 s vs. 1.198 s kept; local accuracy .32 vs. .73). Controls: the probe does not key on length ($\text{corr}(\text{score}, \text{length}) = +.05$; ≈ 0 conditioned on wrongness), and on SD-QA — where wrong answers are short — the same probe selects wrongness identically and produces no fold. The mean is monotonic (1.65→2.06 s); kept-local P50 falls 1.52→0.43 s across tiers (Figure 13). On Reasoning QA the shape inverts usefully: the local chain-of-thought is *spoken*, so every reasoning token enters the response clock, while the expert’s chain is hidden behind a flat ~ 3.7 s round-trip — escalation truncates the latency tail (P90 13.25→10.94 s) rather than fattening it.

Three live sessions, timestamped (balanced arm; every event time is a recorded live-loop timer)




Figure 12: Three real sessions from recorded timers: expert beats the stall phrase (seamless); expert outlives it (dead air, the P99 signature); gate stays closed (first token 0.4 s).

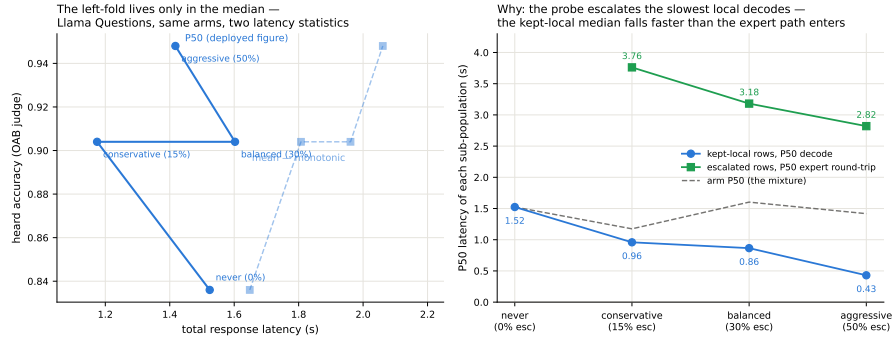


Figure 13: Decomposition of the Llama Questions latency fold: the fold exists only in the median (mean monotonic); kept-local P50 falls 1.52→0.43 s as the probe strips the slowest local decodes while escalated rows pay a flat ~3 s round-trip.

546 **AlpacaEval scale caveats.** The official 4.8 is an offline chat-mode number — our chat-mode
 547 control reproduces it at 4.86 — and the chat-vs-live gap (4.86 vs. 3.99) is the duplex loop's spoken-
 548 reply style plus a token cap, so only live arms are comparable to each other. Selected queries score
 549 3.90 on the never arm vs. 3.98 for skipped: no discrimination; escalated rows gain (3.90→4.71)
 550 because the expert writes better long-form answers — random picks would harvest the same gain.

551 H Full-pool and per-family figures

552 Figures 14 (dual-view) on the full frozen pool, and the complete external-family figure set (Ap-
 553 pendix I); the noise audit behind the ± 0.02 – 0.03 floor is Figure 15.




Figure 14: Dual-view live tradeoff, full pool ($n=240$); companion to Figure 3.




Figure 15: Replication-noise audit: same query, same audio, re-run across arms. Kept-local verdicts flip 2.3–18.8% by pool; repeated escalation 0.7–10.6%; implied paired SE .009–.028.

554 I External-benchmark figure set




Figure 16: Speech TriviaQA, Speech Web Questions, SD-QA: dual-view (left) and accuracy–latency (right).

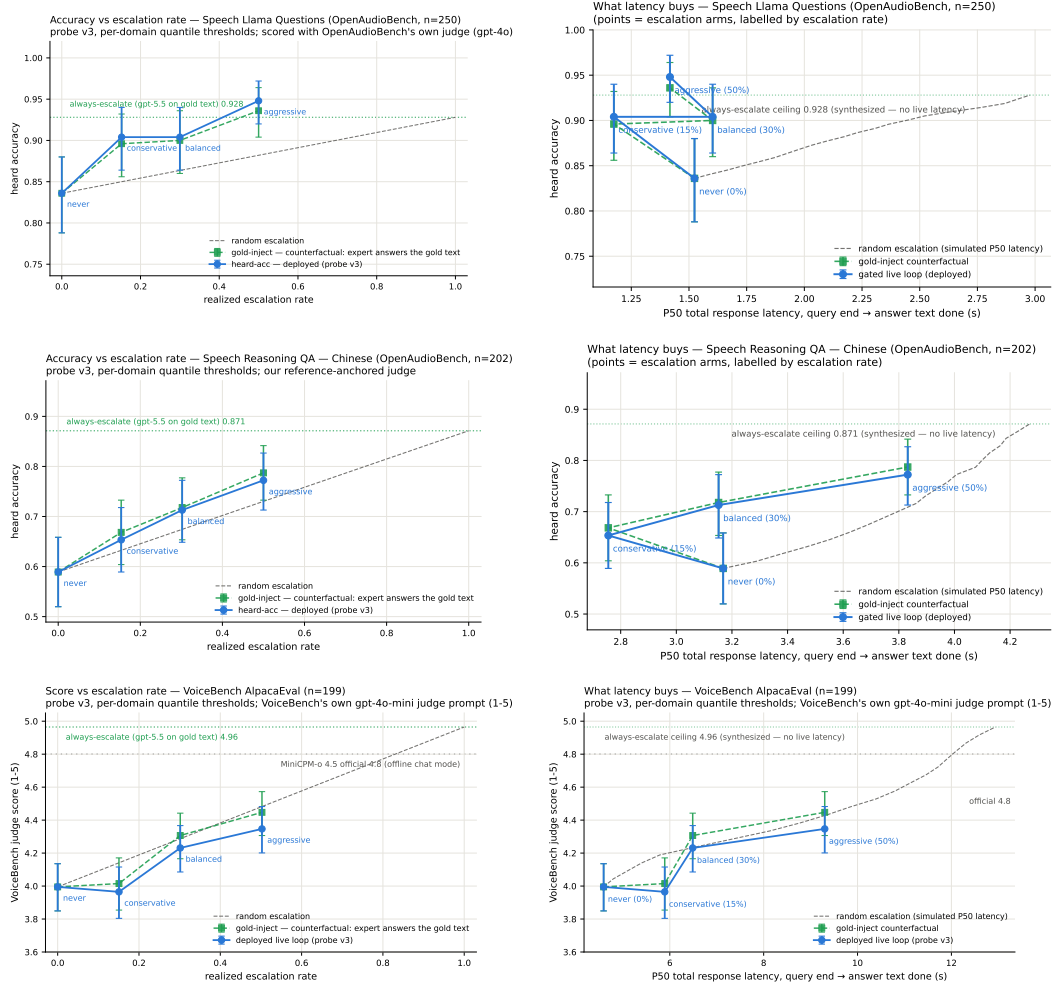


Figure 17: Llama Questions, Reasoning QA (zh), VoiceBench AlpacaEval: dual-view (left) and accuracy-latency (right).

555 J Judge and self-evaluation prompts

556 All adequacy labels use the following fixed system prompt (structured outputs with a strict {reason,
557 adequate} schema; the judge never sees which system produced the answer):

```
558     You are a strict, fair grader of a small language model's answers.  
559     You are given a QUESTION, an optional REFERENCE answer, and a  
560     CANDIDATE answer produced by a small model. Decide whether the  
561     CANDIDATE is ADEQUATE, meaning ALL of: (1) factually correct;  
562     (2) any reasoning shown is valid; (3) it actually and completely  
563     answers the question that was asked. For multiple-choice questions  
564     the candidate must land on the correct option (by letter or by  
565     unambiguously stating the option's content). A partially correct,  
566     hedged, refused, off-topic, or truncated answer is NOT adequate. If  
567     a REFERENCE is provided, grade against it. Be strict but do not  
568     penalize harmless differences in format or extra correct detail.
```

569 The verbalized self-evaluation baselines use one fixed phrasing each; $P(\text{Yes})$ is read from the first
570 generated token's distribution:

```
571     [pre] I am going to show you a question. Do NOT answer it. Judge  
572     honestly whether you yourself would answer it correctly. Question:  
573     {q} Would you answer this question correctly? Reply with exactly one  
574     word: Yes or No.  
575     [post] Here is a question and a proposed answer. Question: {q}  
576     Proposed answer: {a} Is the proposed answer correct? Reply with  
577     exactly one word: Yes or No.
```

578 K Reproducibility

579 All sampling and splits use a fixed seed (42); the calibration/test split is frozen once and the test split
580 is touched only by the evaluation scripts. Environment pins: torch 2.8.0; transformers 4.51.0 for
581 MiniCPM-o 4.5, with per-model pins for comparison backbones; all experiments run on single-H100
582 instances. Probe artifacts (weights, feature configuration, per-pool label-free thresholds), the judge
583 prompt above, all trace files, and the per-experiment execution log — including per-run GPU and
584 API spend, on the order of \$500 total — ship in the project repository, released upon publication.

585 **TODO (process only — delete before submission)**

- 586 • [TODO: P1 — DONE 2026-08-25 (external review applied): main text now ends
587 within 8 pages (refs start mid-page 8). Restructure: 3 contributions + hero figure;
588 merged signal section; live 6pp → 1.2pp; router audit 3 lines; title dropped “Zero-
589 Training”; causal claims softened (“consistent with”); no version numbers / artifact
590 filenames in main text; everything cut lives in the appendix (A–L) or git history]
- 591 • [TODO: P2 — incorporate peer + meeting feedback]
- 592 • [TODO: P3 — submission day (8-28): delete this section; final anonymization grep
593 (names, rhe9527, modal URLs, gmail); decide public-repo → private; OpenReview
594 form + upload (deadline 8-29 AoE, moderation hard limit Fri 8-28)]